



# Quarkus 3

## The Road to Loom for Cheaper, Faster, and Easier Concurrent Apps



Project Loom

# Daniel Oh



X @danieloh30

► @danieloh30

Github danieloh30



- Developer Advocate at Red Hat
  - Cloud Native Runtimes
  - Serverless, Service Mesh, and GitOps
- JAVA Champion
- Cloud Native Computing Foundation (CNCF) Ambassador
- Advisory Board Member of Global Skill Development Council
- Opensource.com Correspondents
- Public Speaker & Published Author



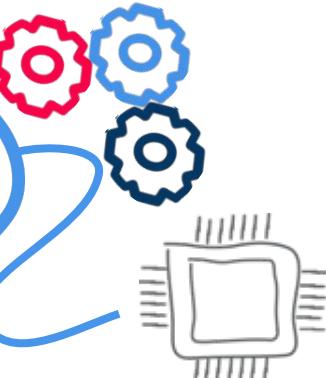
# A about threads

Your program runs  
inside a process

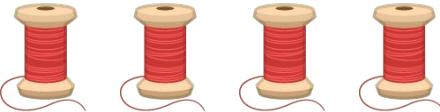
Executed code  
runs on a  
process thread

The OS schedules  
threads to run them on  
CPU

Your program /  
framework can spawn  
new process threads



# Imperative



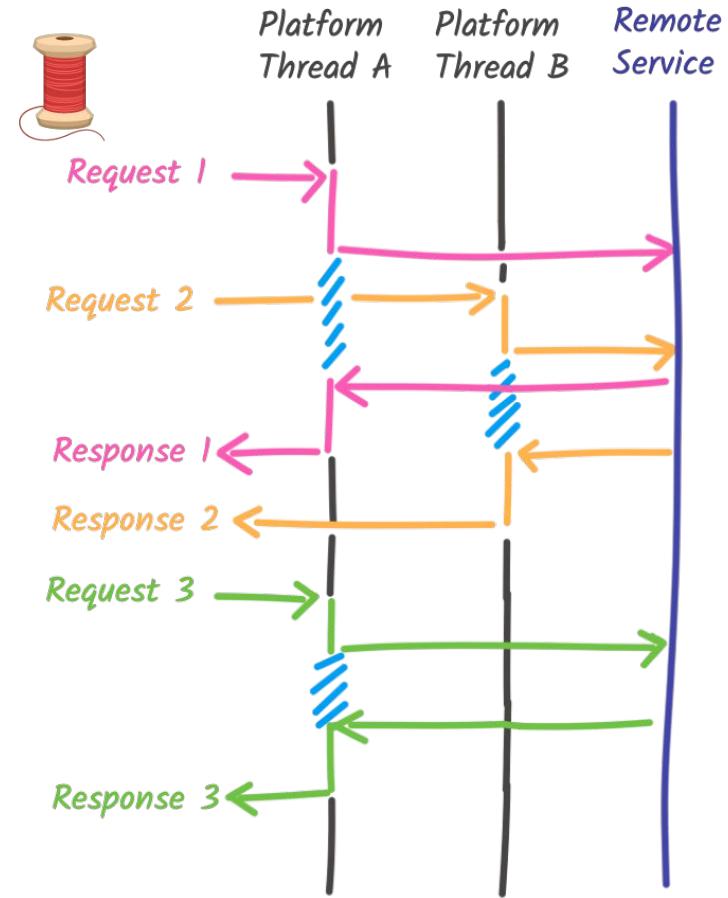
```
@Path("/greetings")
public class ImperativeApp {

    @RestClient RemoteService service;

    @GET
    public String process() {
        // Runs on a worker (platform) thread because the
        // method uses a synchronous signature

        // `service` is a rest client, it executes an I/O operation
        // (send an HTTP request and waits for the response)
        var response = service.greetings(); // Blocking, it waits for the response
                                            // The OS thread is also blocked

        return response.toUpperCase();
    }
}
```



Thread blocked

# Reactive

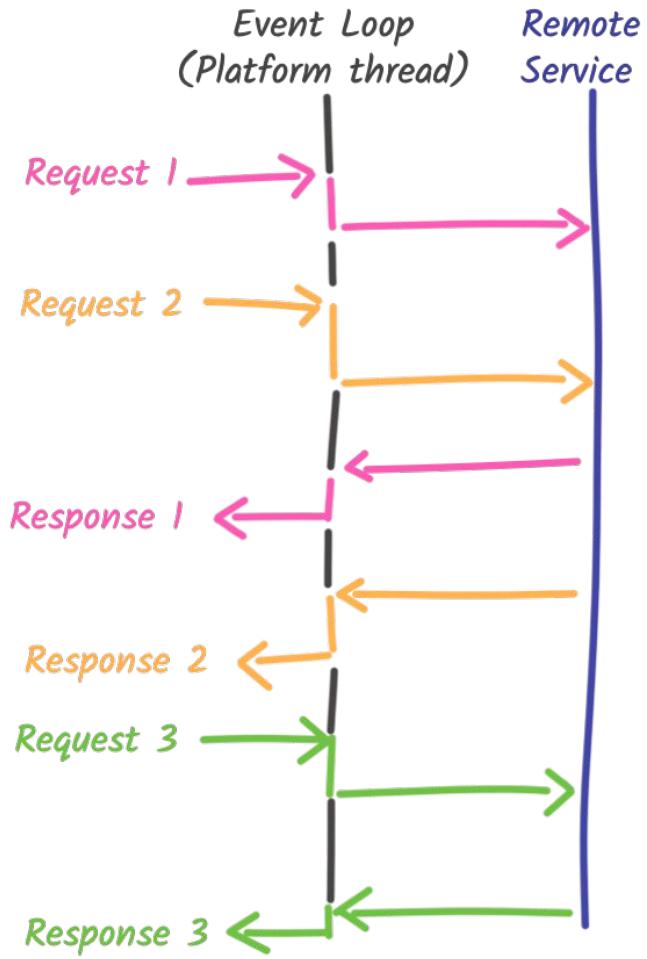


```
@Path("/greetings")
public class ReactiveApp {

    @RestClient RemoteService service;

    @GET
    public Uni<String> process() {
        // Runs on an event loop (platform) thread because the
        // method uses a asynchronous signature

        // `service` is a rest client, it executes an I/O operation
        // but this time it returns a Uni<String>, so it's not blocking
        return service.asyncGreetings() // Non-blocking
            .map(resp -> {
                // This is the continuation, the code executed once the
                // response is received
                return resp.toUpperCase();
            });
    }
}
```



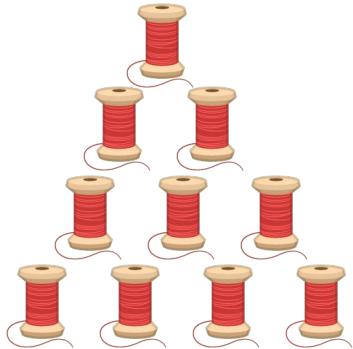
*Reactive Golden Rule:*

*> Do not block the I/O thread.*



# Imperative

- Simple development model
- Expensive: 1Mb / thread + 10  $\mu$ s
- Limited concurrency



# Reactive

- More complex development model
- Resource efficient
- High **concurrency**



*Virtual Threads (Loom):  
The pros without the cons!*



# So, what about Virtual Threads (Loom)?

## Imperative

- Simple development model
- Expensive: 1Mb / thread + 10 µs
- Limit concurrency

## Virtual Threads

- Simple development model
- Resource efficient
- High concurrency

## Reactive

- More complex development model
- Resource efficient
- High concurrency



# Virtual Threads: 1 ↪ , n

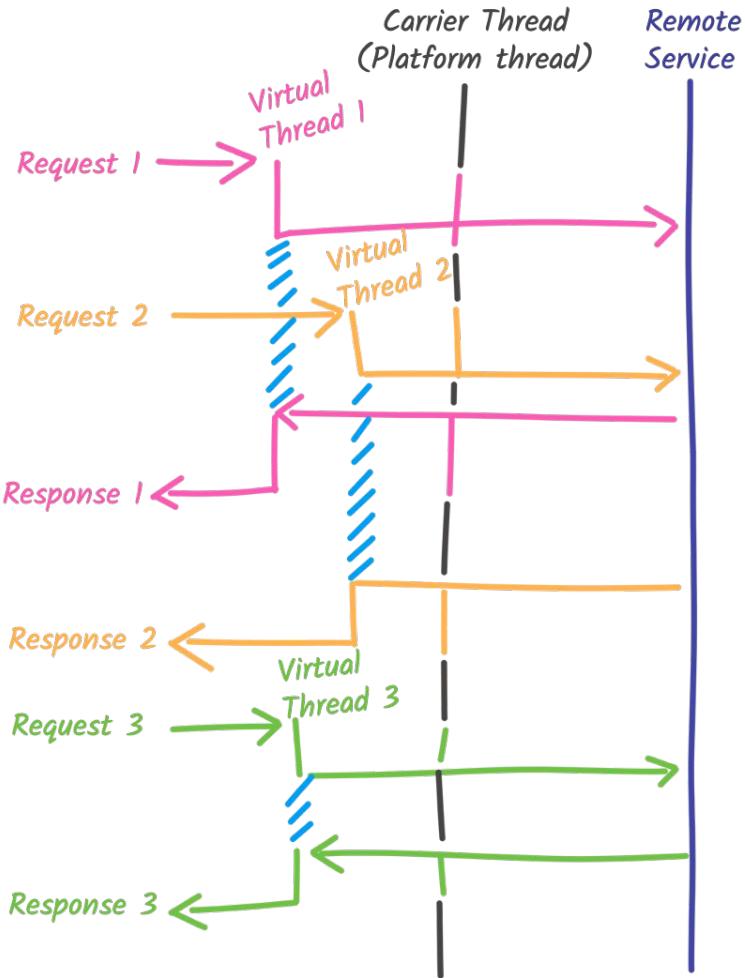
```
@Path("/greetings")
public class VirtualThreadApp {

    @RestClient RemoteService service;

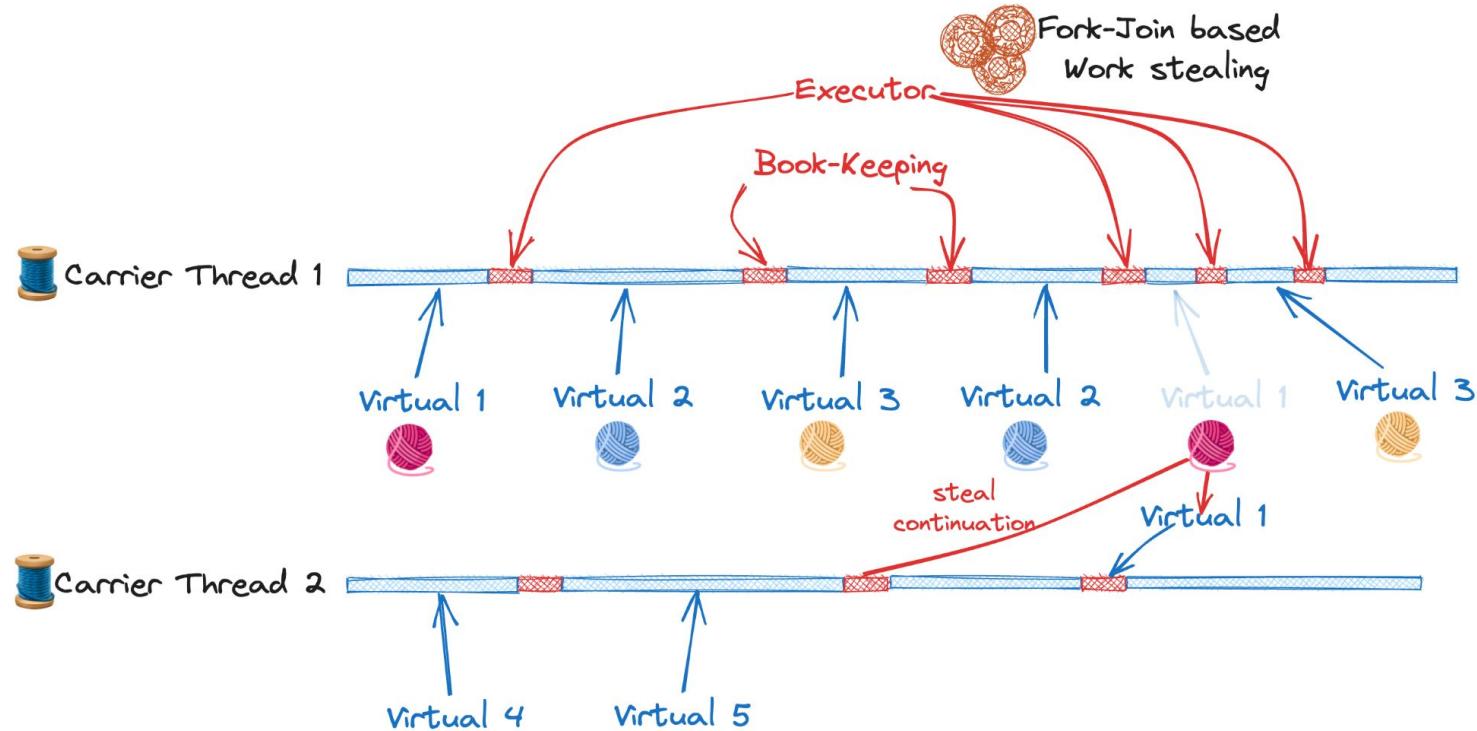
    @GET
    @RunOnVirtualThread
    public String process() {
        // Runs on a virtual thread because the
        // method uses the @RunOnVirtualThread annotation.

        // `service` is a rest client, it executes an I/O operation
        var response = service.greetings(); // Blocking, but this time, it
                                            // does neither block the carrier thread
                                            // nor the OS thread.
                                            // Only the virtual thread is blocked.

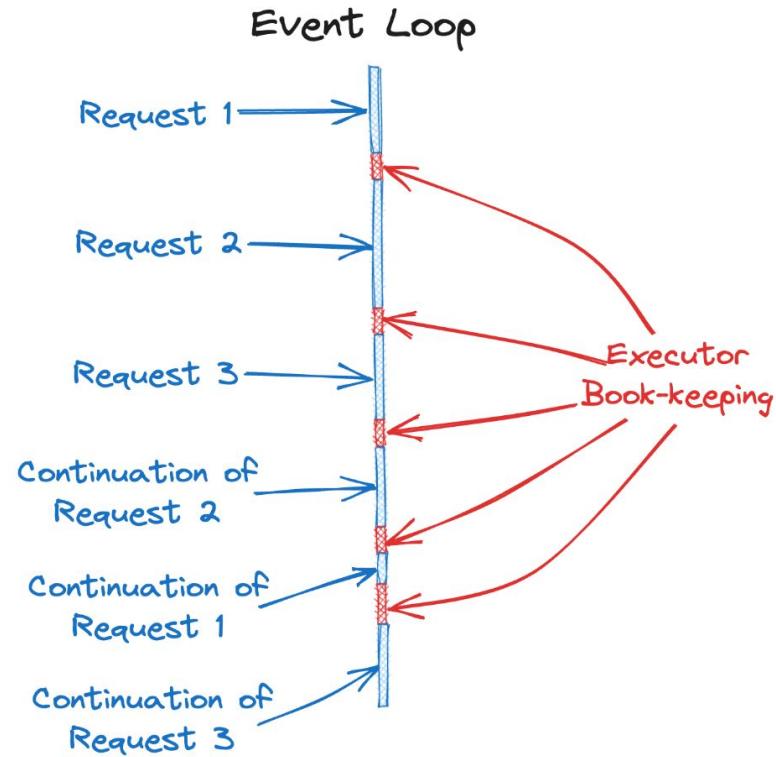
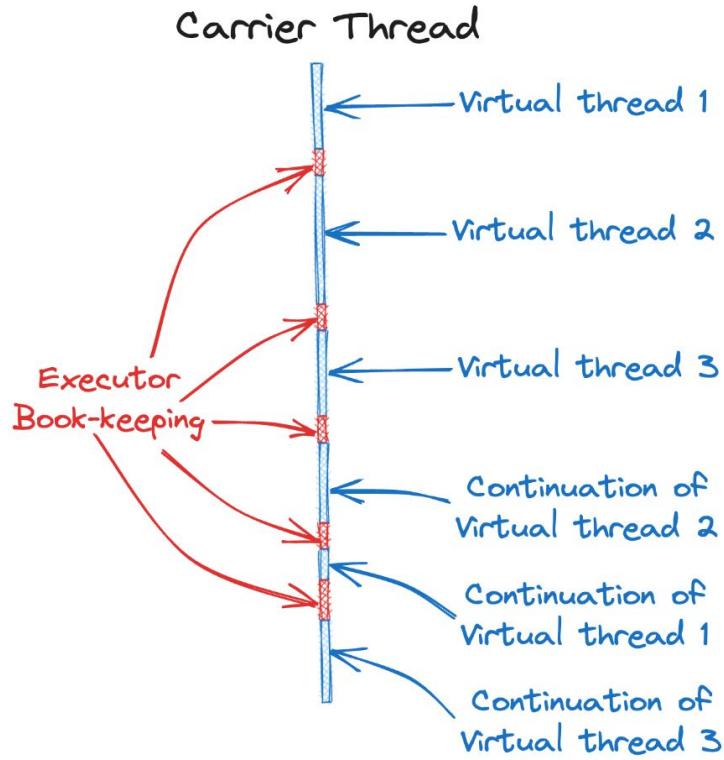
        return response.toUpperCase();
    }
}
```



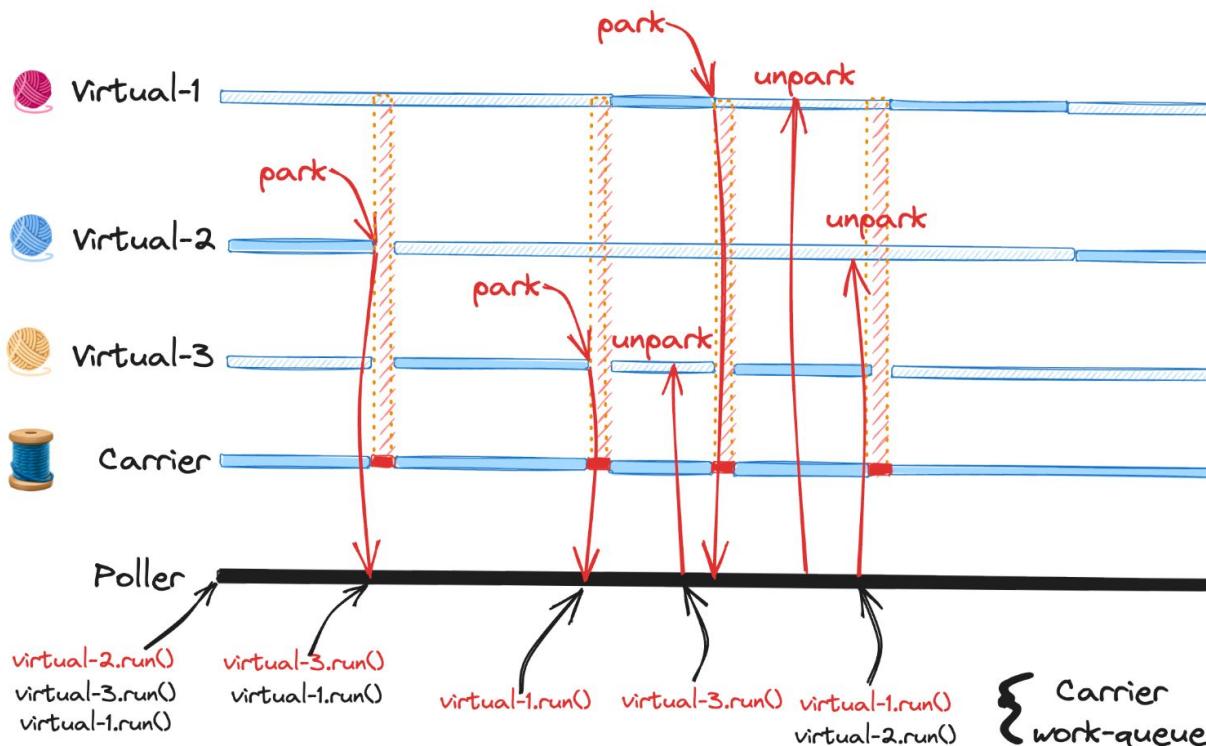
# Virtual Threads: 1 , n: Scheduling made in JVM



# 1 , n ≈ Reactive?



# Suspending threads, how to?



## Park

- store stack
- register unpark()
- unmount

## Unpark

- submit run() to carrier thread



*Bored with the theory,  
let's see some code!*



# Integration of Virtual Threads in Quarkus

## Imperative / Blocking

```
@GET  
public List<Todo> getAll() {  
    return  
Todo.listAll(Sort.by("order"));  
}
```

## Virtual Threads

```
@GET  
@RunOnVirtualThread  
public List<Todo> getAll() {  
    return  
Todo.listAll(Sort.by("order"))  
};  
}
```

## Reactive / Non-Blocking

```
@GET  
public Uni<List<Todo>>  
getAll() {  
    return  
Panache.withTransaction(()  
->  
Todo.findAll(Sort.by("order"  
)).list());  
}
```



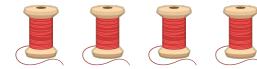
<https://github.com/danieloh30/quarkus-todo-app-main>

# Under the hood

At **build time**, Quarkus analyzes the endpoints:

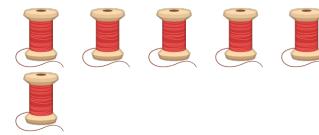
- Return types (reactive vs. imperative types)
- Method annotations
  - **@RunOnVirtualThread**
  - @Blocking
  - @NonBlocking

Reactive types or @NonBlocking



*fixed* number of event-loops

@Blocking and  
@RunOnVirtualThread

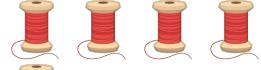
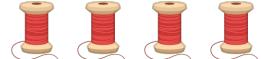


*dynamic* number of  
virtual threads



+1 per request

@Blocking or  
imperative types



# DEMO



<https://github.com/danieloh30/quarkus-todo-app-main>

*One annotation and  
everything is great?*



*Why not run everything on  
virtual threads?*



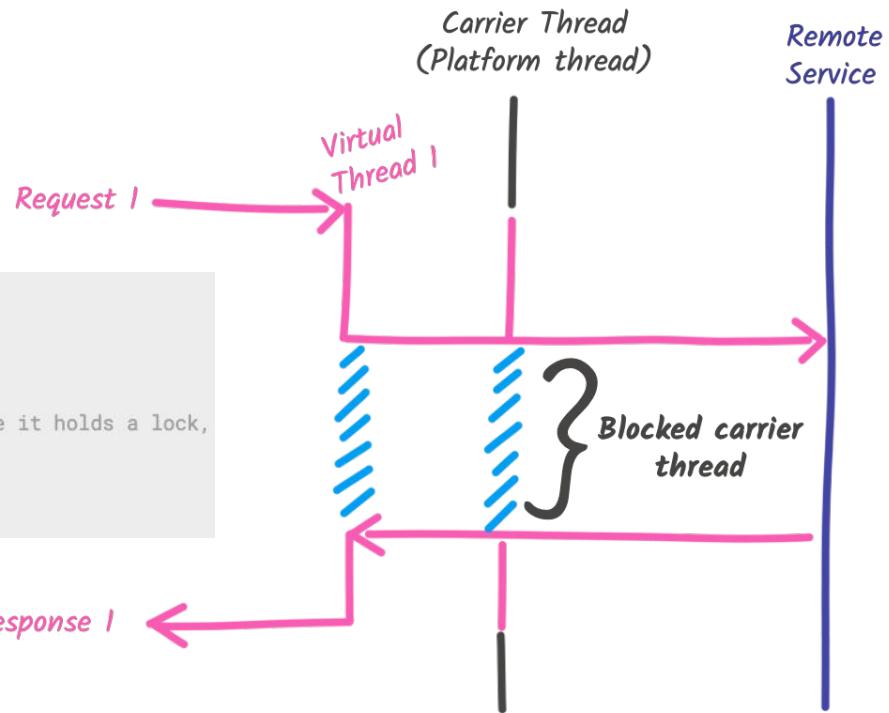
# 5 things you need to know on virtual threads



## Pinning

- Synchronized block or method
- Native method or Foreign function

```
Object monitor = new Object();
//...
public void aMethodThatPinTheCarrierThread() throws Exception {
    synchronized(monitor) {
        Thread.sleep(1000); // The virtual thread cannot be unmounted because it holds a lock,
                           // so the carrier thread is blocked.
    }
}
```

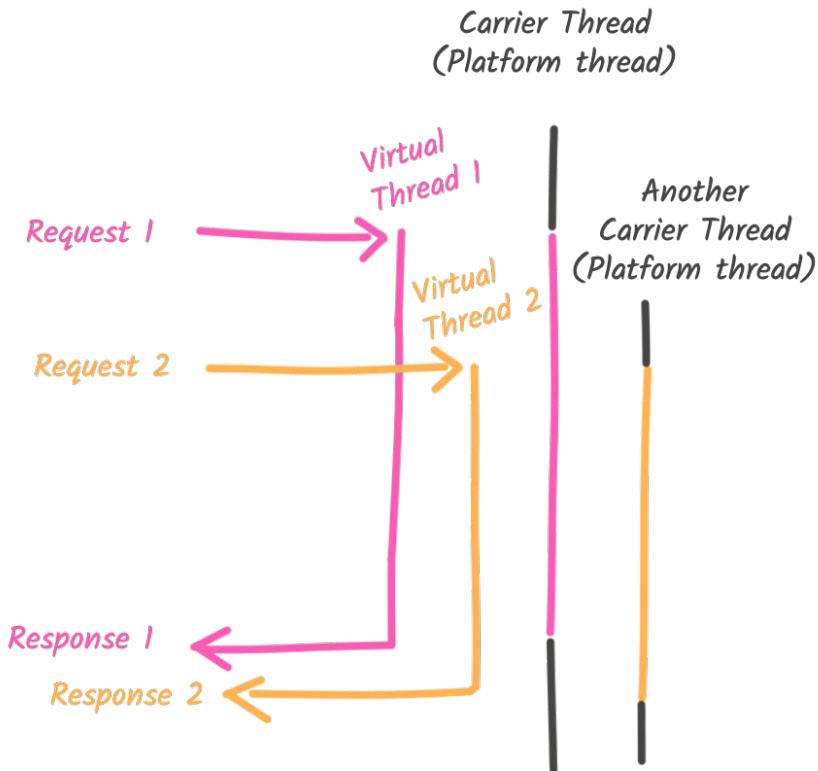


# 5 things you need to know on virtual threads



## Monopolization

- Lead to high-memory usage
- Avoid with using a dedicated platform thread pool



# 5 things you need to know on virtual threads



## Stressing thread safety

- Ensure to develop thread-safe code
- Lots of libraries and frameworks do not allow concurrent access

```
@GET  
@RunOnVirtualThread  
public String structuredConcurrencyExample() throws InterruptedException, ExecutionException {  
    var someState = ... // Must be thread-safe, as multiple virtual thread will access  
                      // it concurrently  
    try (var scope = new StructuredTaskScope.ShutdownOnFailure()) {  
        var task1 = scope.fork(() -> {  
            // Run in another virtual thread  
            return someState.touch();  
        });  
        var task2 = scope.fork(() -> {  
            // Run in another virtual thread  
            return someState.touch();  
        });  
  
        scope.join().throwIfFailed();  
  
        return task1.get() + "/" + task2.get();  
    }  
}
```



# 5 things you need to know on virtual threads



## Object pooling

- Using ThreadLocal (e.g., Jackson or Netty)
- Limited number of threads (recycled thread pool, unrelated but sequential tasks)



<https://quarkus.io/blog/virtual-thread-1/>

# 5 things you need to know on virtual threads



## Carrier thread pool elasticity

- Avoid having too many unscheduled virtual threads
- If you don't, in a container with memory constraints, the application can be killed



<https://quarkus.io/blog/virtual-thread-1/>

**A lot of the Java ecosystem would need a rewrite...**



### **More Java ecosystems**

- JMS, Rest Client, Mailer, Event Bus, Scheduled Methods, and so on



*What about numbers...*



# The application

## Techempower

- Multiple scenarios
- Run on bare metal and Cloud
- Multi languages, Multi frameworks

## The *Todo\** scenario

1. Fetch all rows from a DB (quotes)
2. Add one quote in the returned list
3. Sort the list
4. *Return the list as JSON\**



<https://github.com/danieloh30/quarkus-todo-app-main>

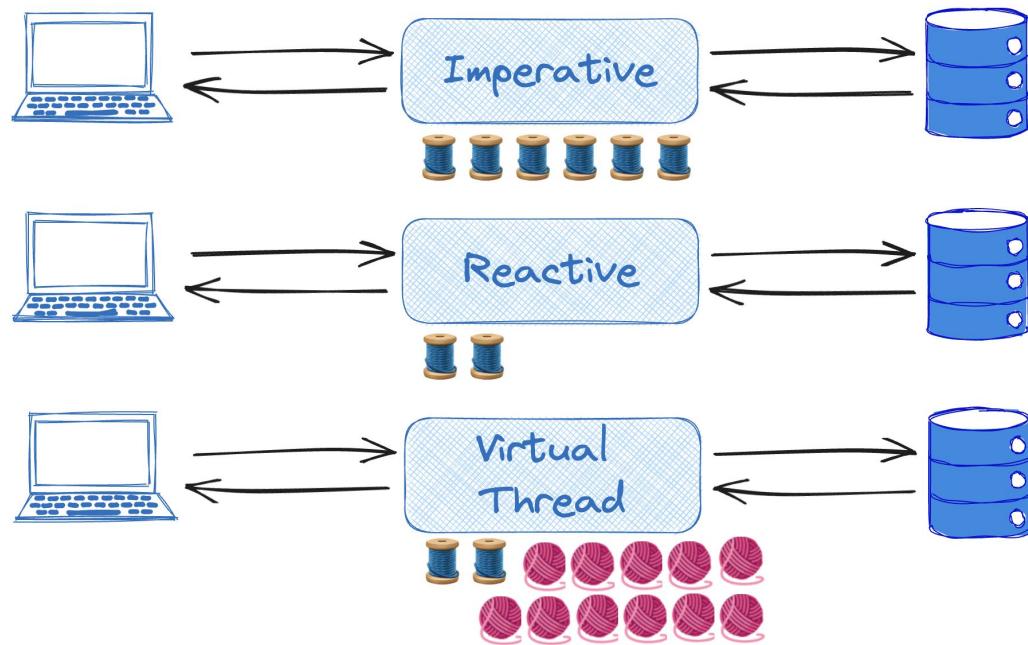
# The benchmark

## Variants

- Blocking,
- Non-Blocking (Reactive)
- @RunOnVirtualThread (Loom)

## Environments

- Container (0.5vCPU, 512Mb Ram)



<https://github.com/danieloh30/quarkus-todo-app-main>

# The benchmark

## Concurrency

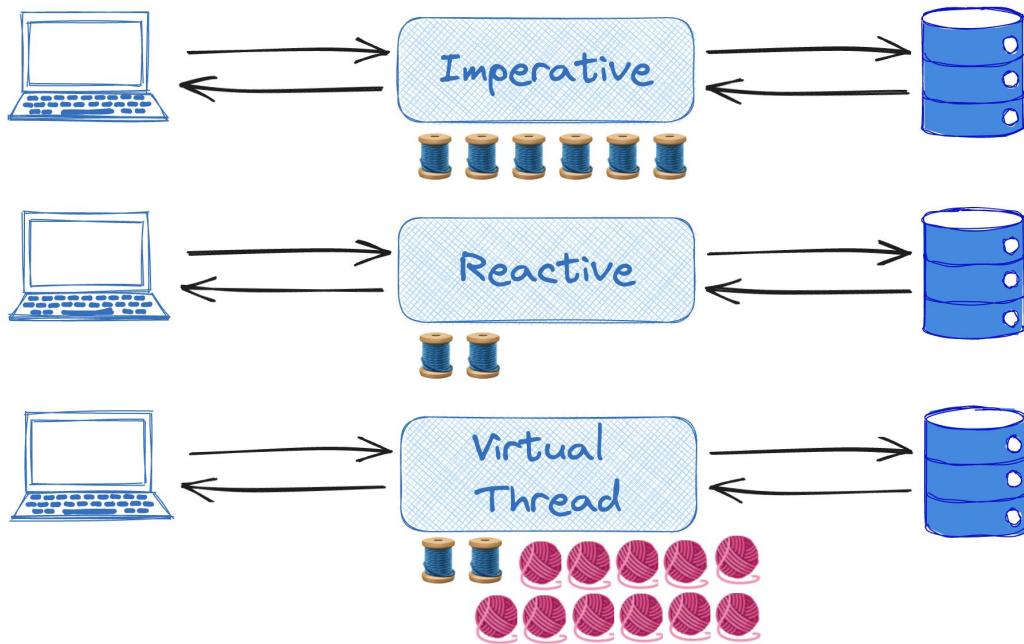
- 1200 => 4400 req/sec

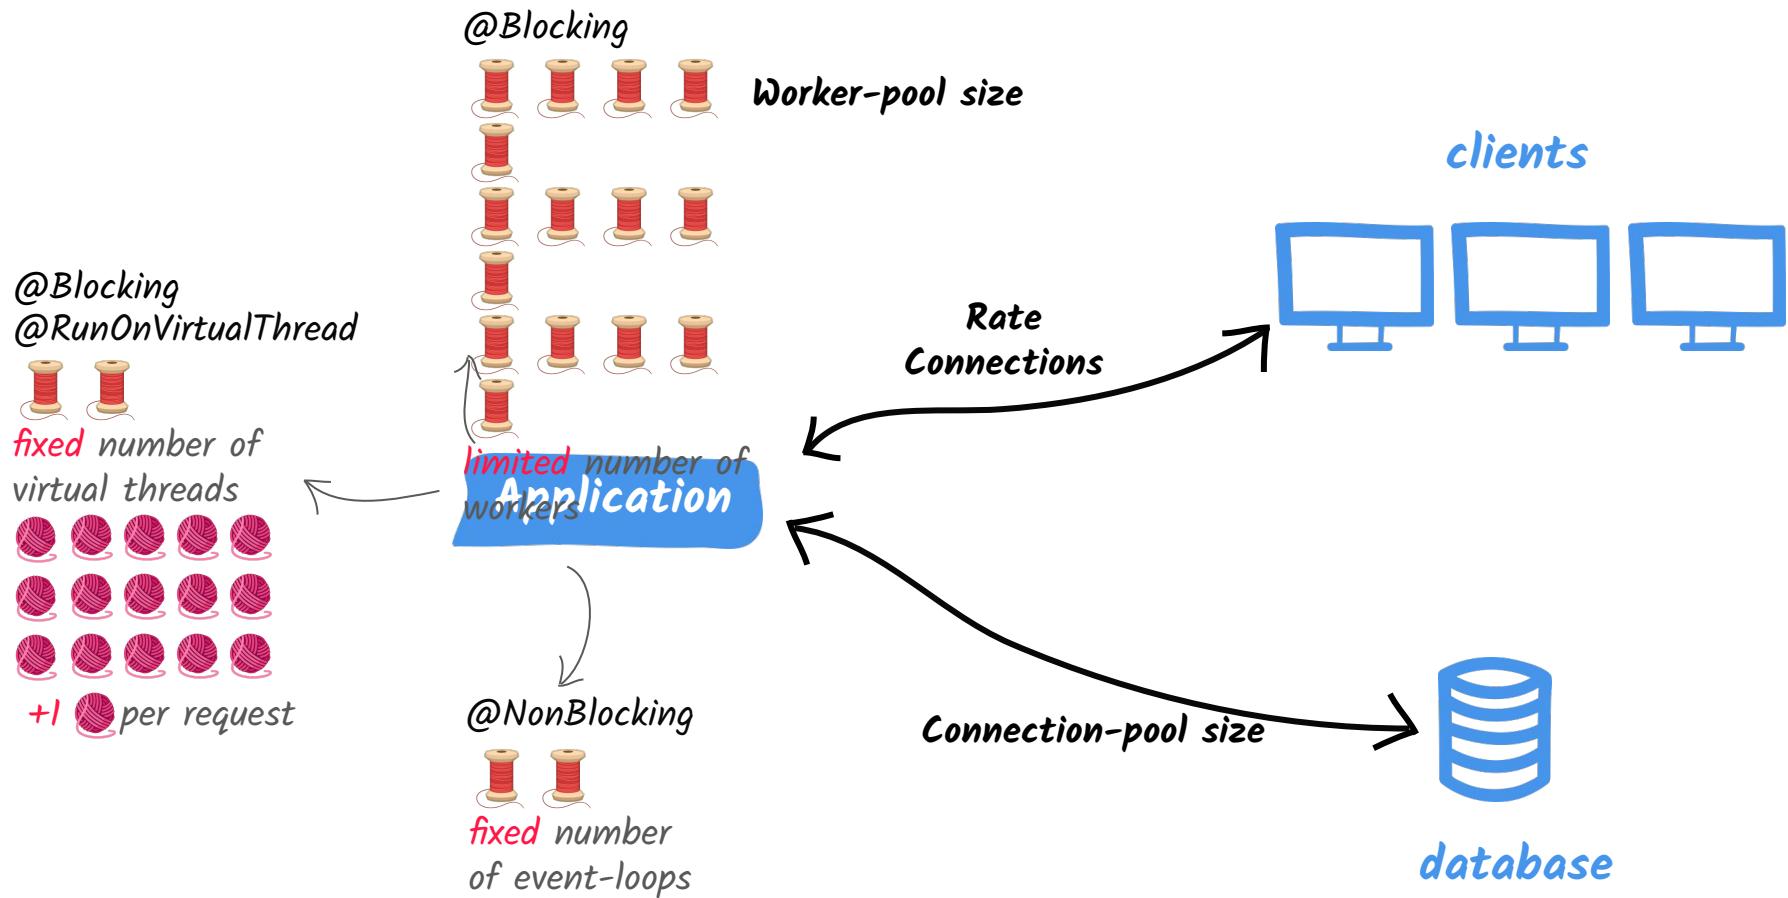
## Latency

- 0.2s of latency between application and DB

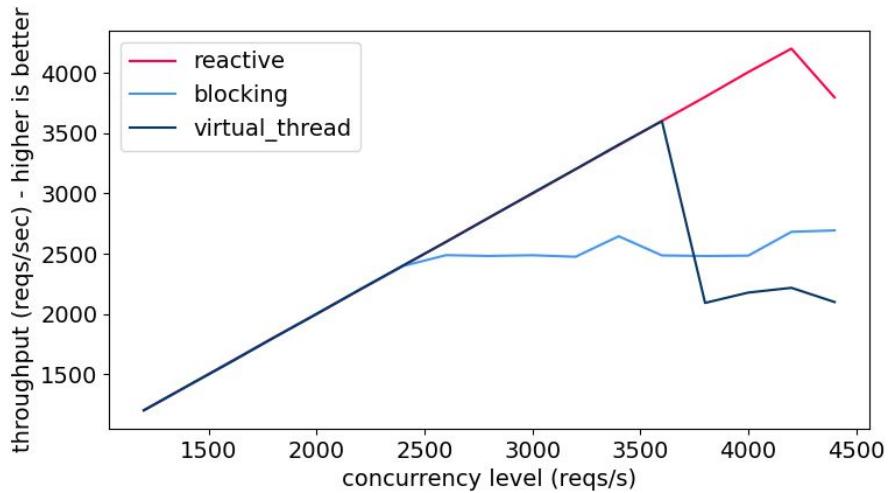
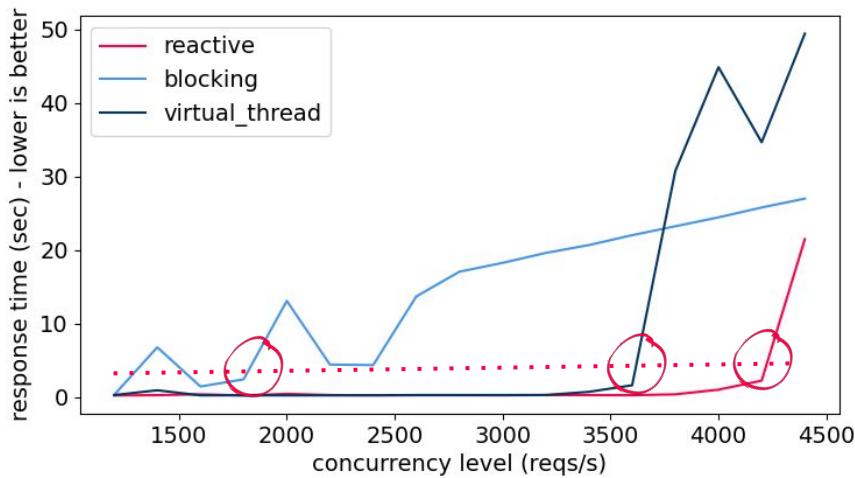
## Measures

- Response time (max)
- Throughput
- Resident State Size (RSS)





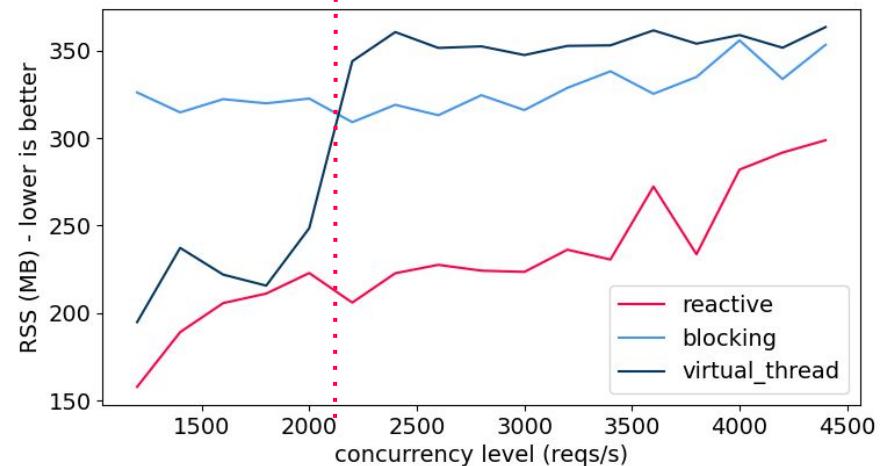
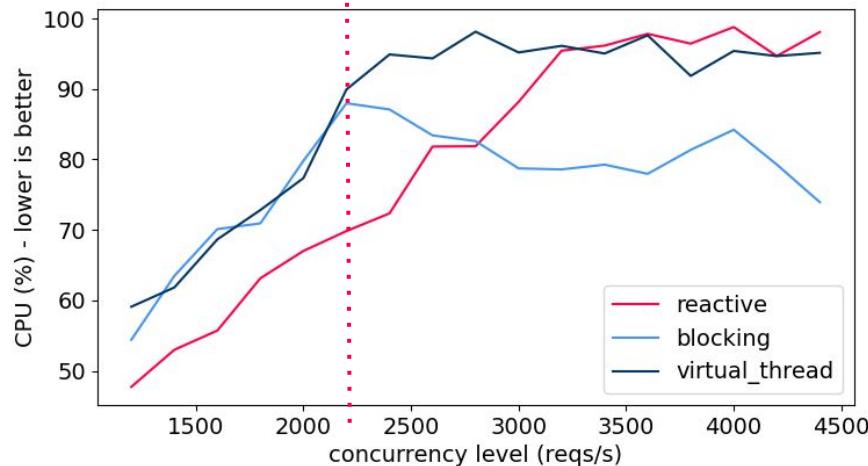
# Response time and Throughput



Scale **better** than worker threads,  
but not as good as reactive.



# Resource Usage (CPU and RSS)

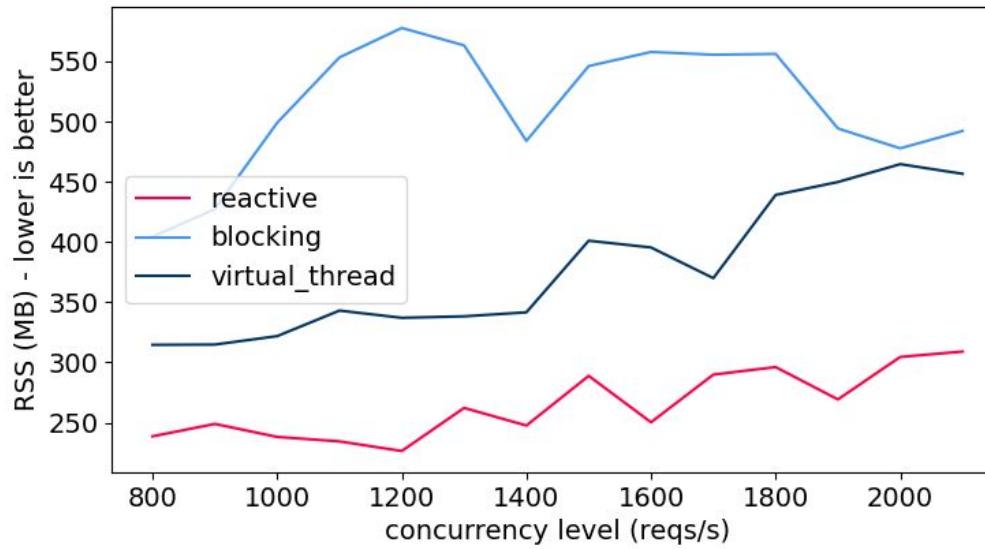


**High memory/CPU usage  
(higher than worker threads)**



<https://github.com/danieloh30/quarkus-todo-app-main>

# Memory Usage - Container



**High memory usage  
(higher than reactive)**



<https://github.com/danieloh30/quarkus-todo-app-main>

# So, what about Virtual Threads?

## Imperative

- Simple development model
- Expensive: 1 Mb / thread + 10 µs
- Limit concurrency

## Virtual Threads

- Simple development model 🌟
- Resource efficient 🤝
- High concurrency 🙌

## Reactive

- More complex development model
- Resource efficient
- High concurrency



# *Key Takeaways*



# Virtual Threads in Quarkus

## Every reactive APIs can be used

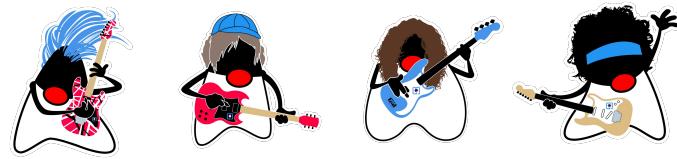
- Databases
- Network client (Rest client, gRPC, GraphQL)
- Kafka, Pulsar, AMQP
- Template engine, Email...

## Roadmap

- Web Sockets
- CDI Async Observer

## How do use Virtual Threads in Quarkus

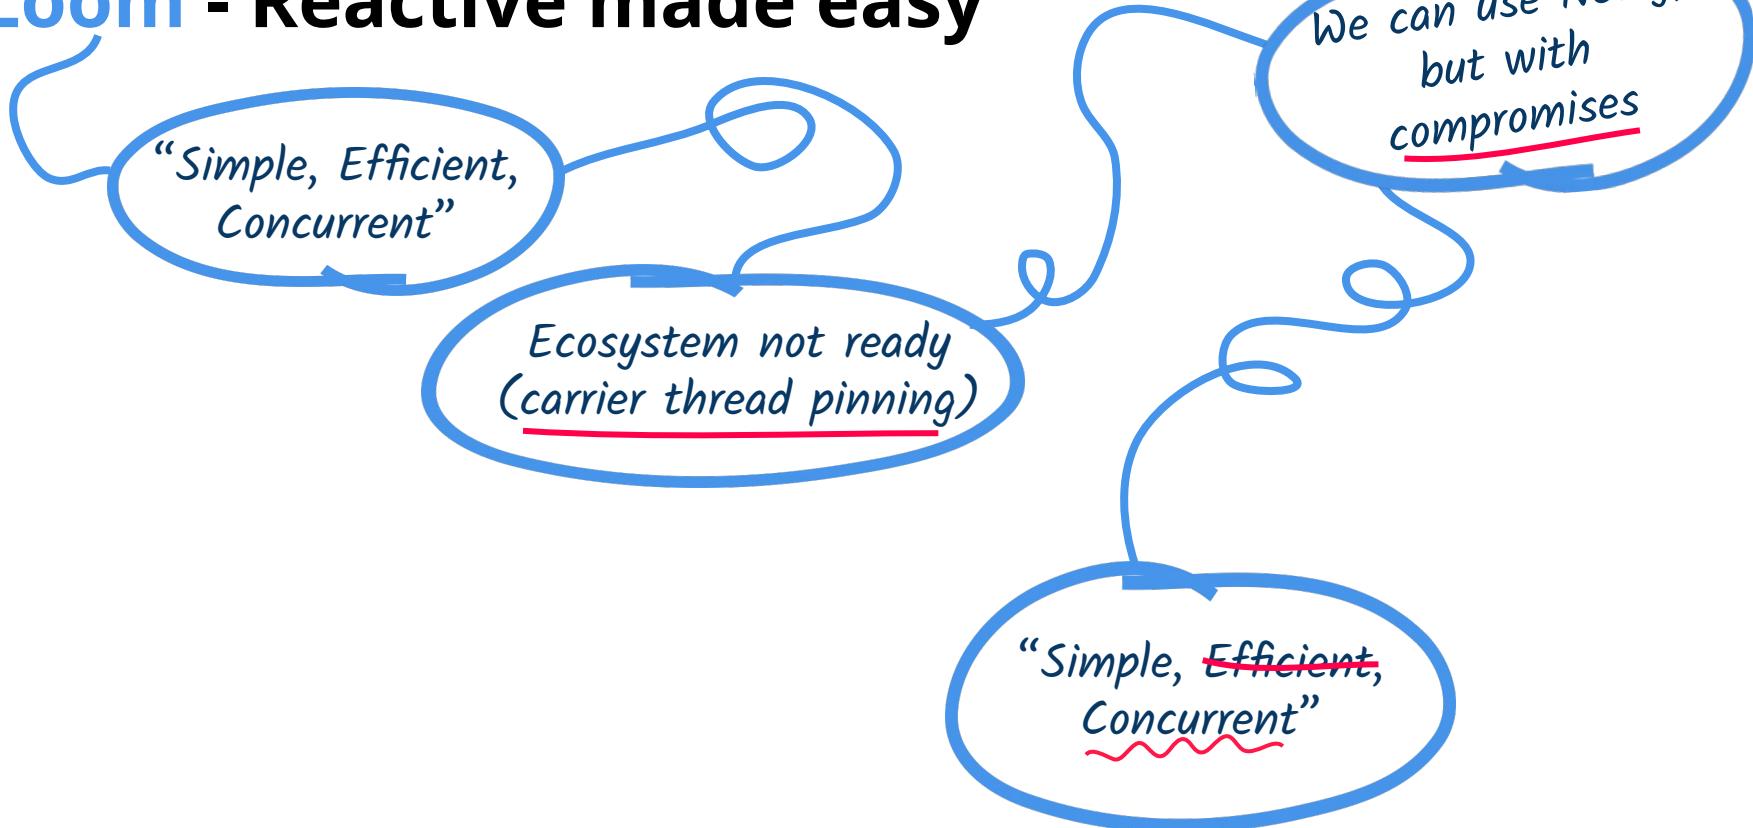
- JDK 21
- Use **@RunOnVirtualThread**



# Quarkus 3: The Road to Loom for Cheaper, Faster, and Easier Concurrent Applications?



# Loom - Reactive made easy



*Learn More*

## Quarkus Virtual Threads

- REST / HTTP
- Hibernate, Bean Validation, Transactions
- Kafka, AMQP, JMS
- gRPC
- Scheduled tasks
- Event Bus
- ...



# [youtube.com/@danieloh30](https://youtube.com/@danieloh30)

# [bit.ly/danielohtv](https://bit.ly/danielohtv)



The screenshot shows the Daniel Oh YouTube channel interface. It features three main playlists:

- KUBERNETES LEARN BY EXAMPLE**: A series of 11 videos titled "Kubernetes Learn by Example #11" through "#1".
- QUARKUS**: A series of 8 videos related to Quarkus, including "Monitoring Native Executables with Quarkus and JFR", "Quarkus 3 natively Integrates Azure Functions", and "Flexible Data Development with DynamoDB".
- SERVERLESS & FUNCTION**: A series of 7 videos related to serverless and functions, including "Quarkus 3 natively Integrates Azure Functions" (repeated), "Distributed Tracing Integration with OpenTelemetry, Knative, and Quarkus", and "Build your first Java Serverless Function using Quarkus Quick start".

Each video thumbnail includes the title, duration, and view count. The channel has over 600,000 subscribers and 10 million views.





# Thank you! Questions?

