

What the CRaC...

Coordinated Restore at Checkpoint
on the Java Virtual Machine

ABOUTME.



Gerrit Grunwald | Developer Advocate | Azul

SLIDES

...



JAVAIS

GREAT



VIBRANT
COMMUNITY. . .

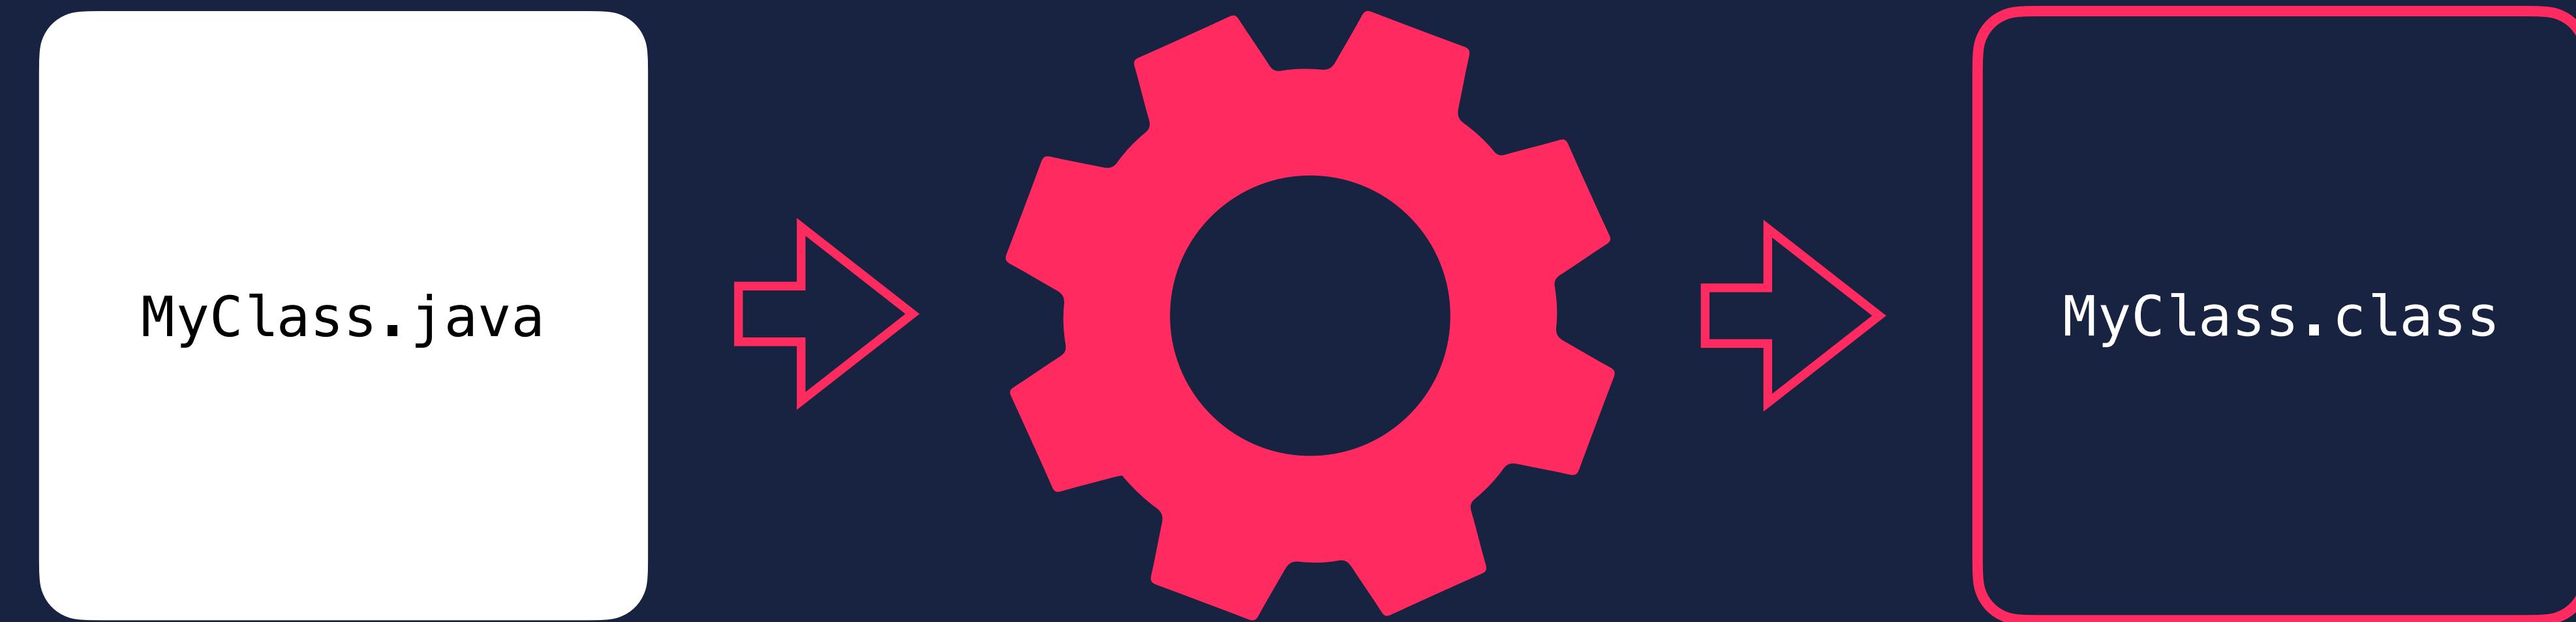
HUNDREDS OF
JUGS....

THOUSANDS OF
FOSS PROJECTS...
ooo



**JAVA VIRTUAL
MACHINE**

HOW DOES
IT WORK. . .



SOURCE CODE

COMPILER

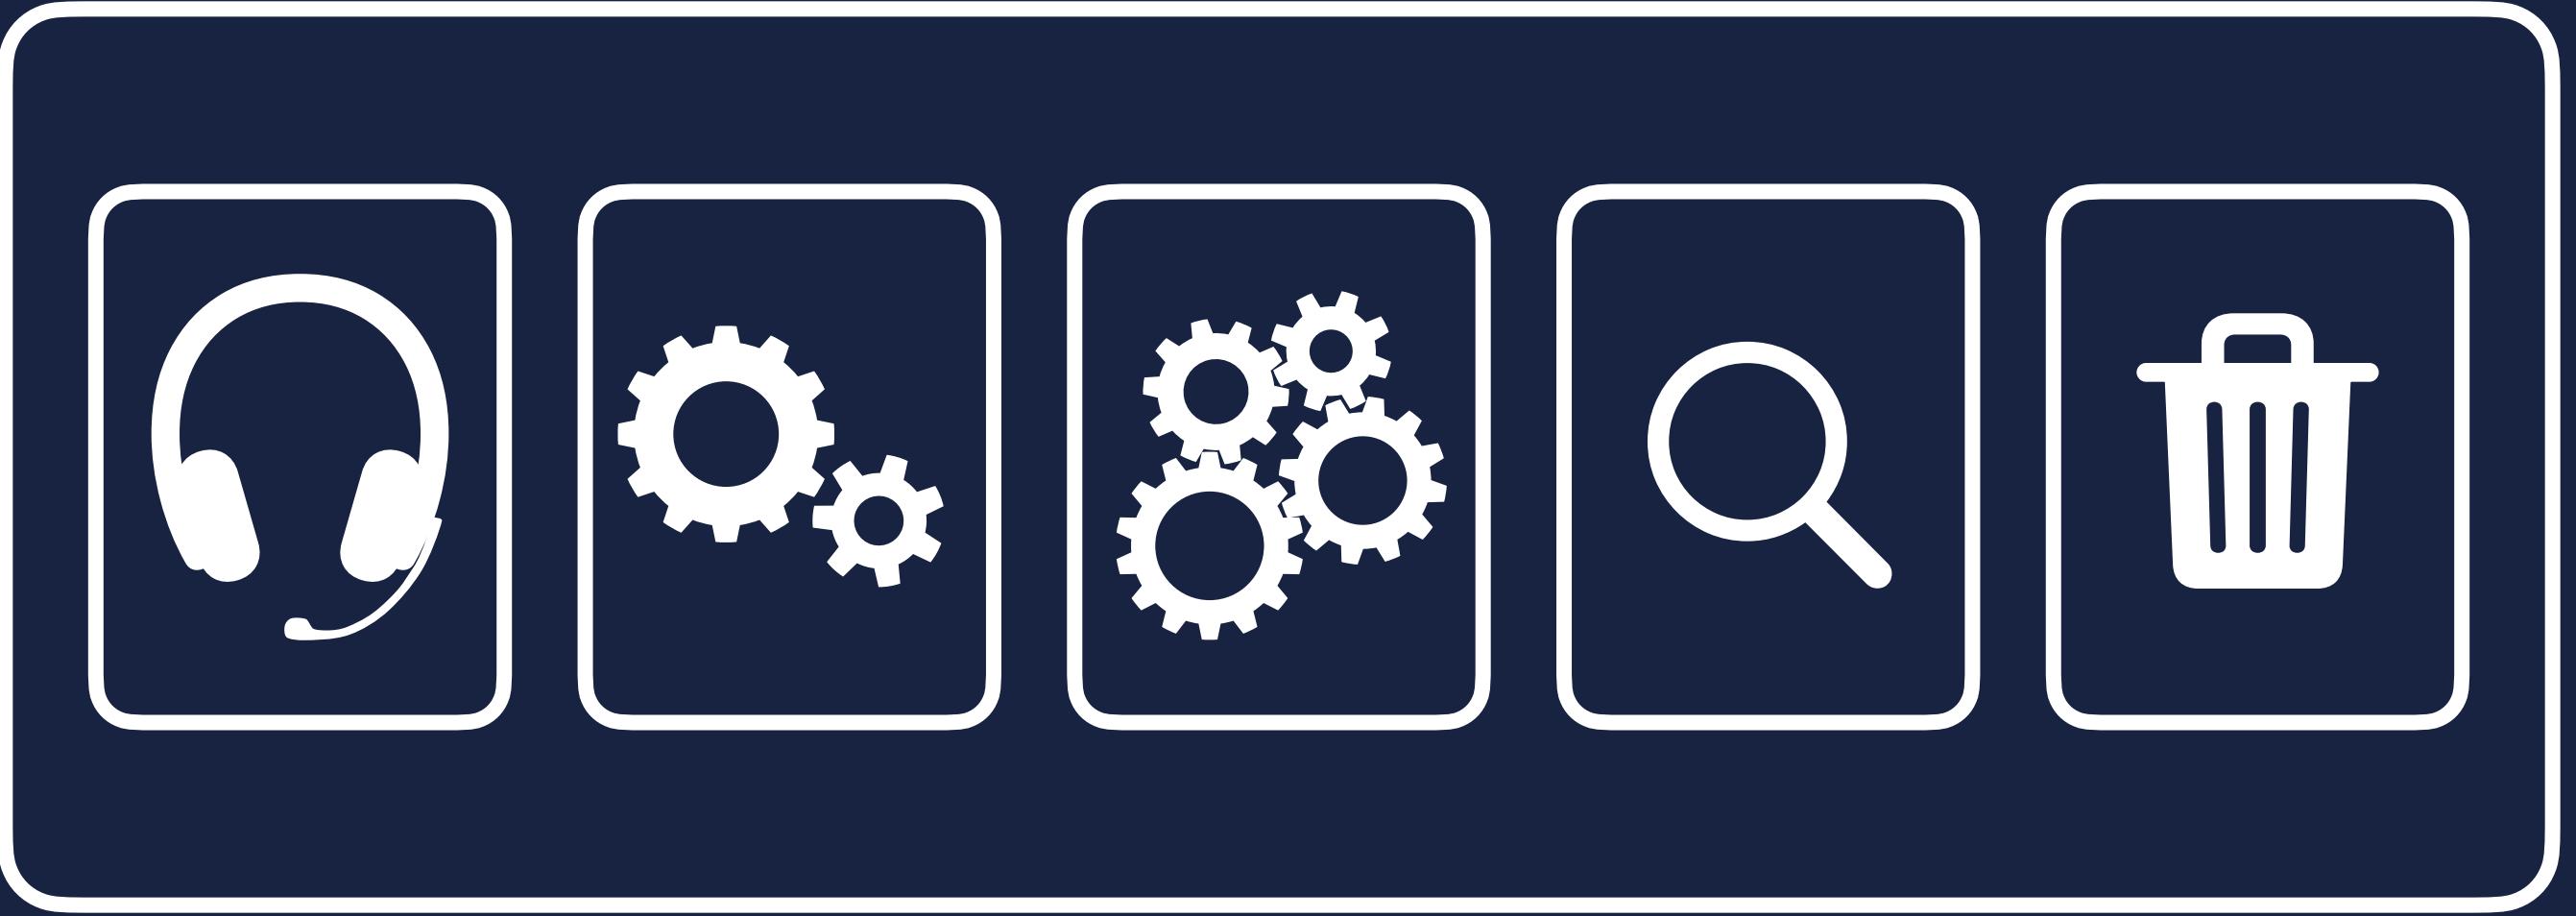
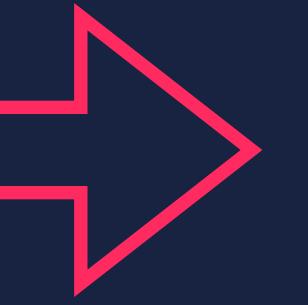
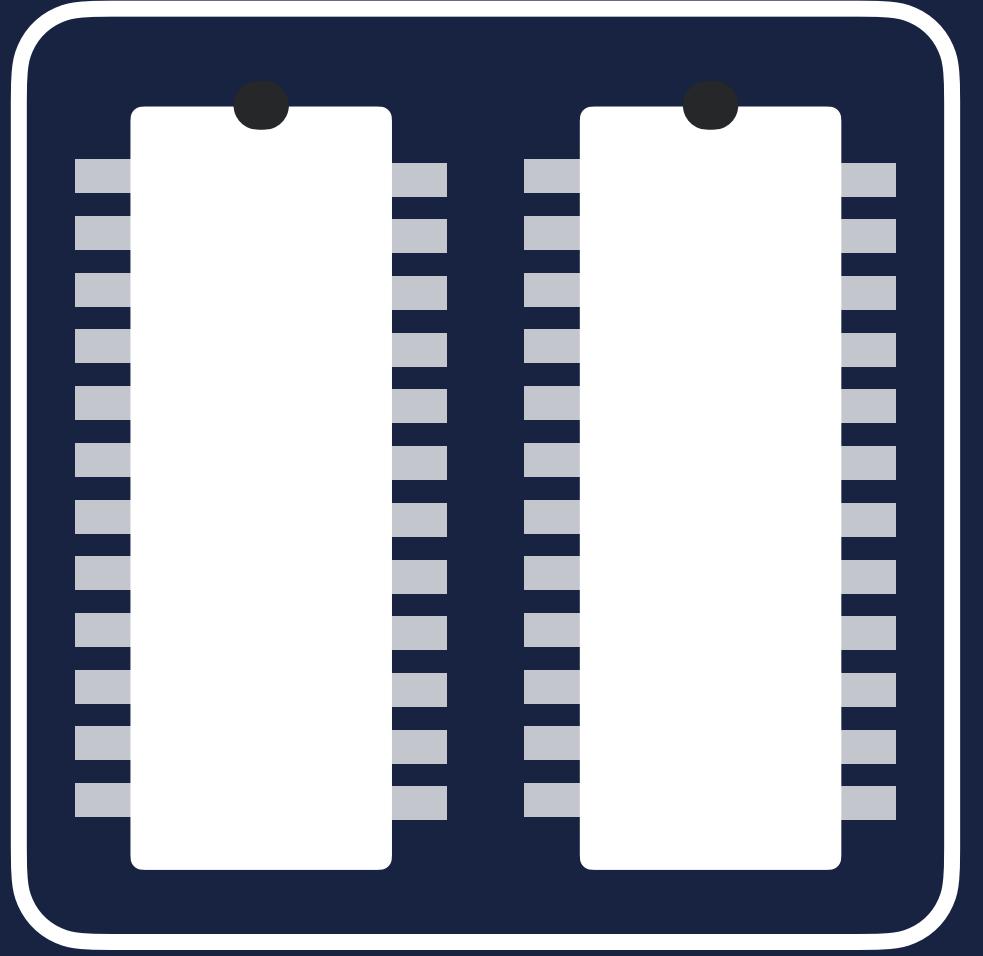
BYTE CODE



BYTE CODE

CLASS LOADER

JVM MEMORY



JVM MEMORY

EXECUTION ENGINE

EXECUTION ENGINE



Interpreter



C1 JIT
Compiler
(client)



C2 JIT
Compiler
(server)



Profiler



Garbage
Collector

EXECUTION ENGINE



Interpreter



C1 JIT
Compiler
(client)



C2 JIT
Compiler
(server)



Profiler



Garbage
Collector

Converts ByteCode into
instruction set of CPU



INTERPRETER

Detects hot spots by
counting method calls

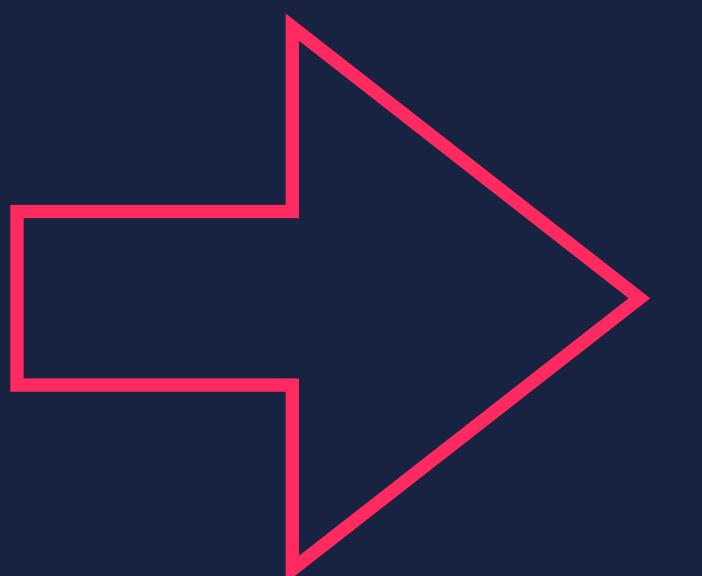


JVM

Pass the hot spot methods
to C1 JIT Compiler



JVM



Compiles code as quickly
as possible with low optimisation



C1 JIT
COMPILER

Compiles code as quickly
as possible with low optimisation



C1 JIT
COMPILER

THRESHOLD
REACHED



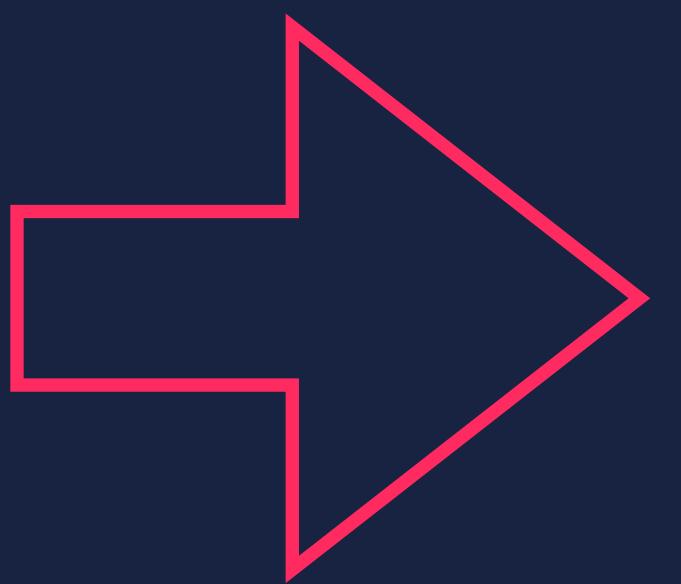
JVM

Profiles the running code
(detecting hot code)

Pass the "hot" code
to C2 JIT Compiler



JVM



Compiles code with best
optimisation possible (slower)



C2 JIT
COMPILER

TIERED

COMPILATION

LEVELS OF EXECUTION

- Level 0 - Interpreted code
- Level 1 - C1 compiled code (no profiling)
- Level 2 - C1 compiled code (basic profiling)
- Level 3 - C1 compiled code (full profiling, ~35% overhead)
- Level 4 - C2 compiled code (uses profile data from previous steps)

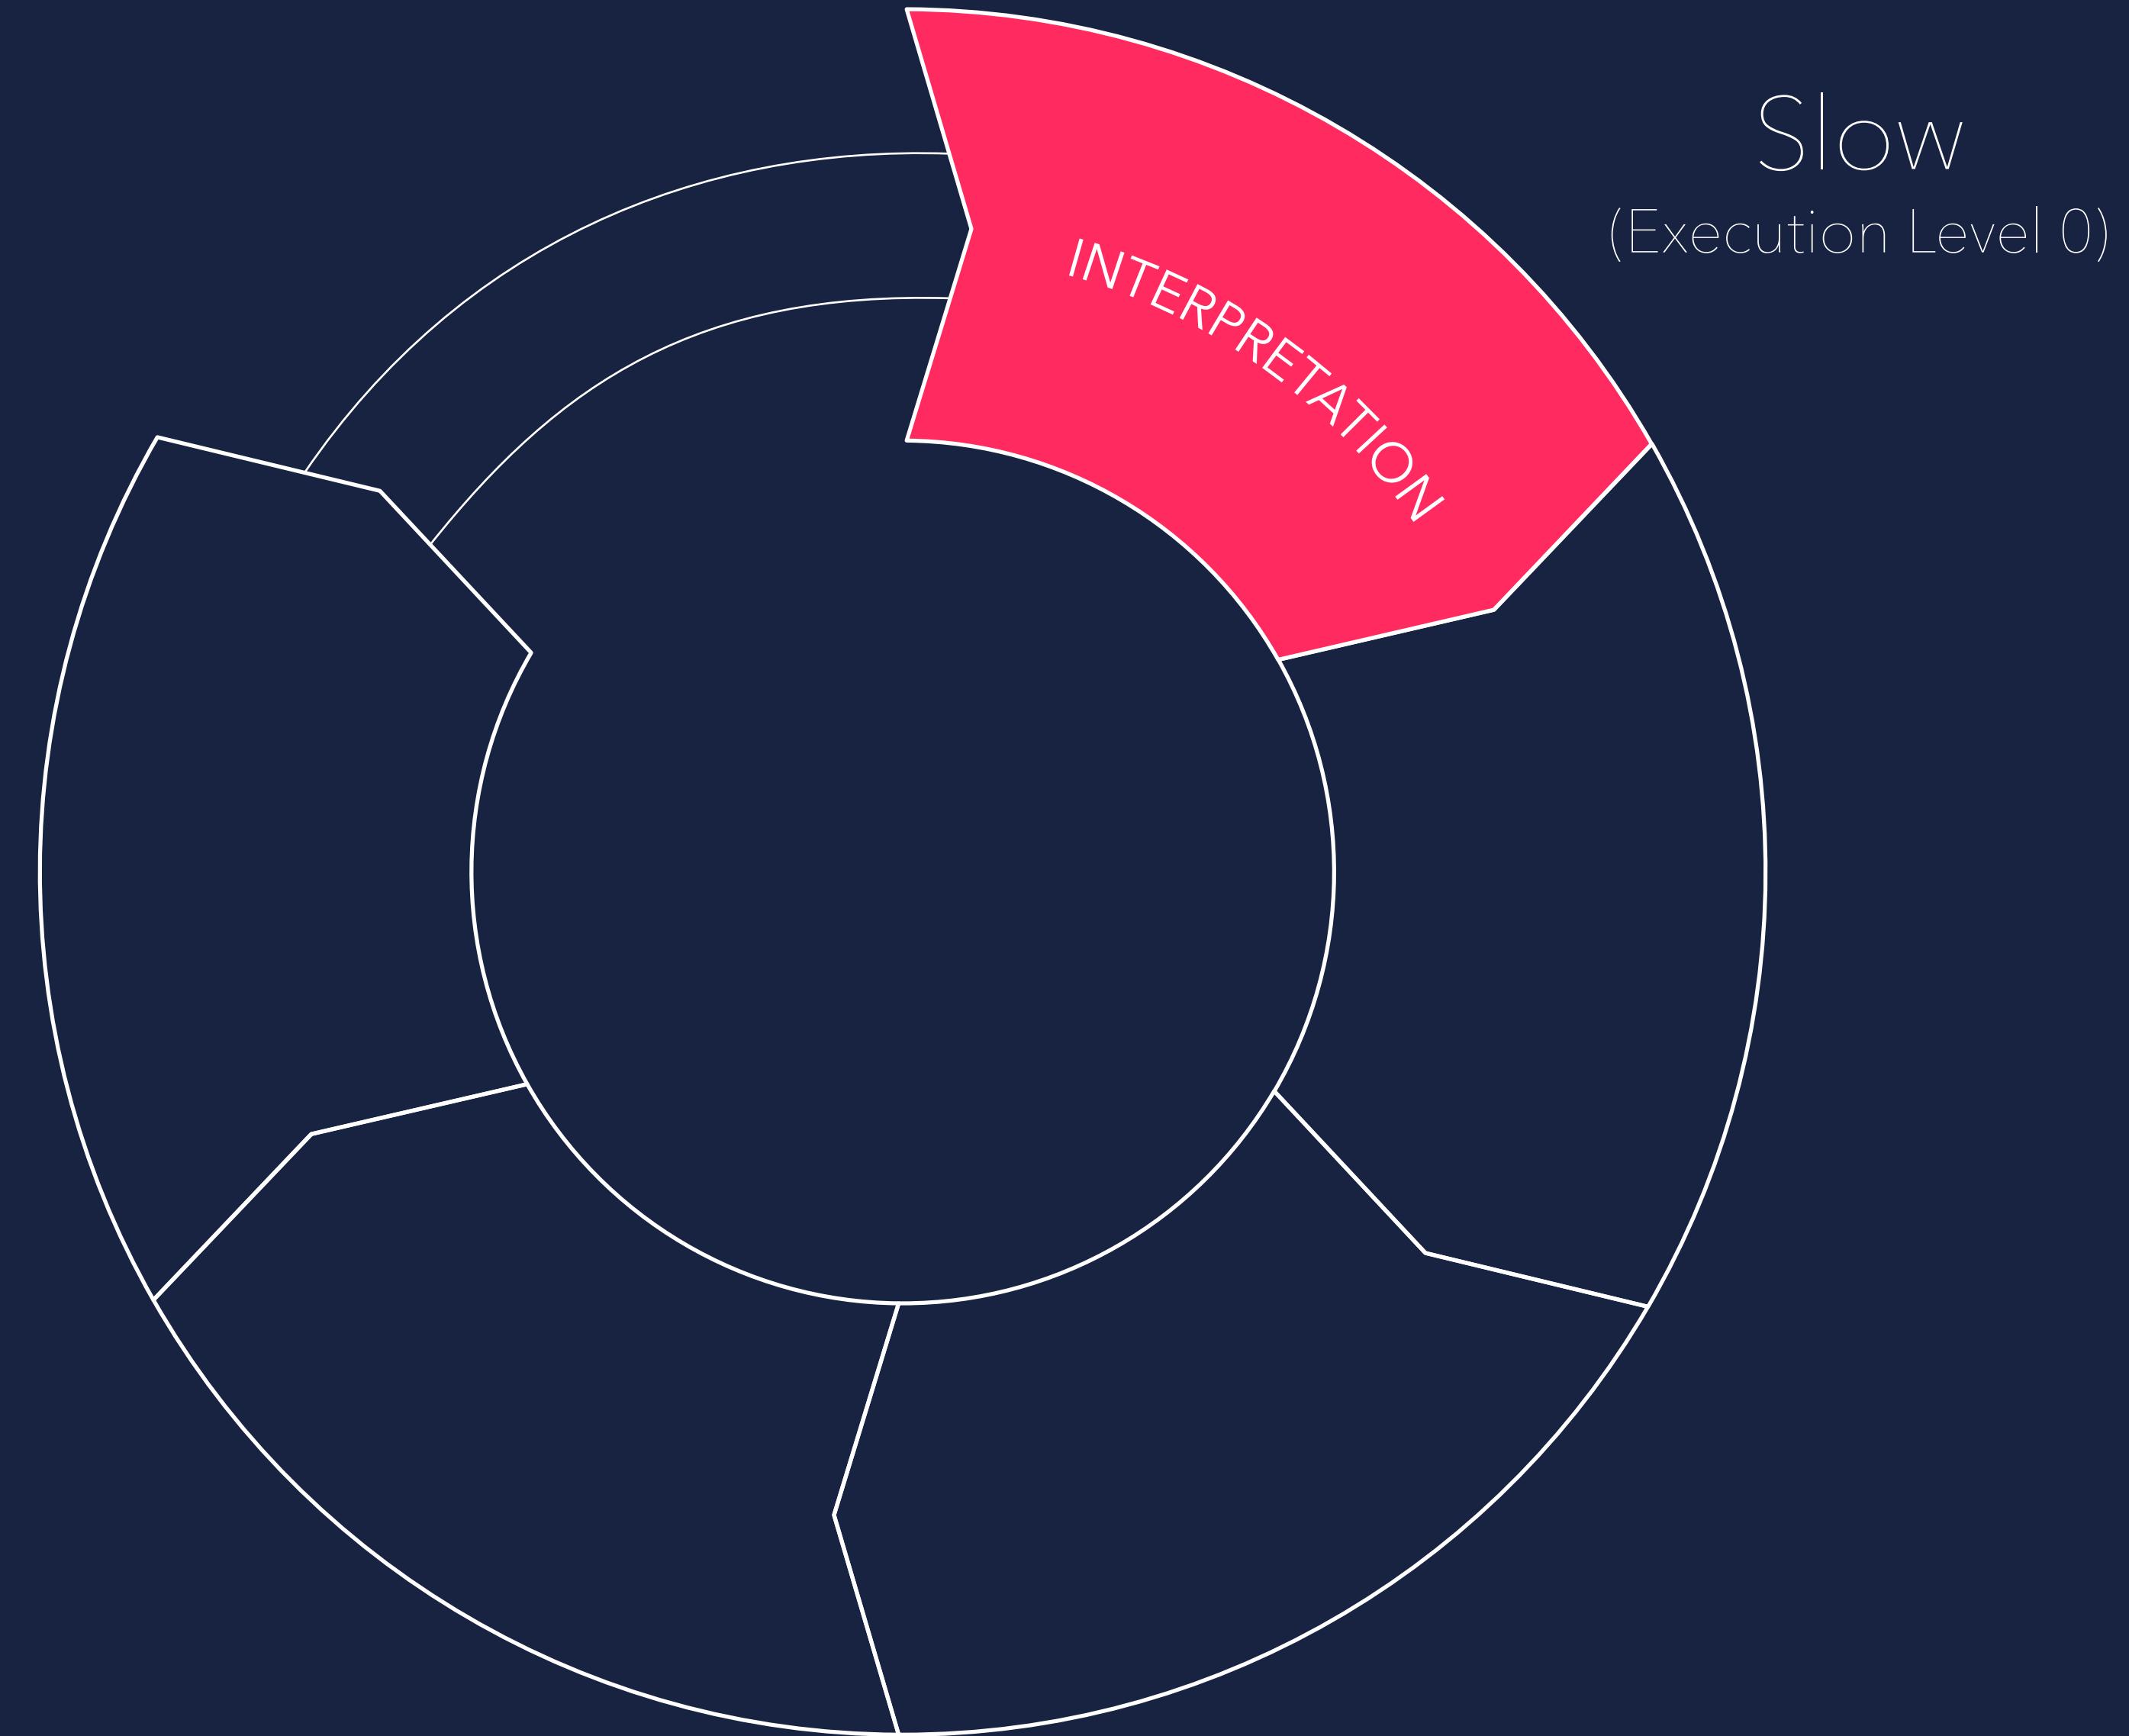
LEVELS OF EXECUTION

- Level 0 - Interpreted code
- Level 1 - C1 compiled code (no profiling)
- Level 2 - C1 compiled code (basic profiling)
- Level 3 - C1 compiled code (full profiling, ~35% overhead)
- Level 4 - C2 compiled code (uses profile data from previous steps)

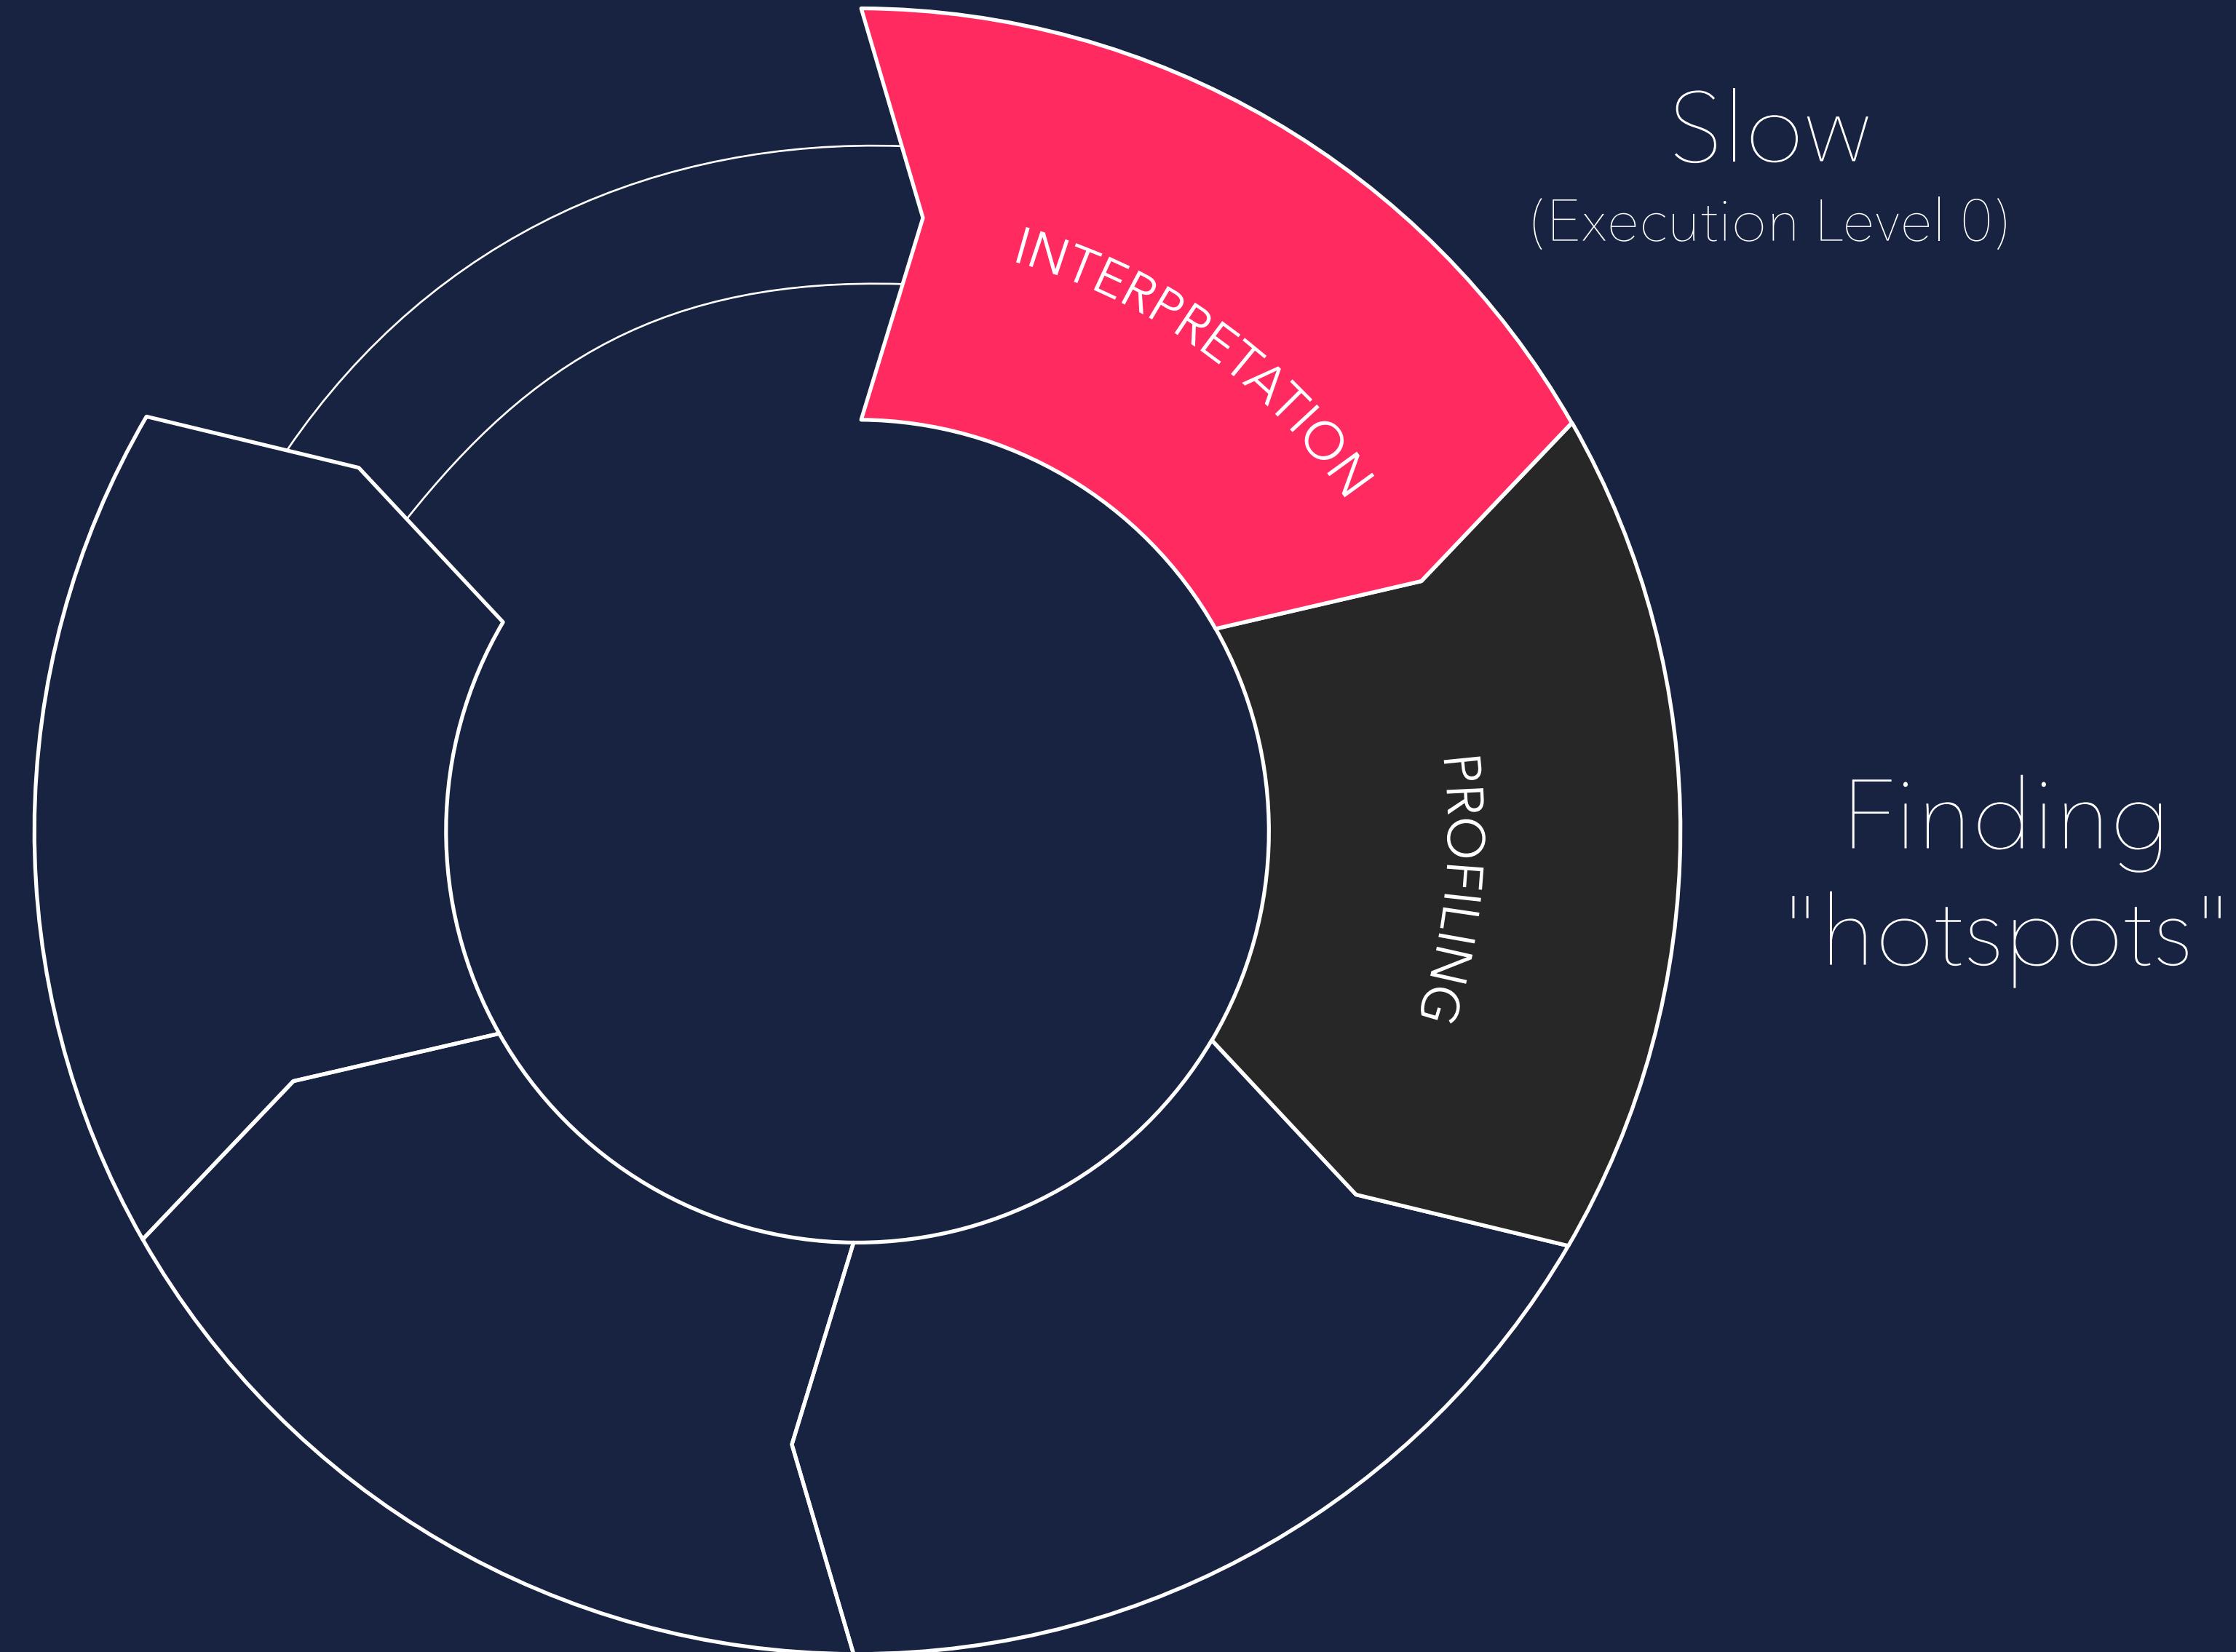
PREFERRED

EXECUTION
CYCLE

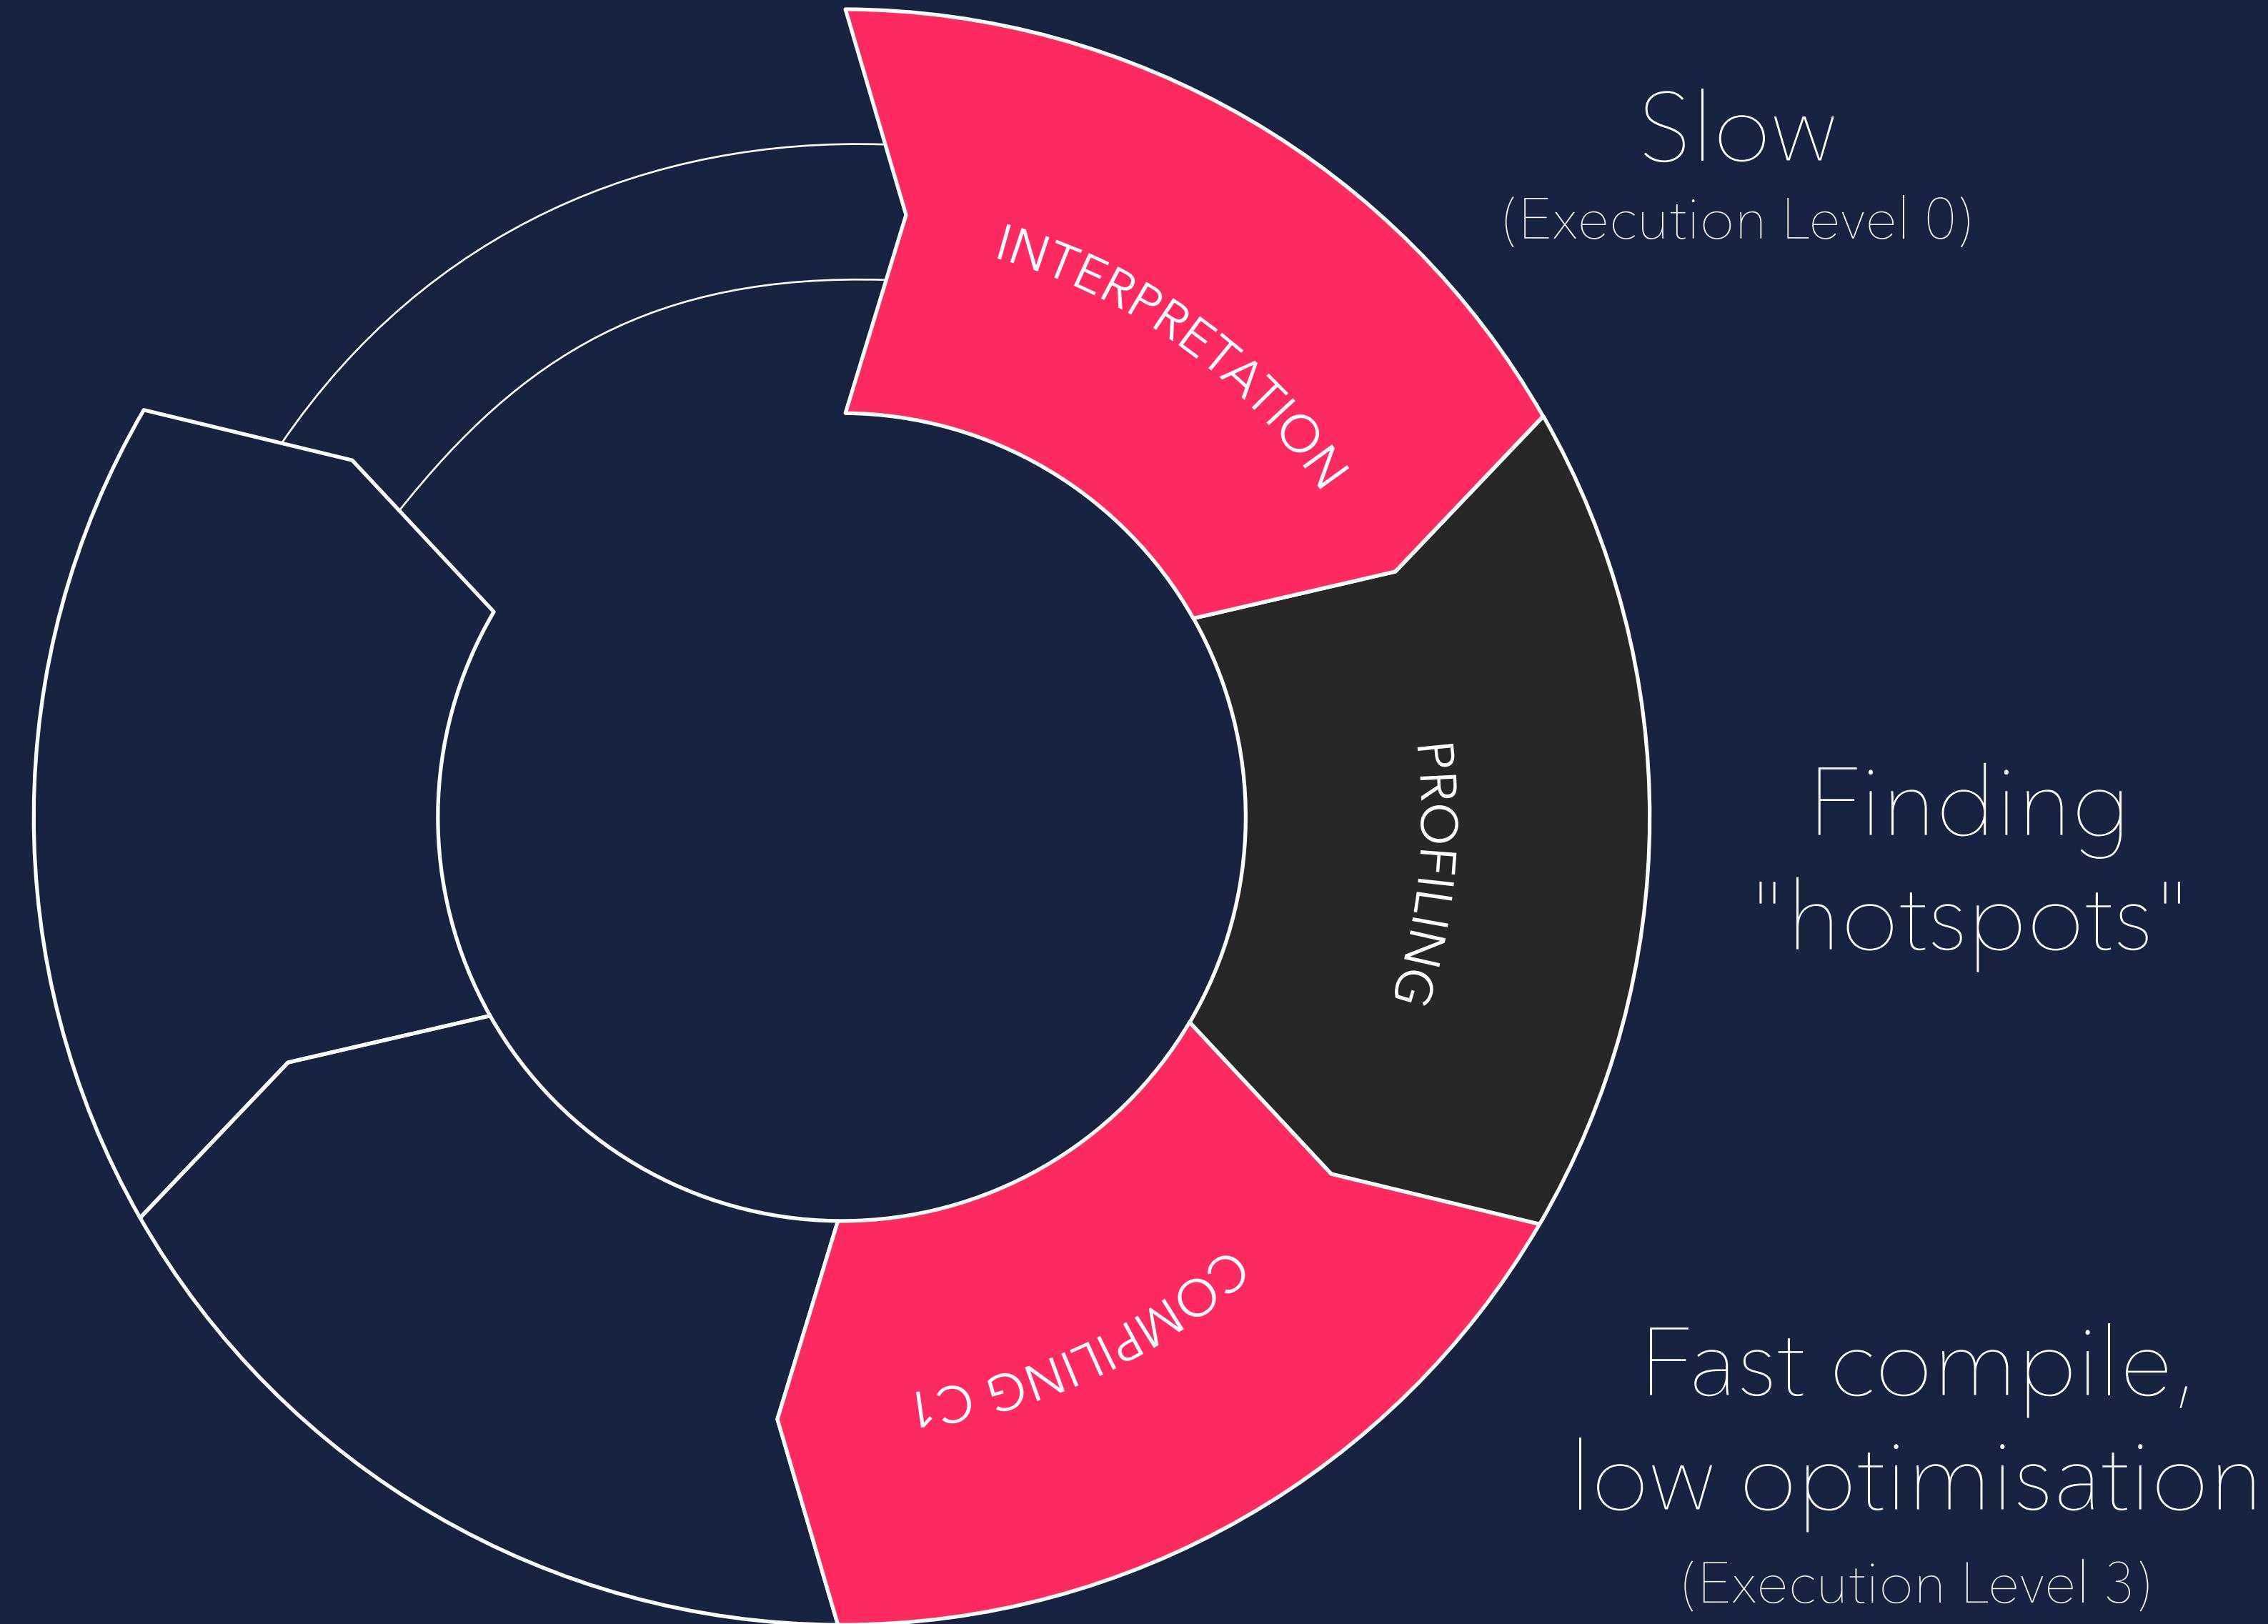
EXECUTION CYCLE



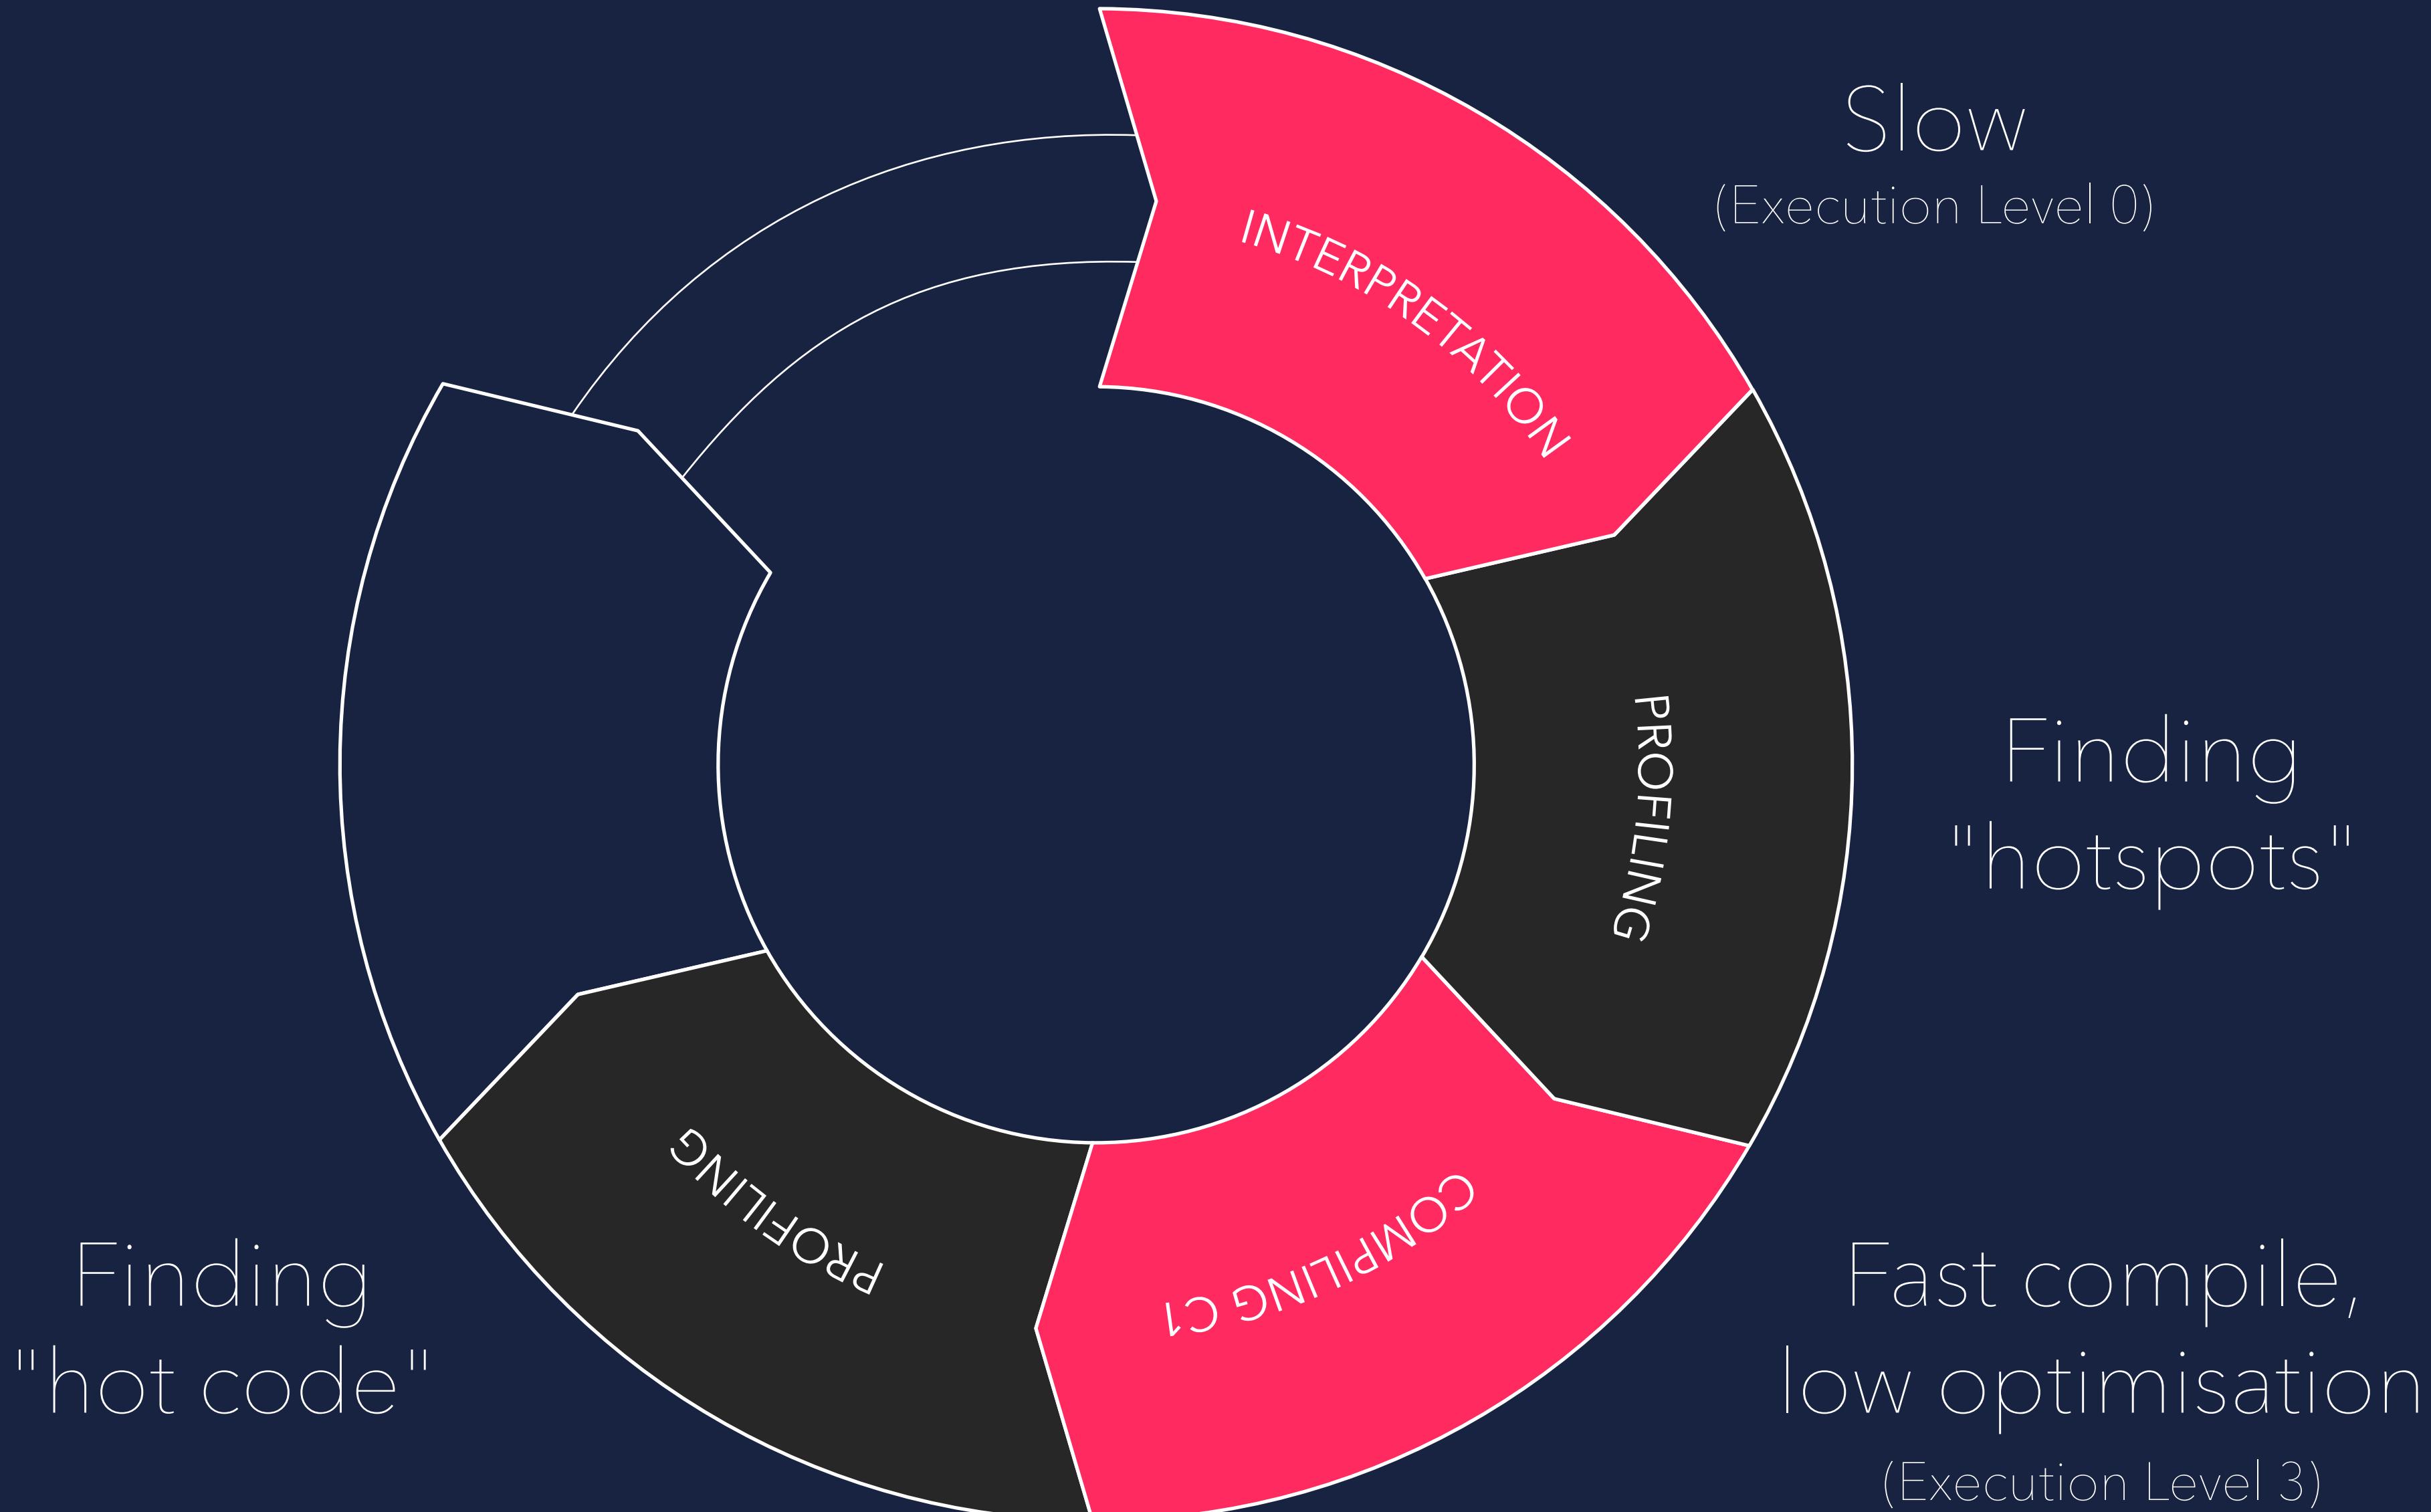
EXECUTION CYCLE



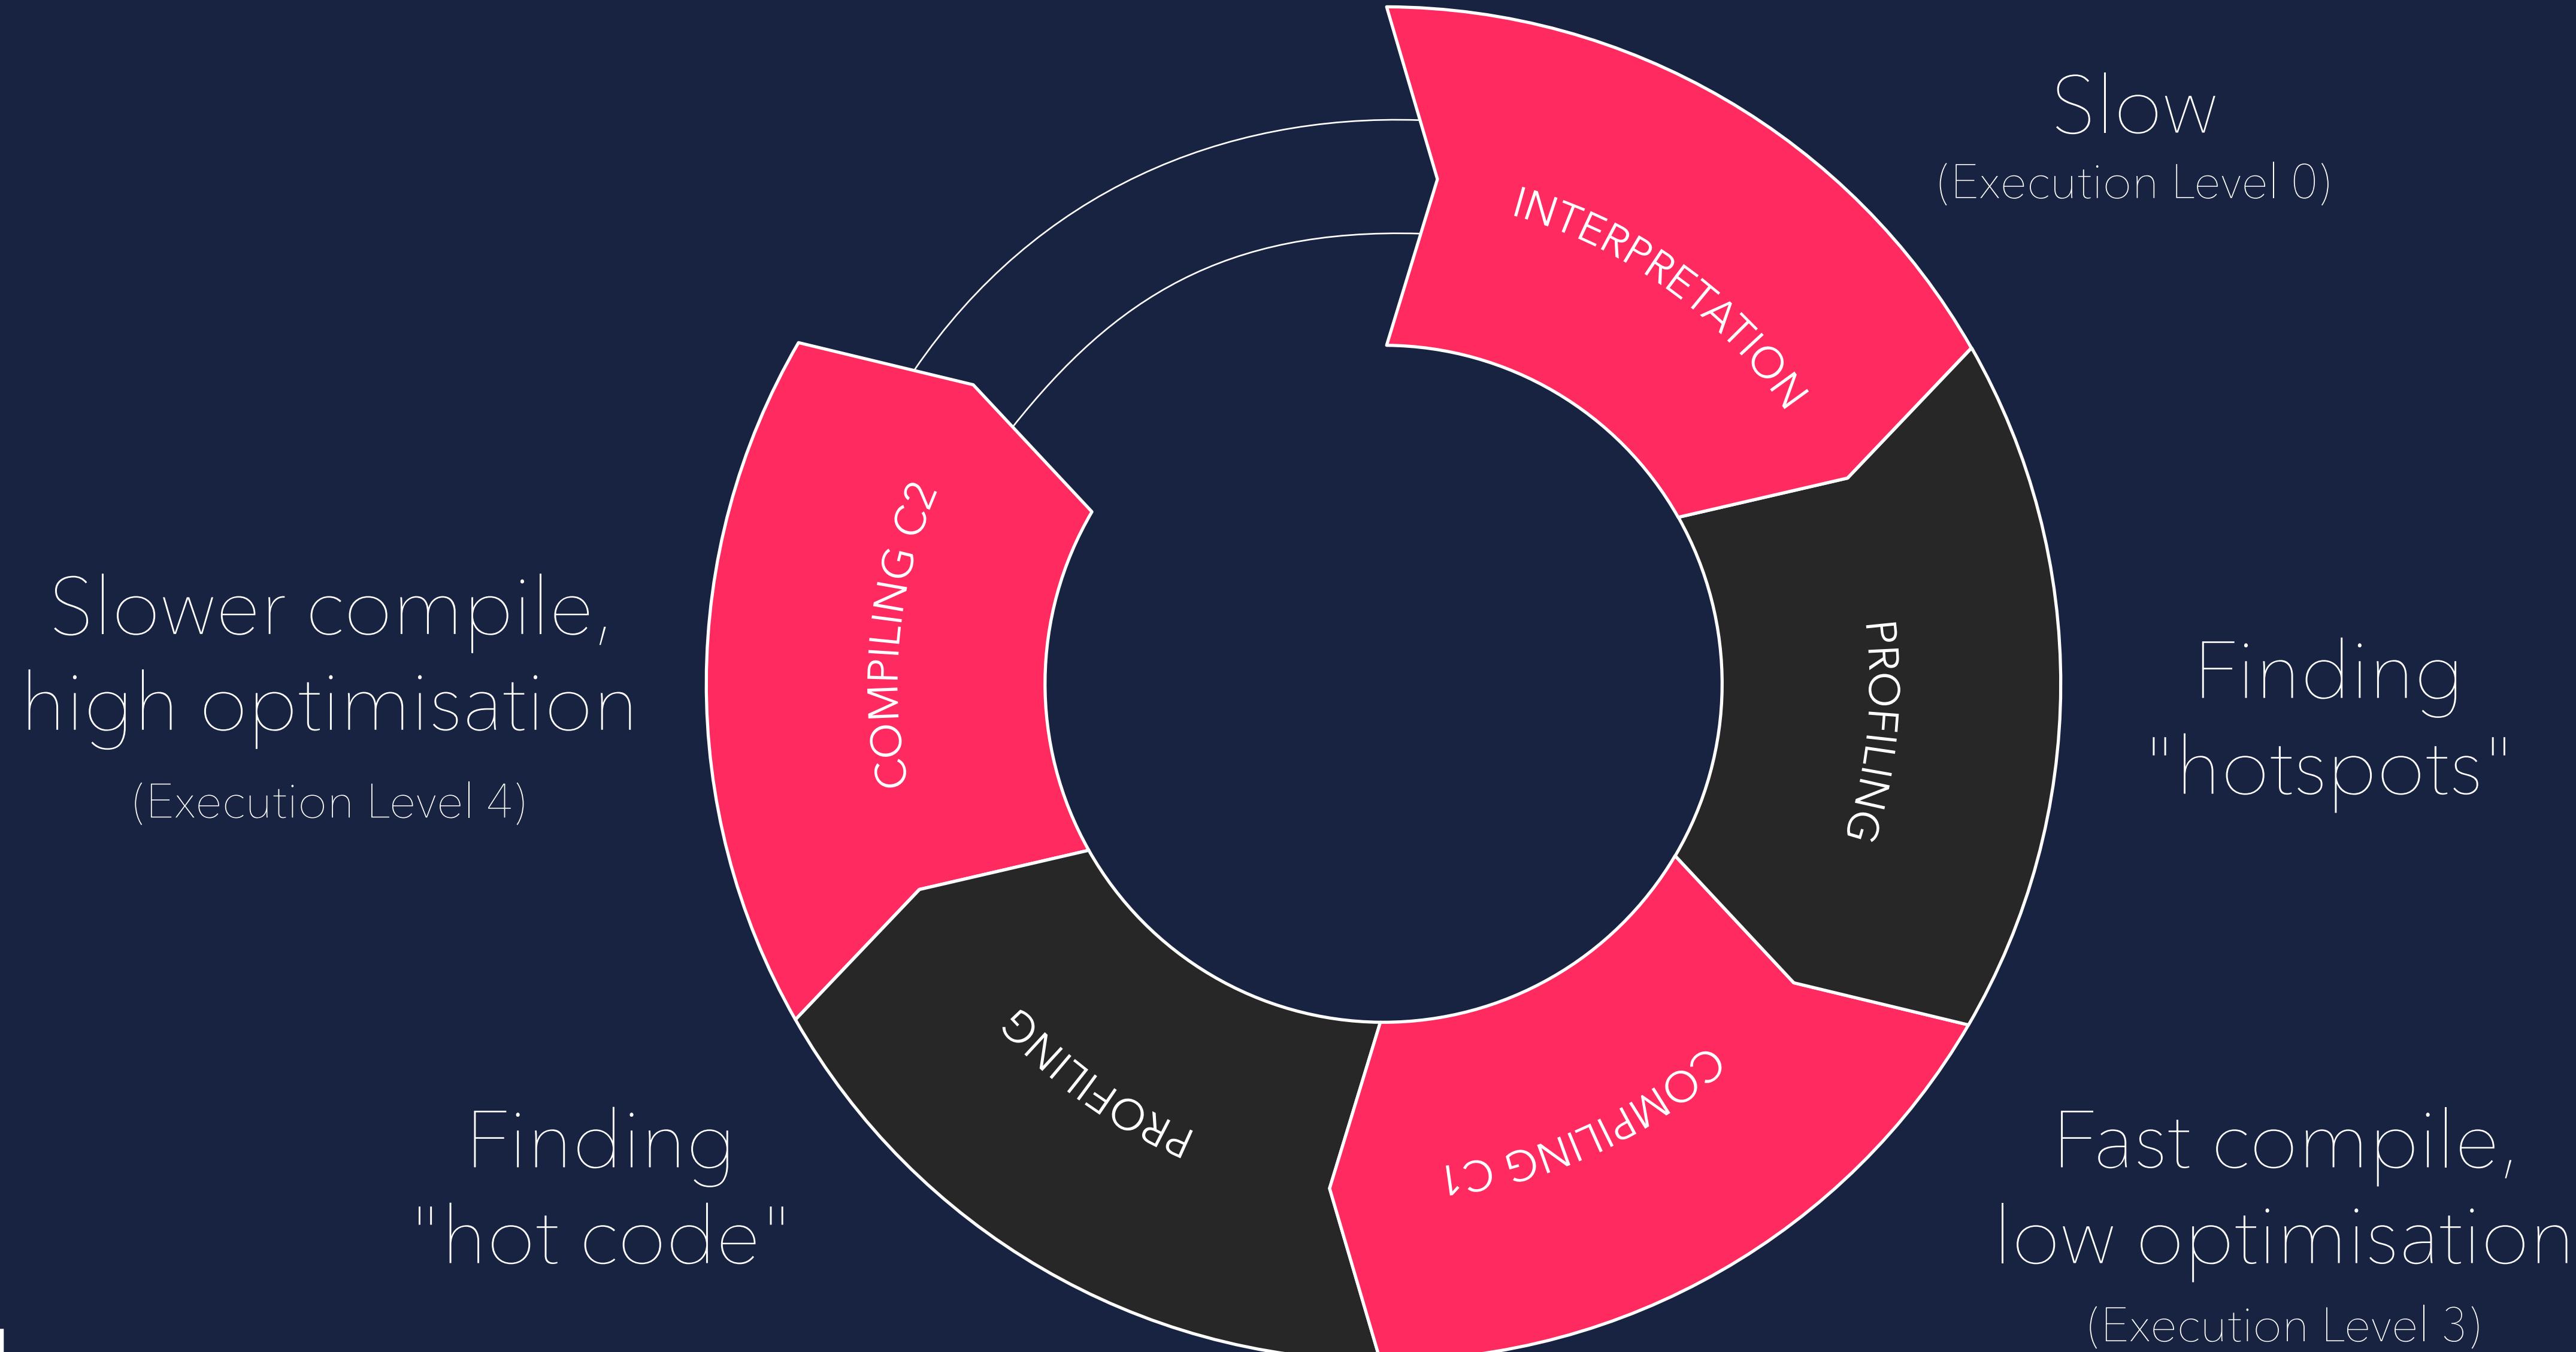
EXECUTION CYCLE



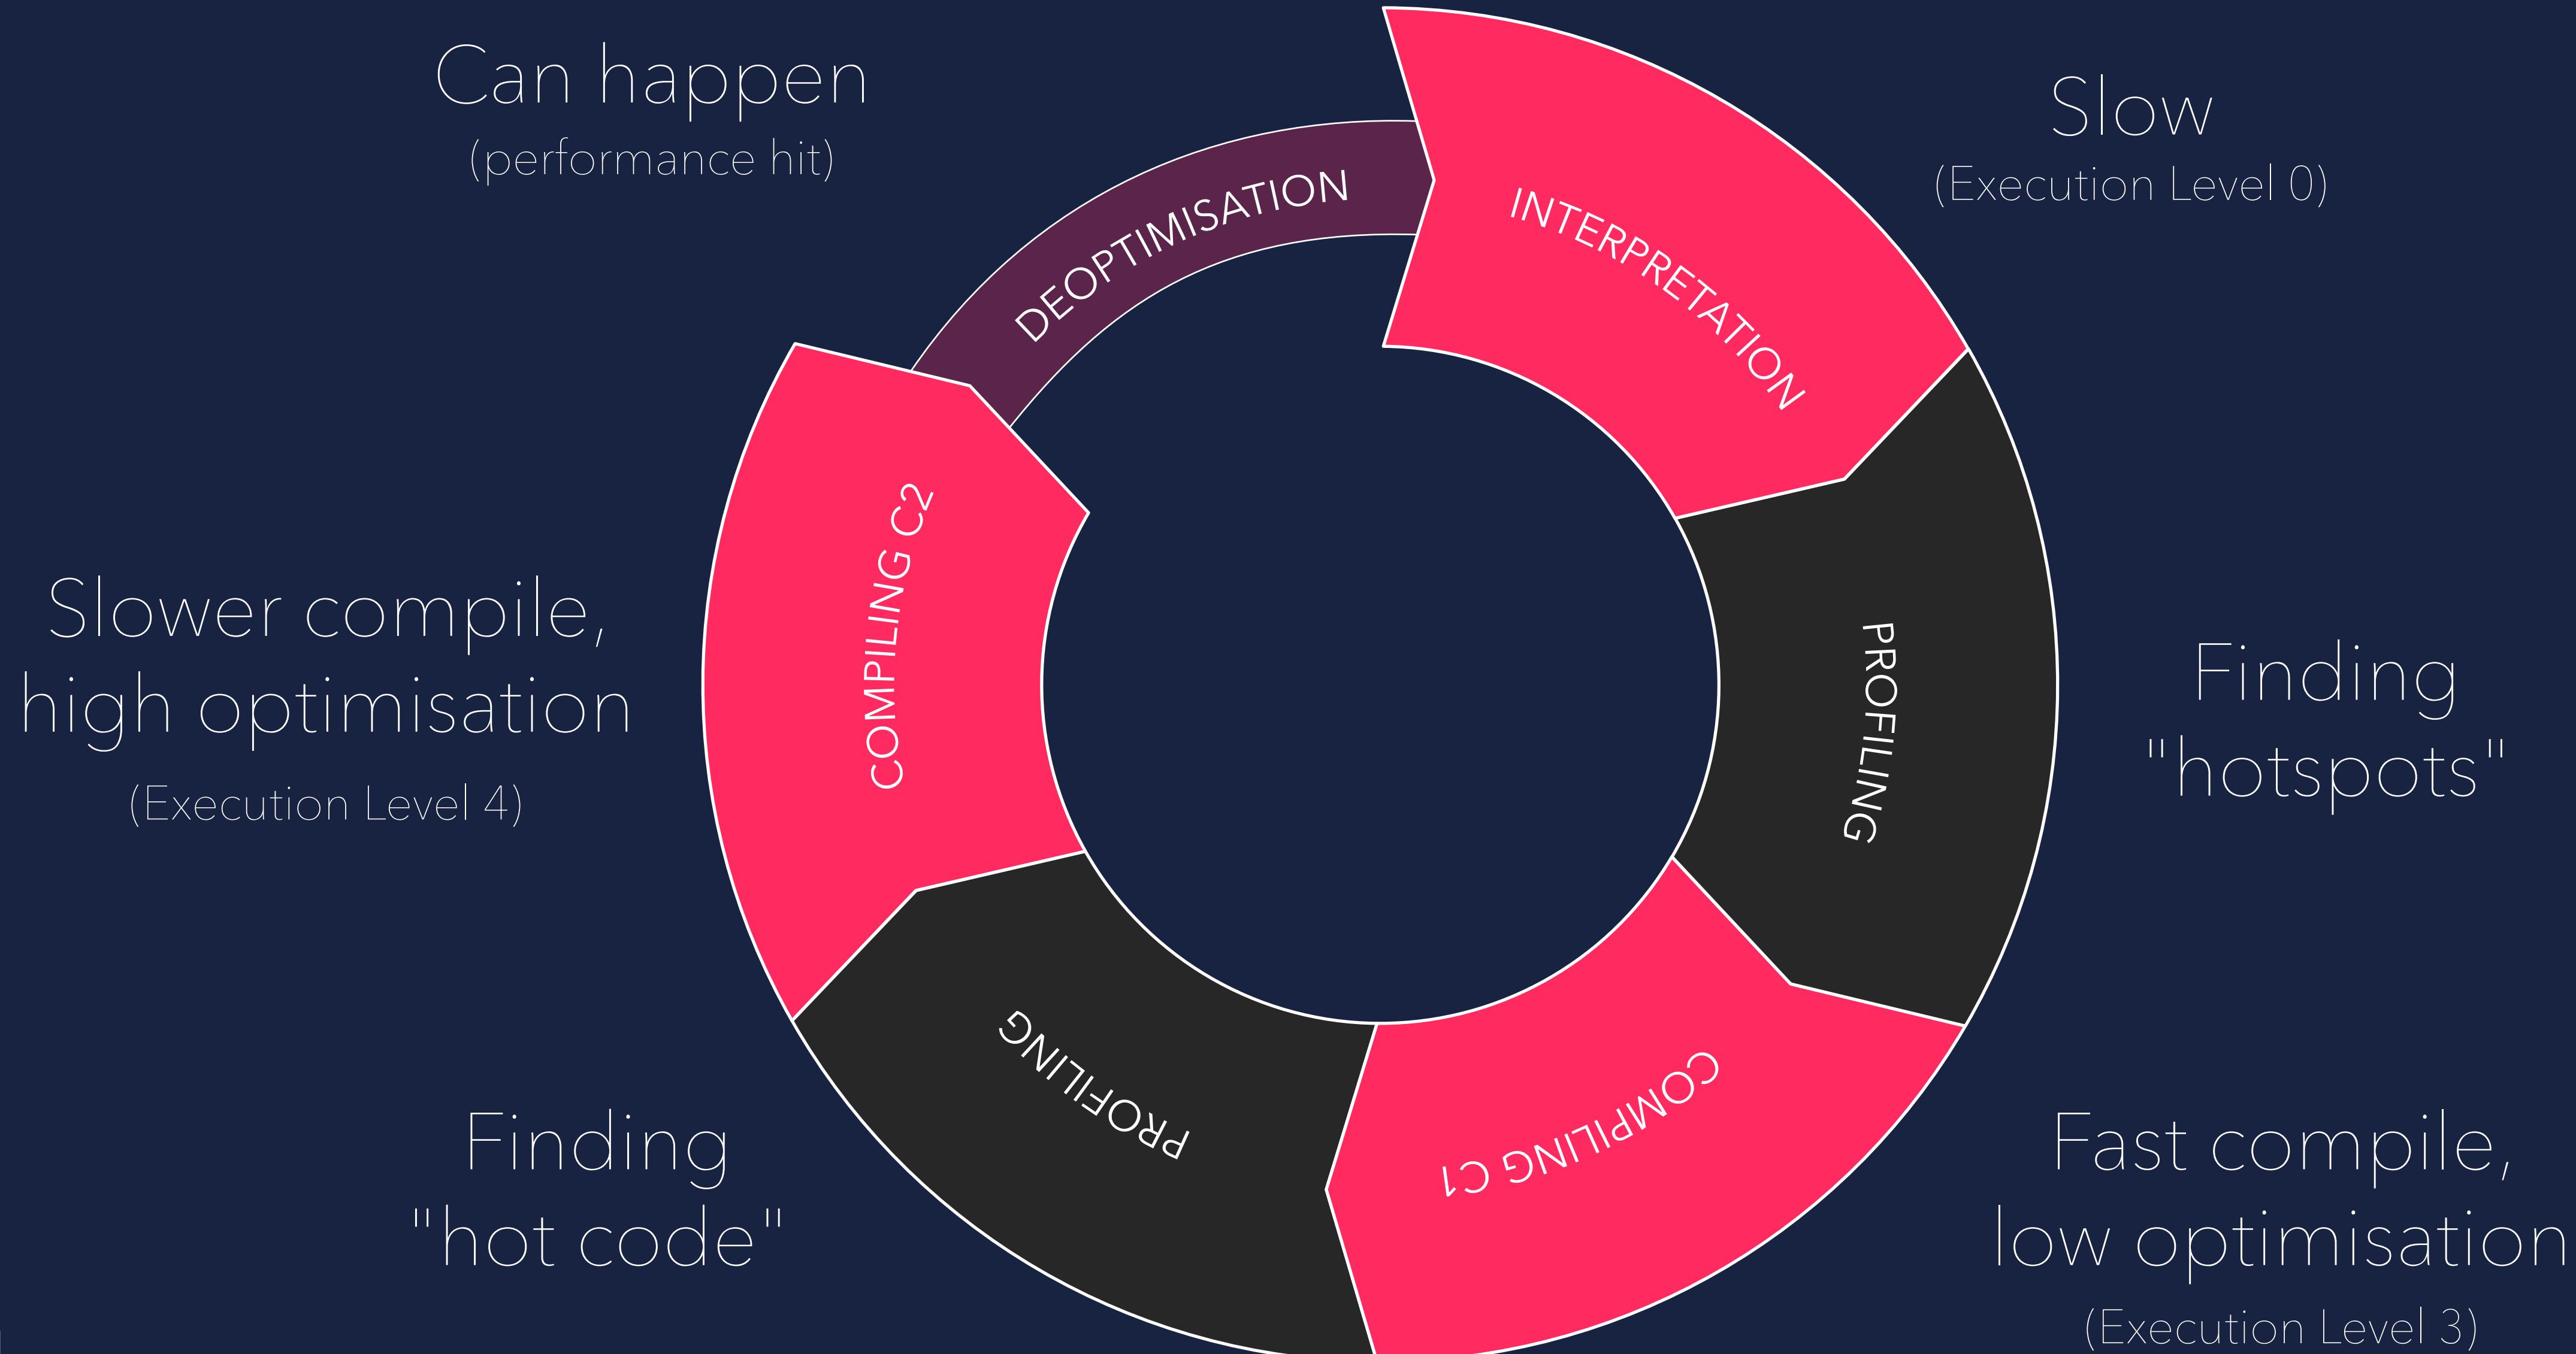
EXECUTION CYCLE



EXECUTION CYCLE



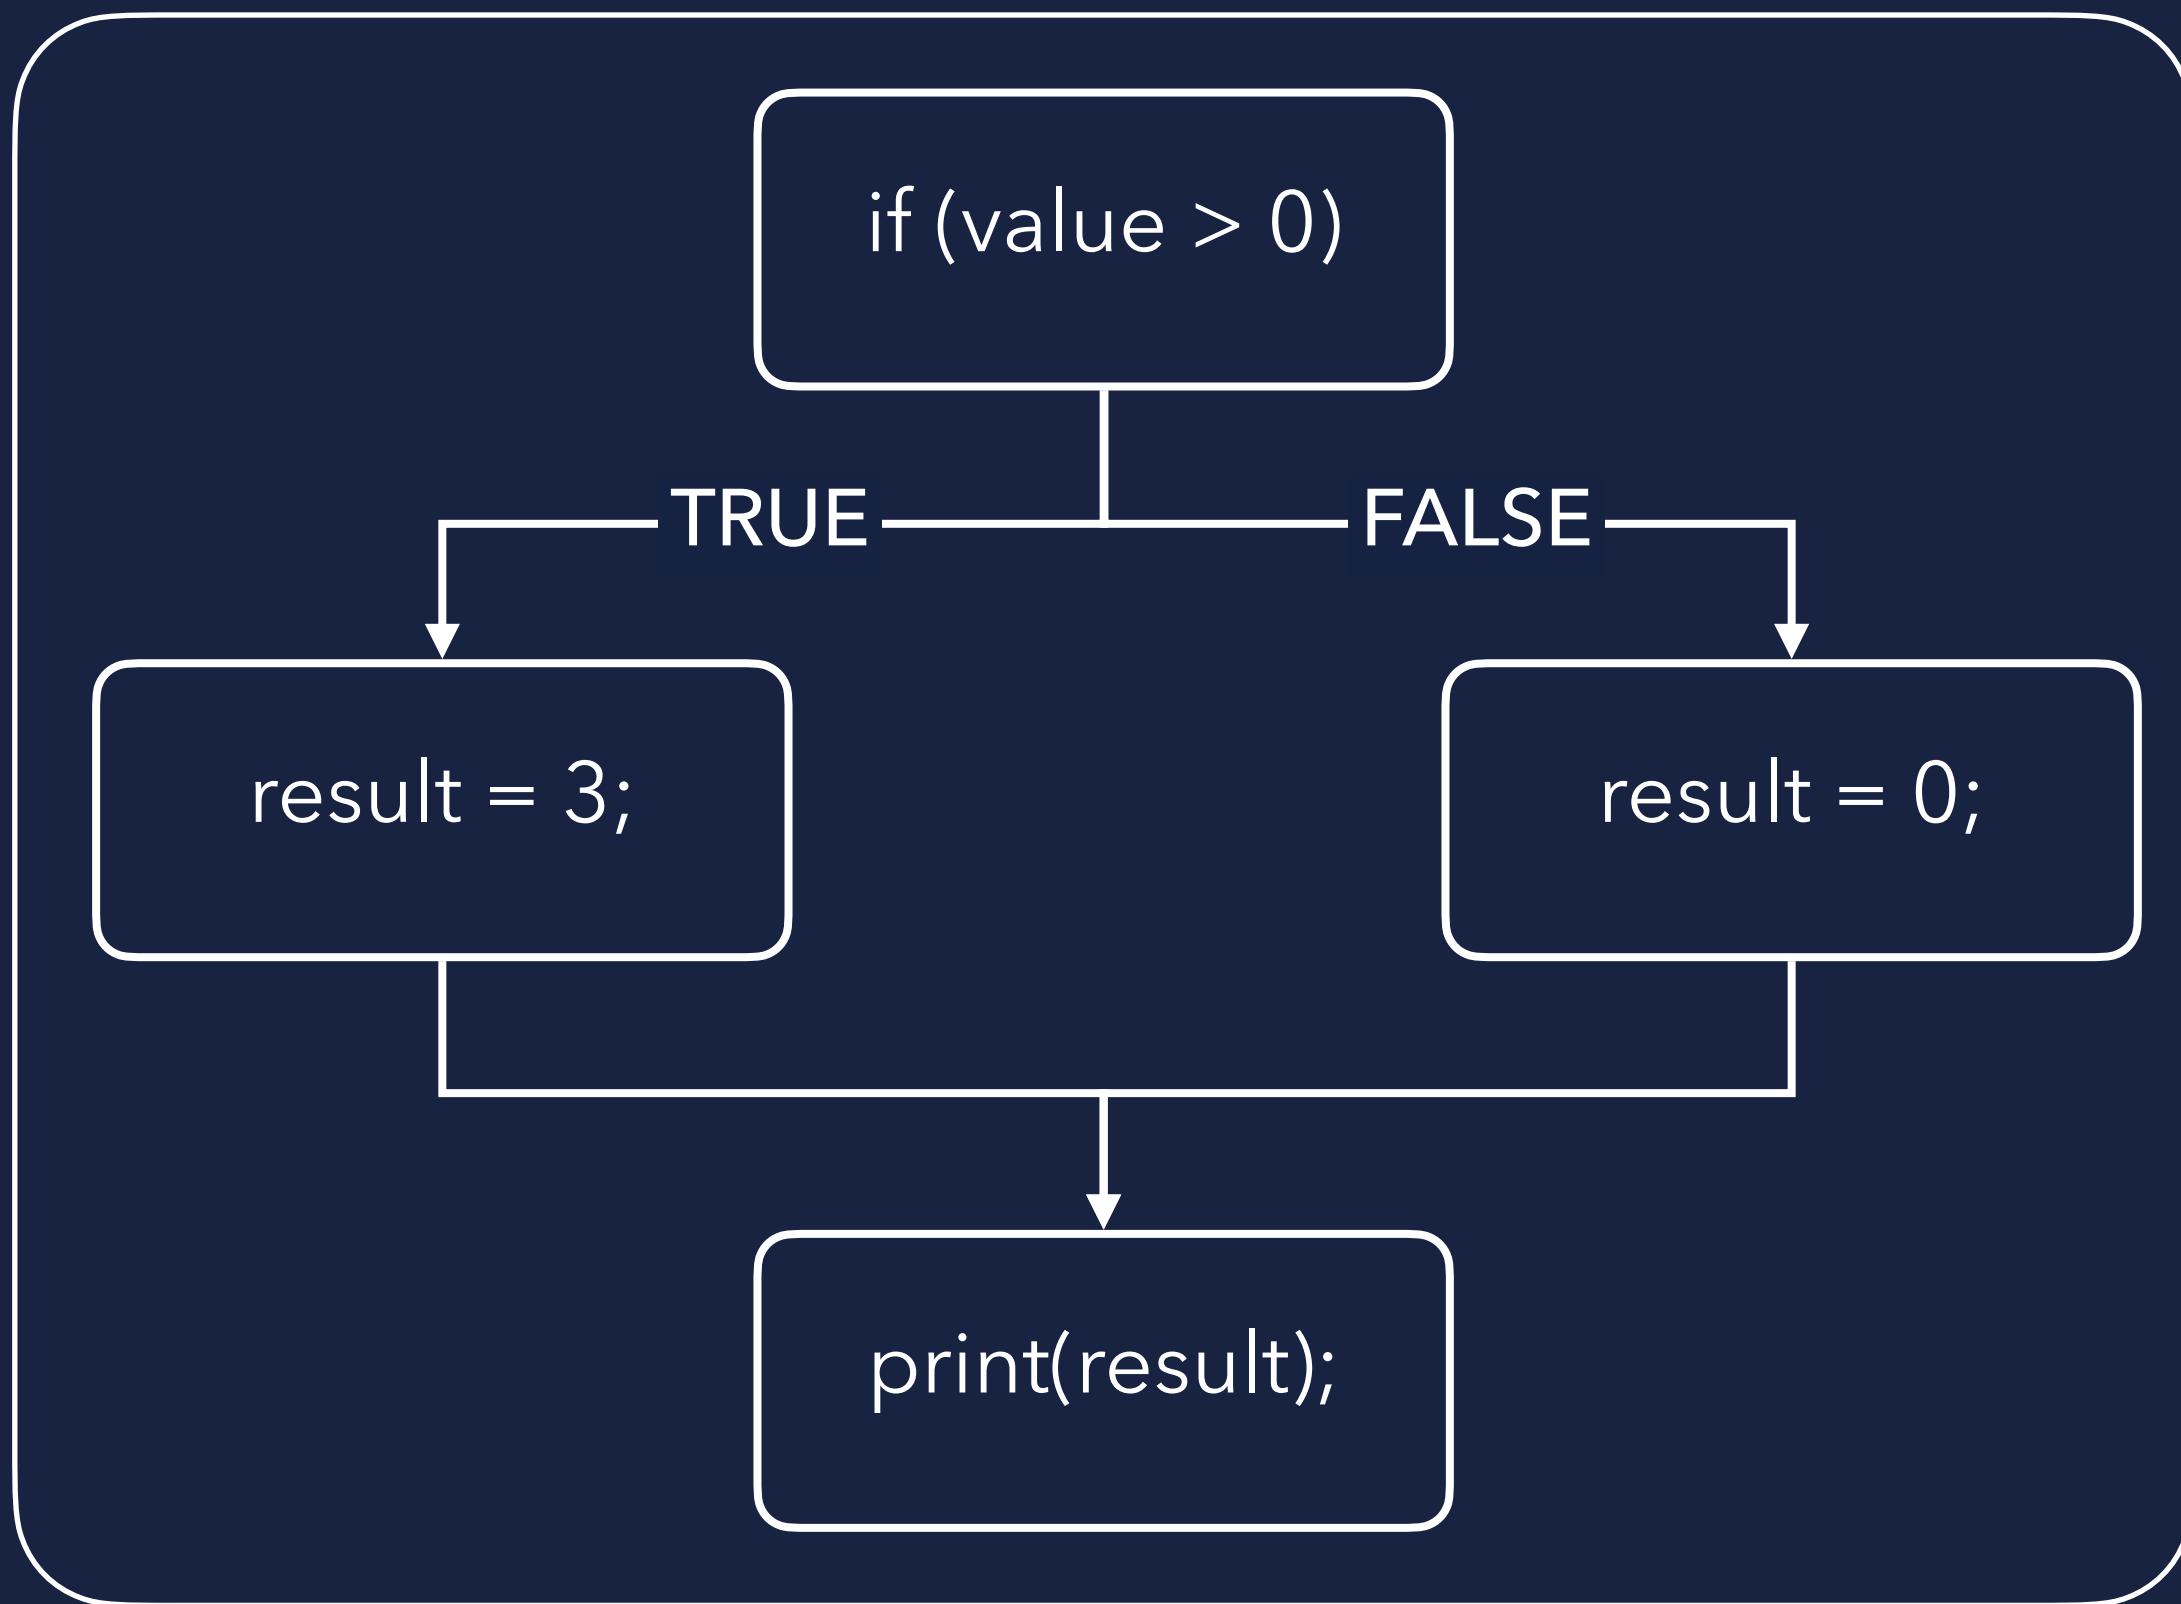
EXECUTION CYCLE



DEOPTIMISATION

DEOPTIMISATION

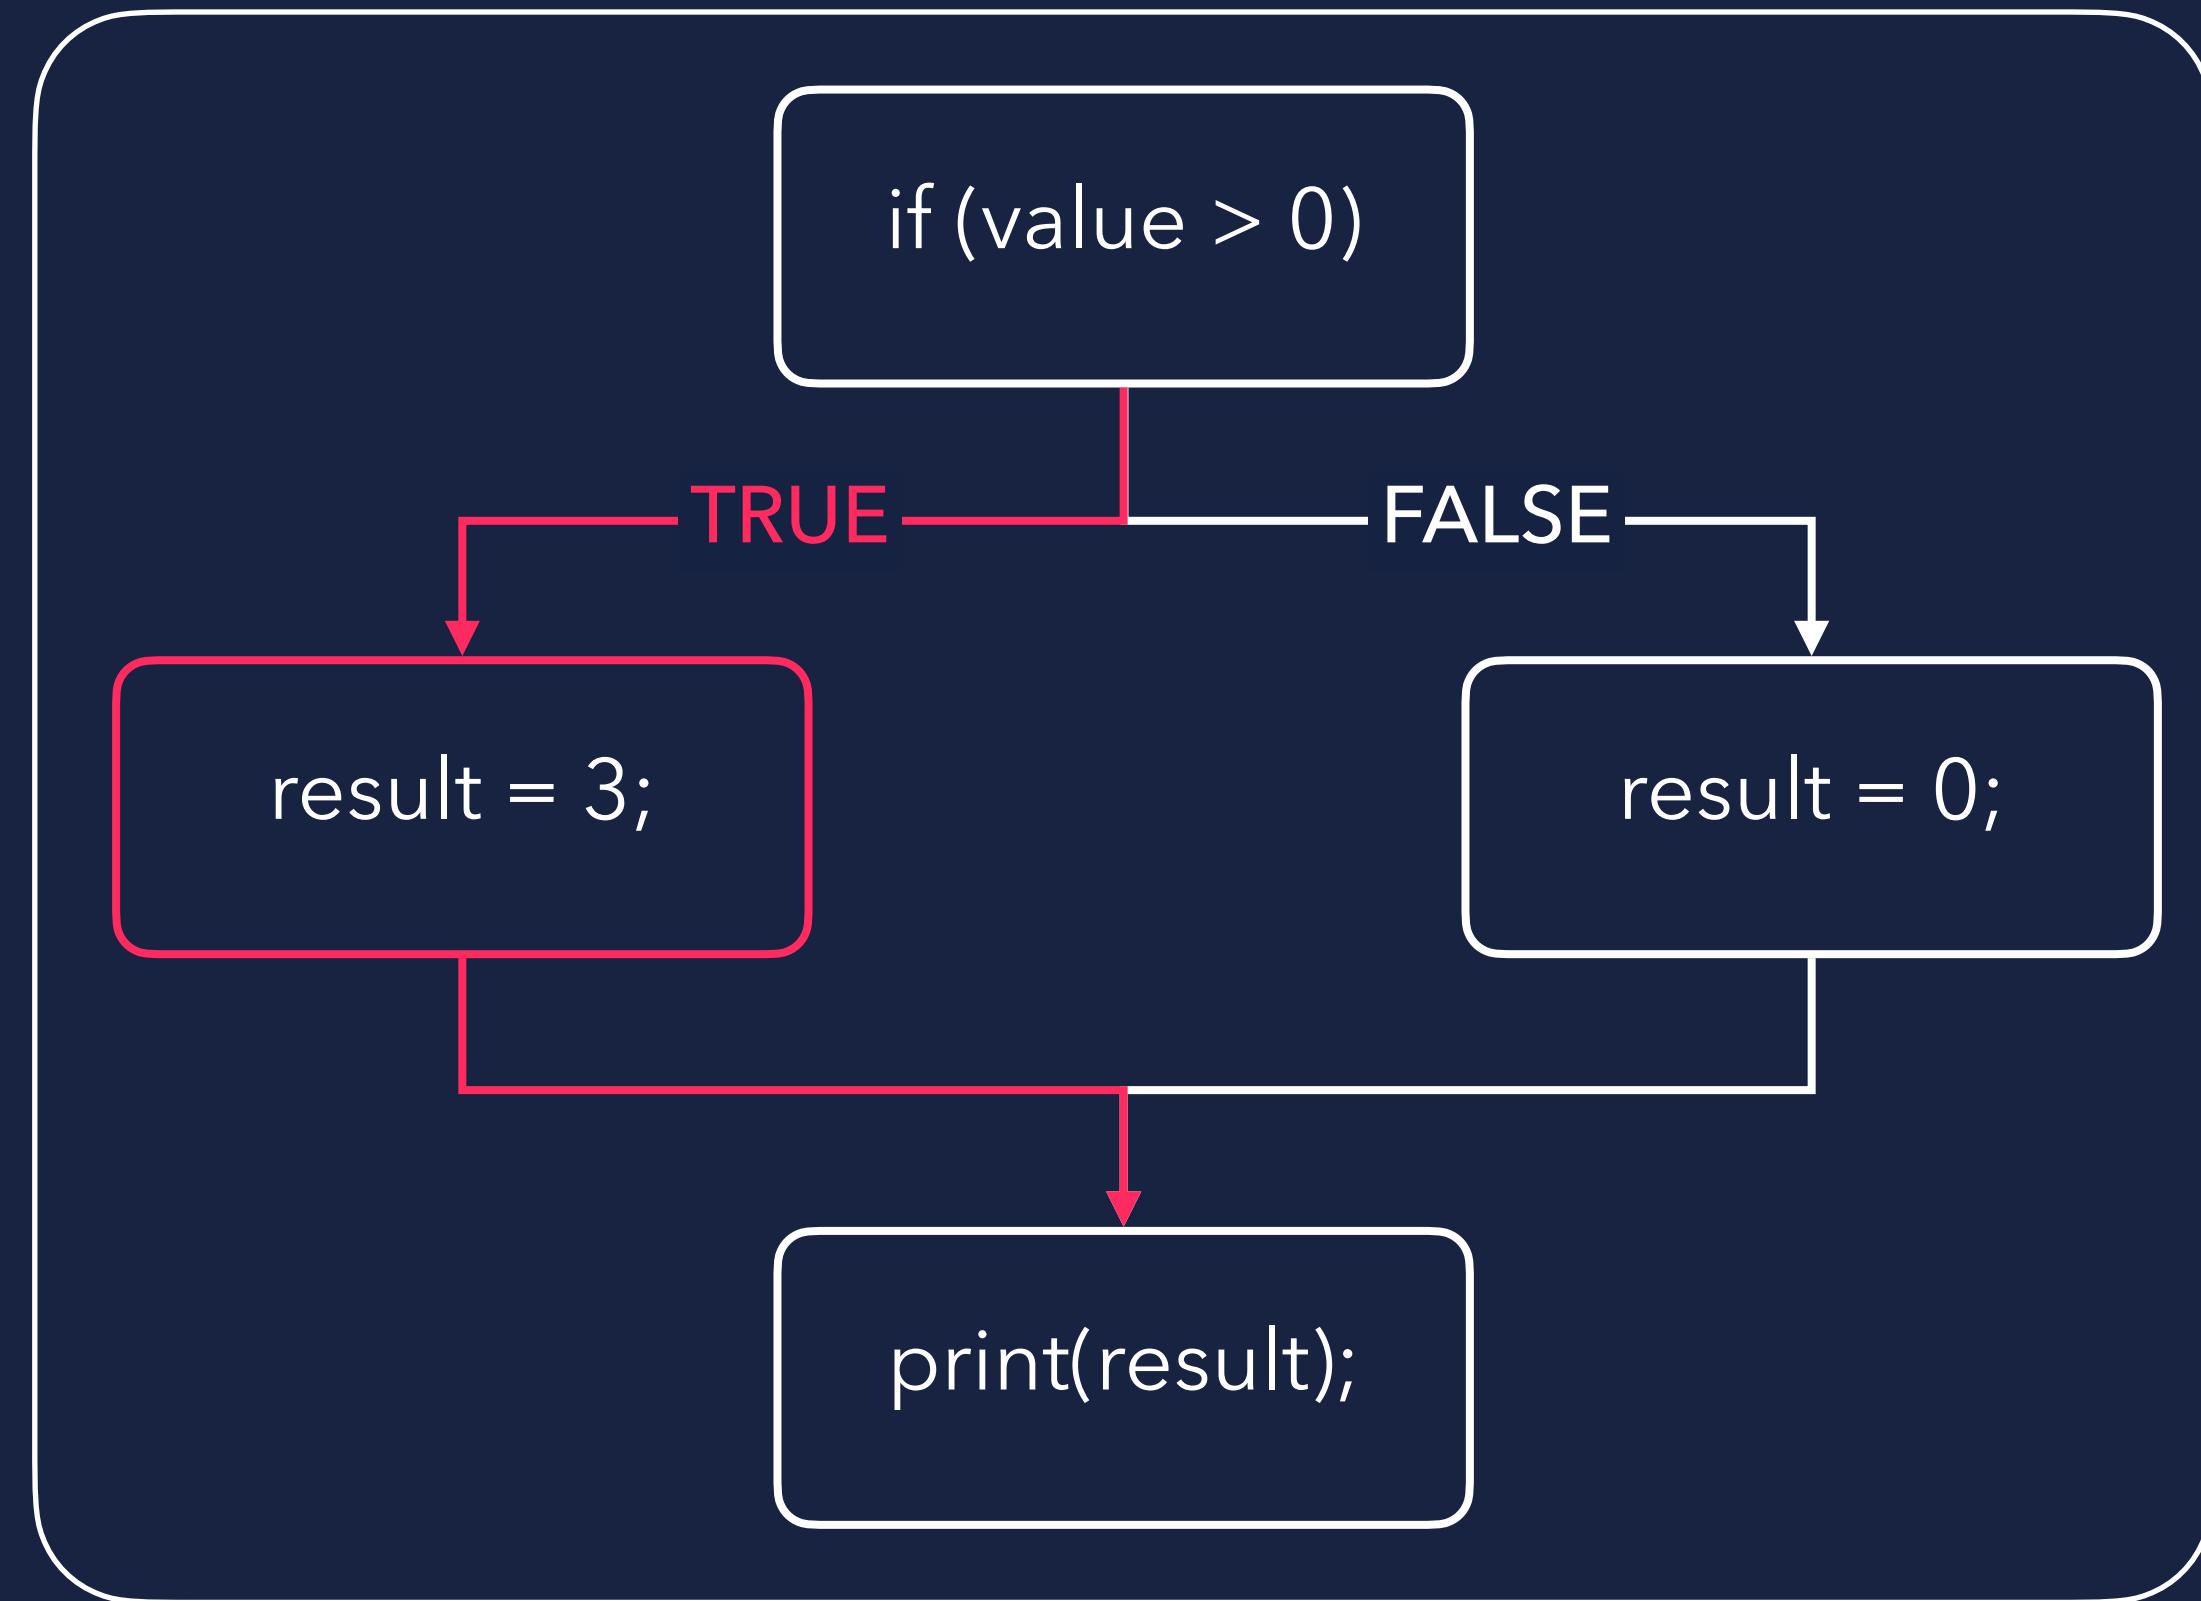
CODE



DEOPTIMISATION

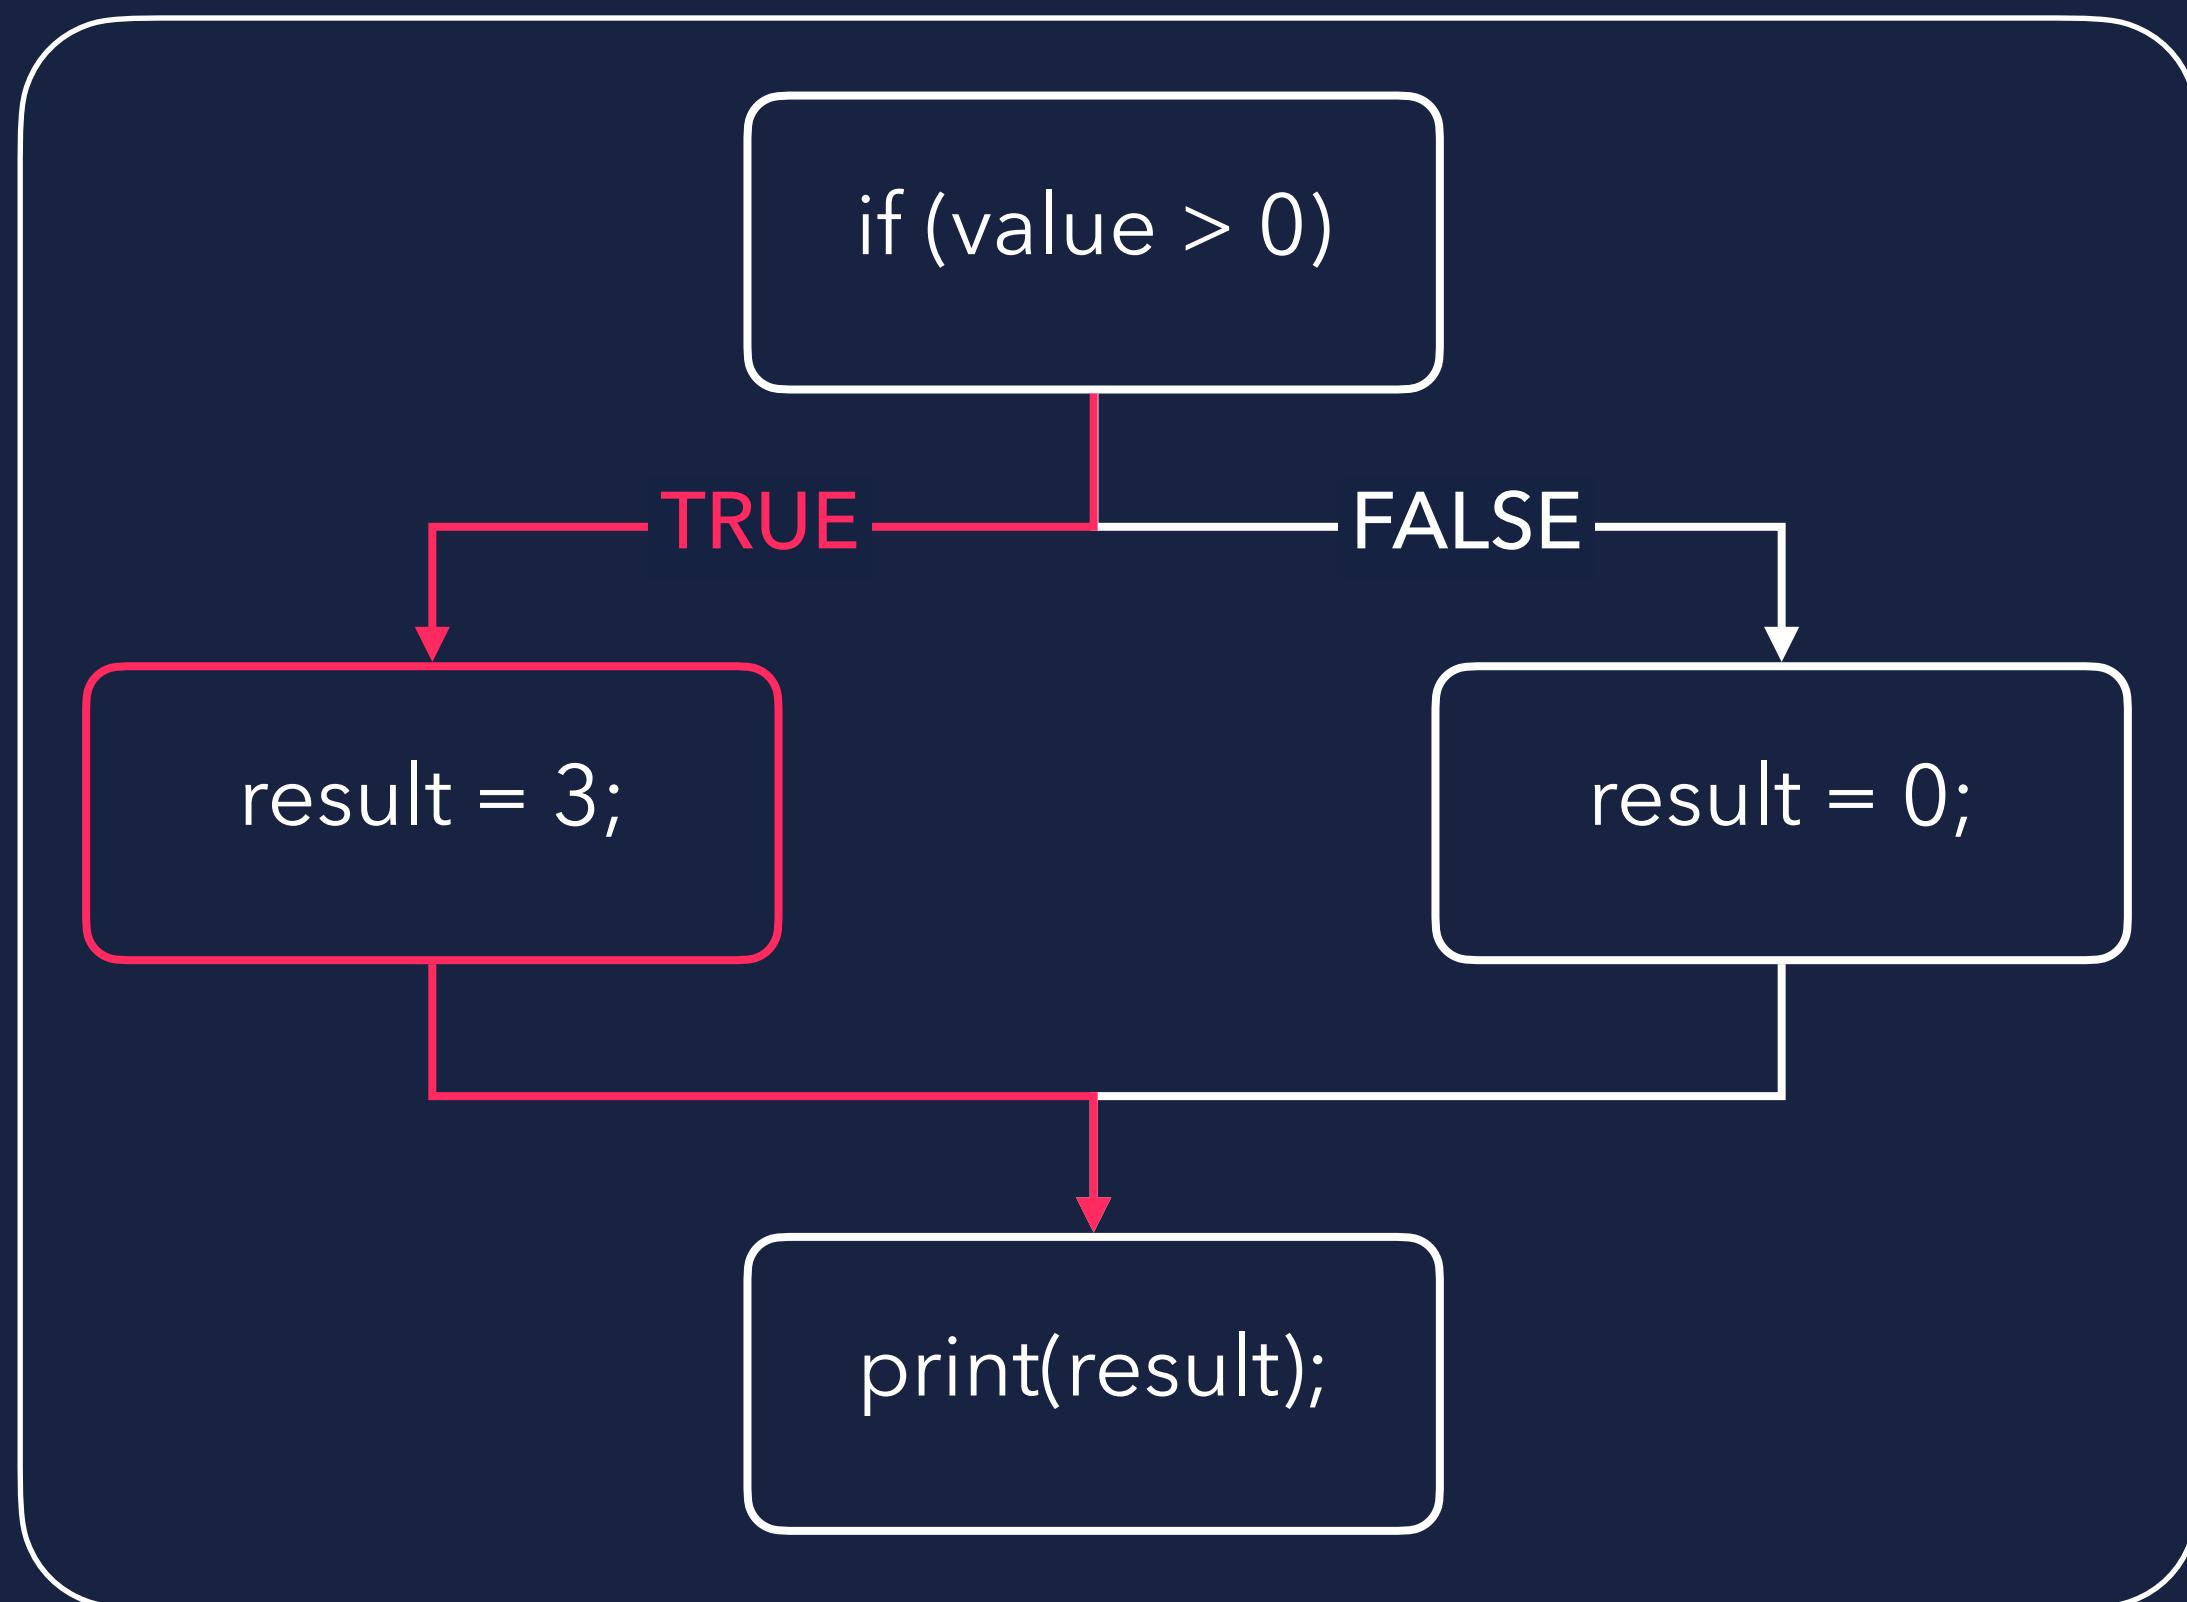
CODE

Hot Path (profiled)

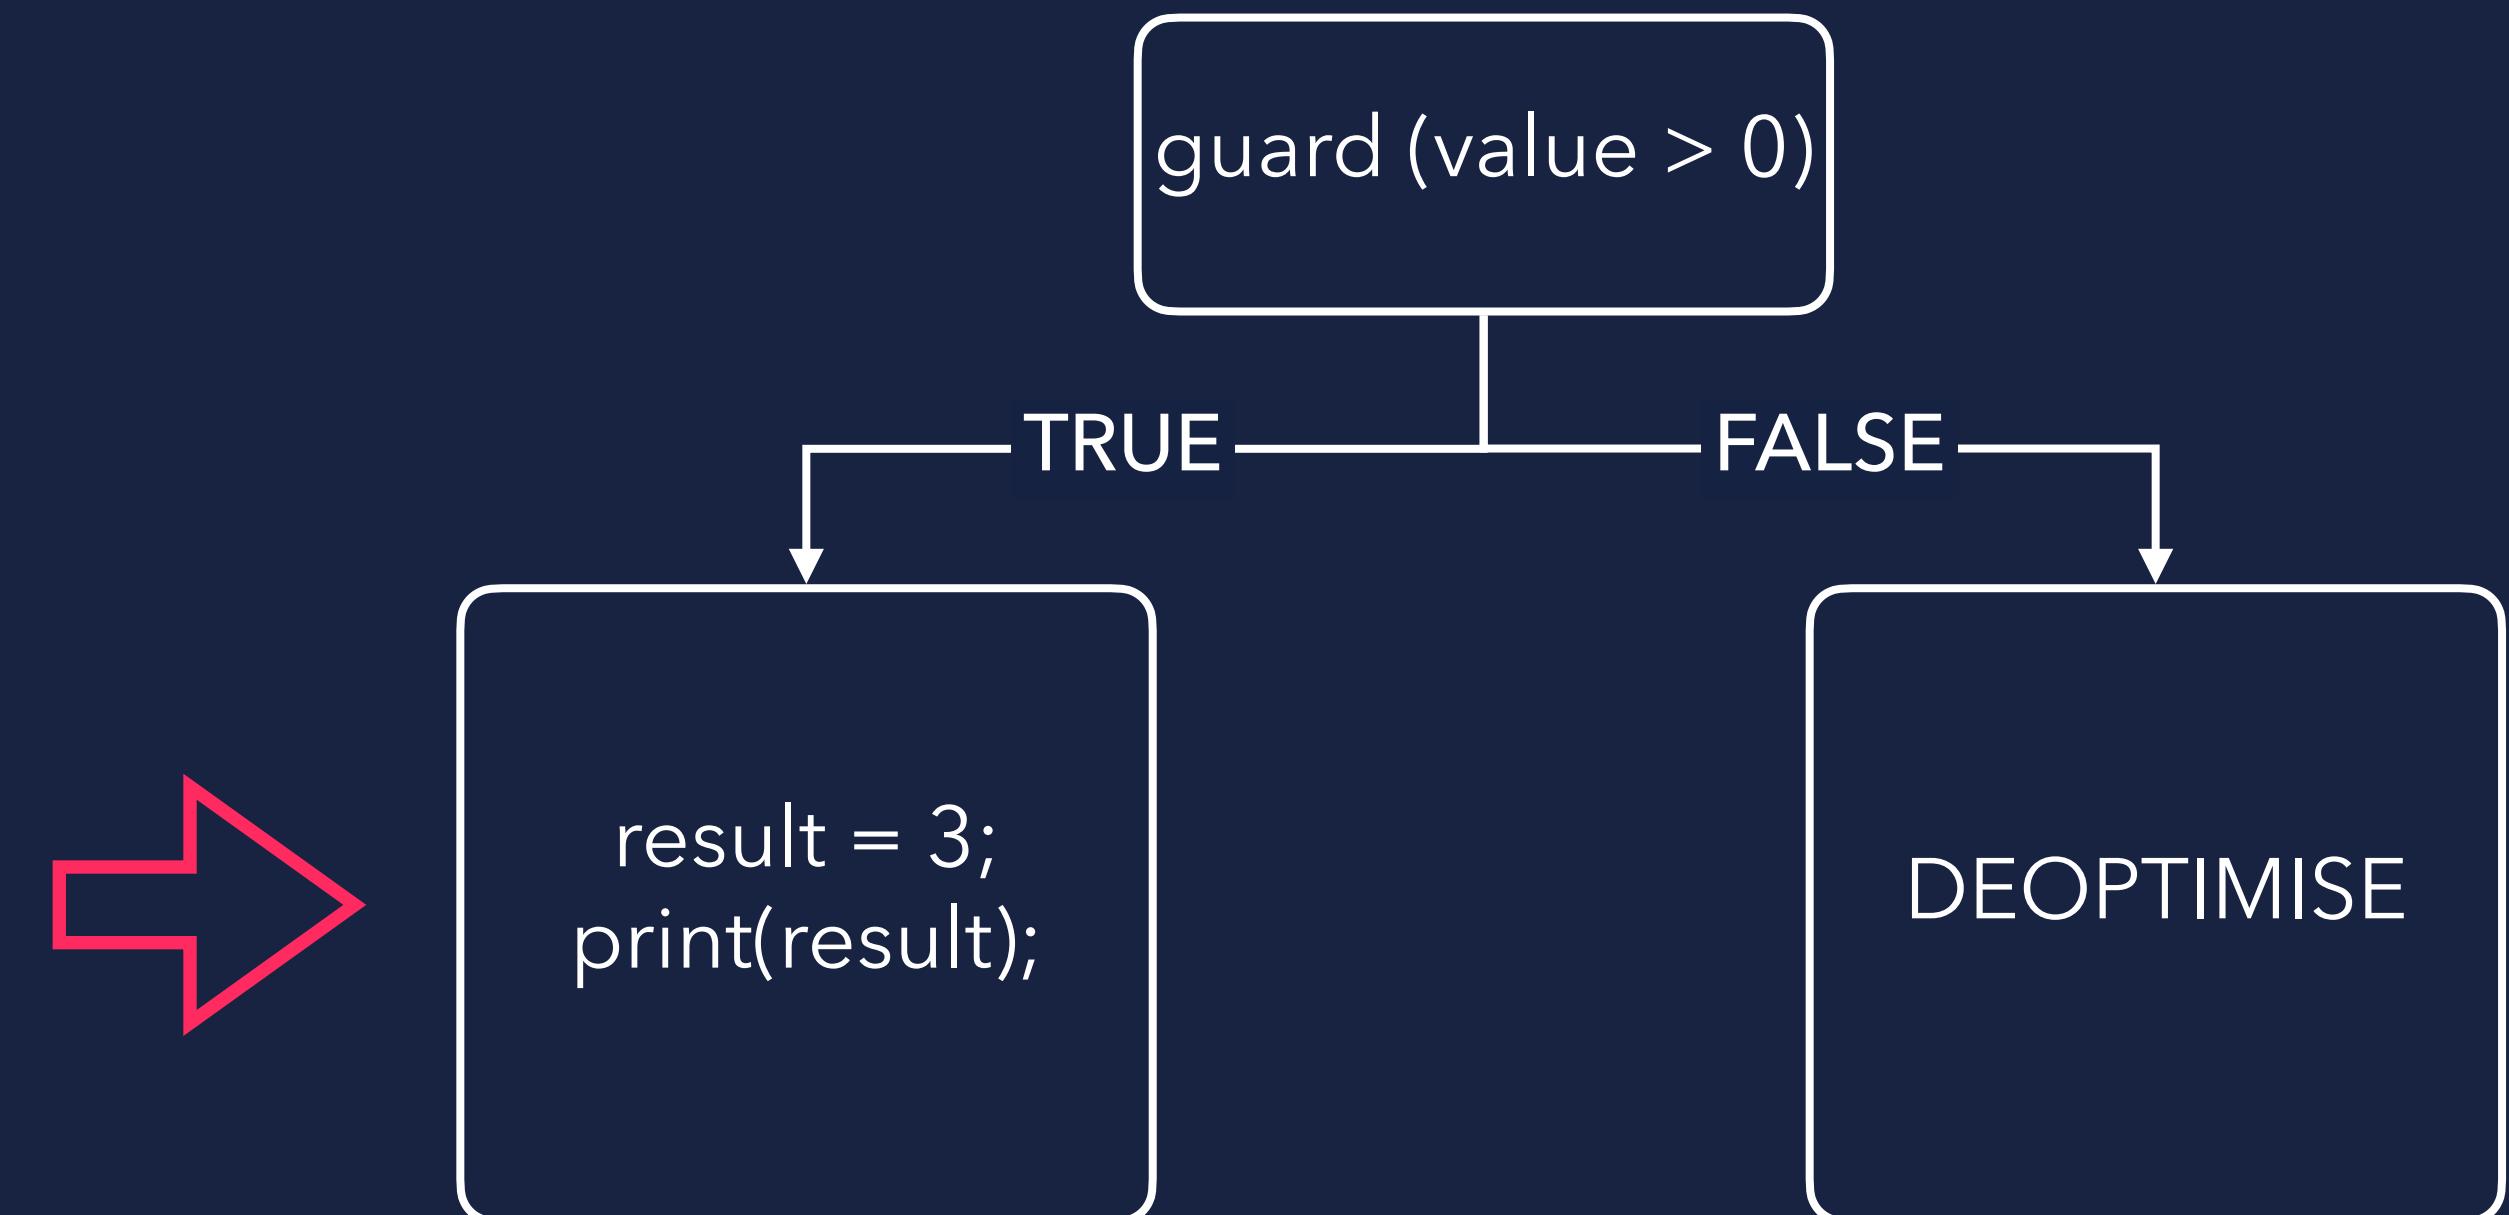


DEOPTIMISATION

CODE

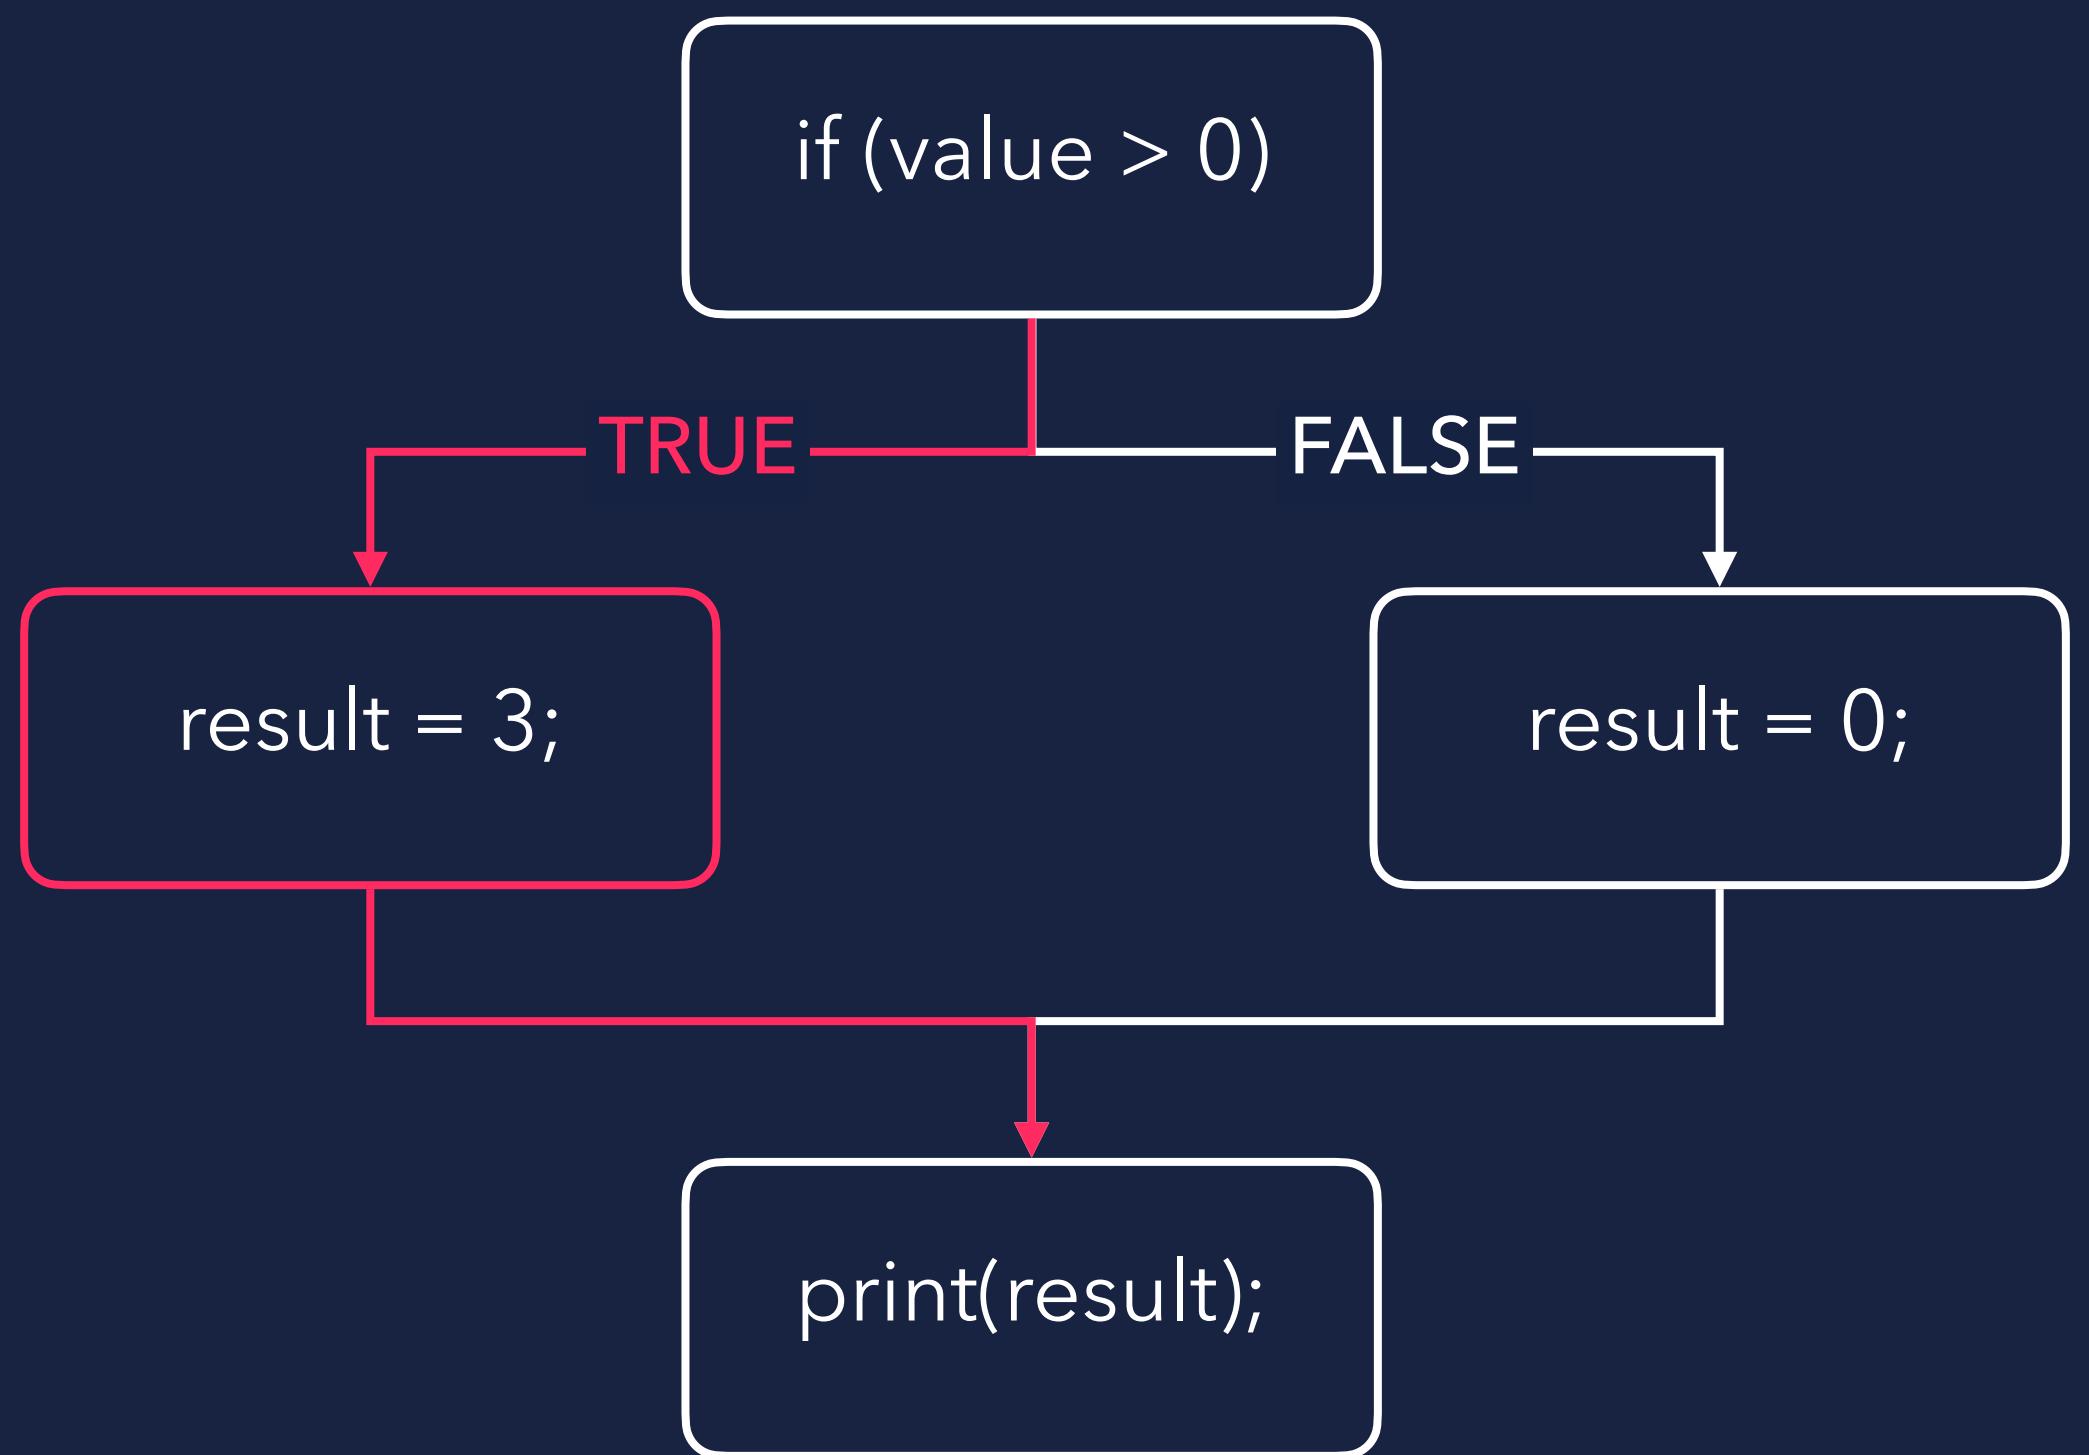


OPTIMIZED CODE



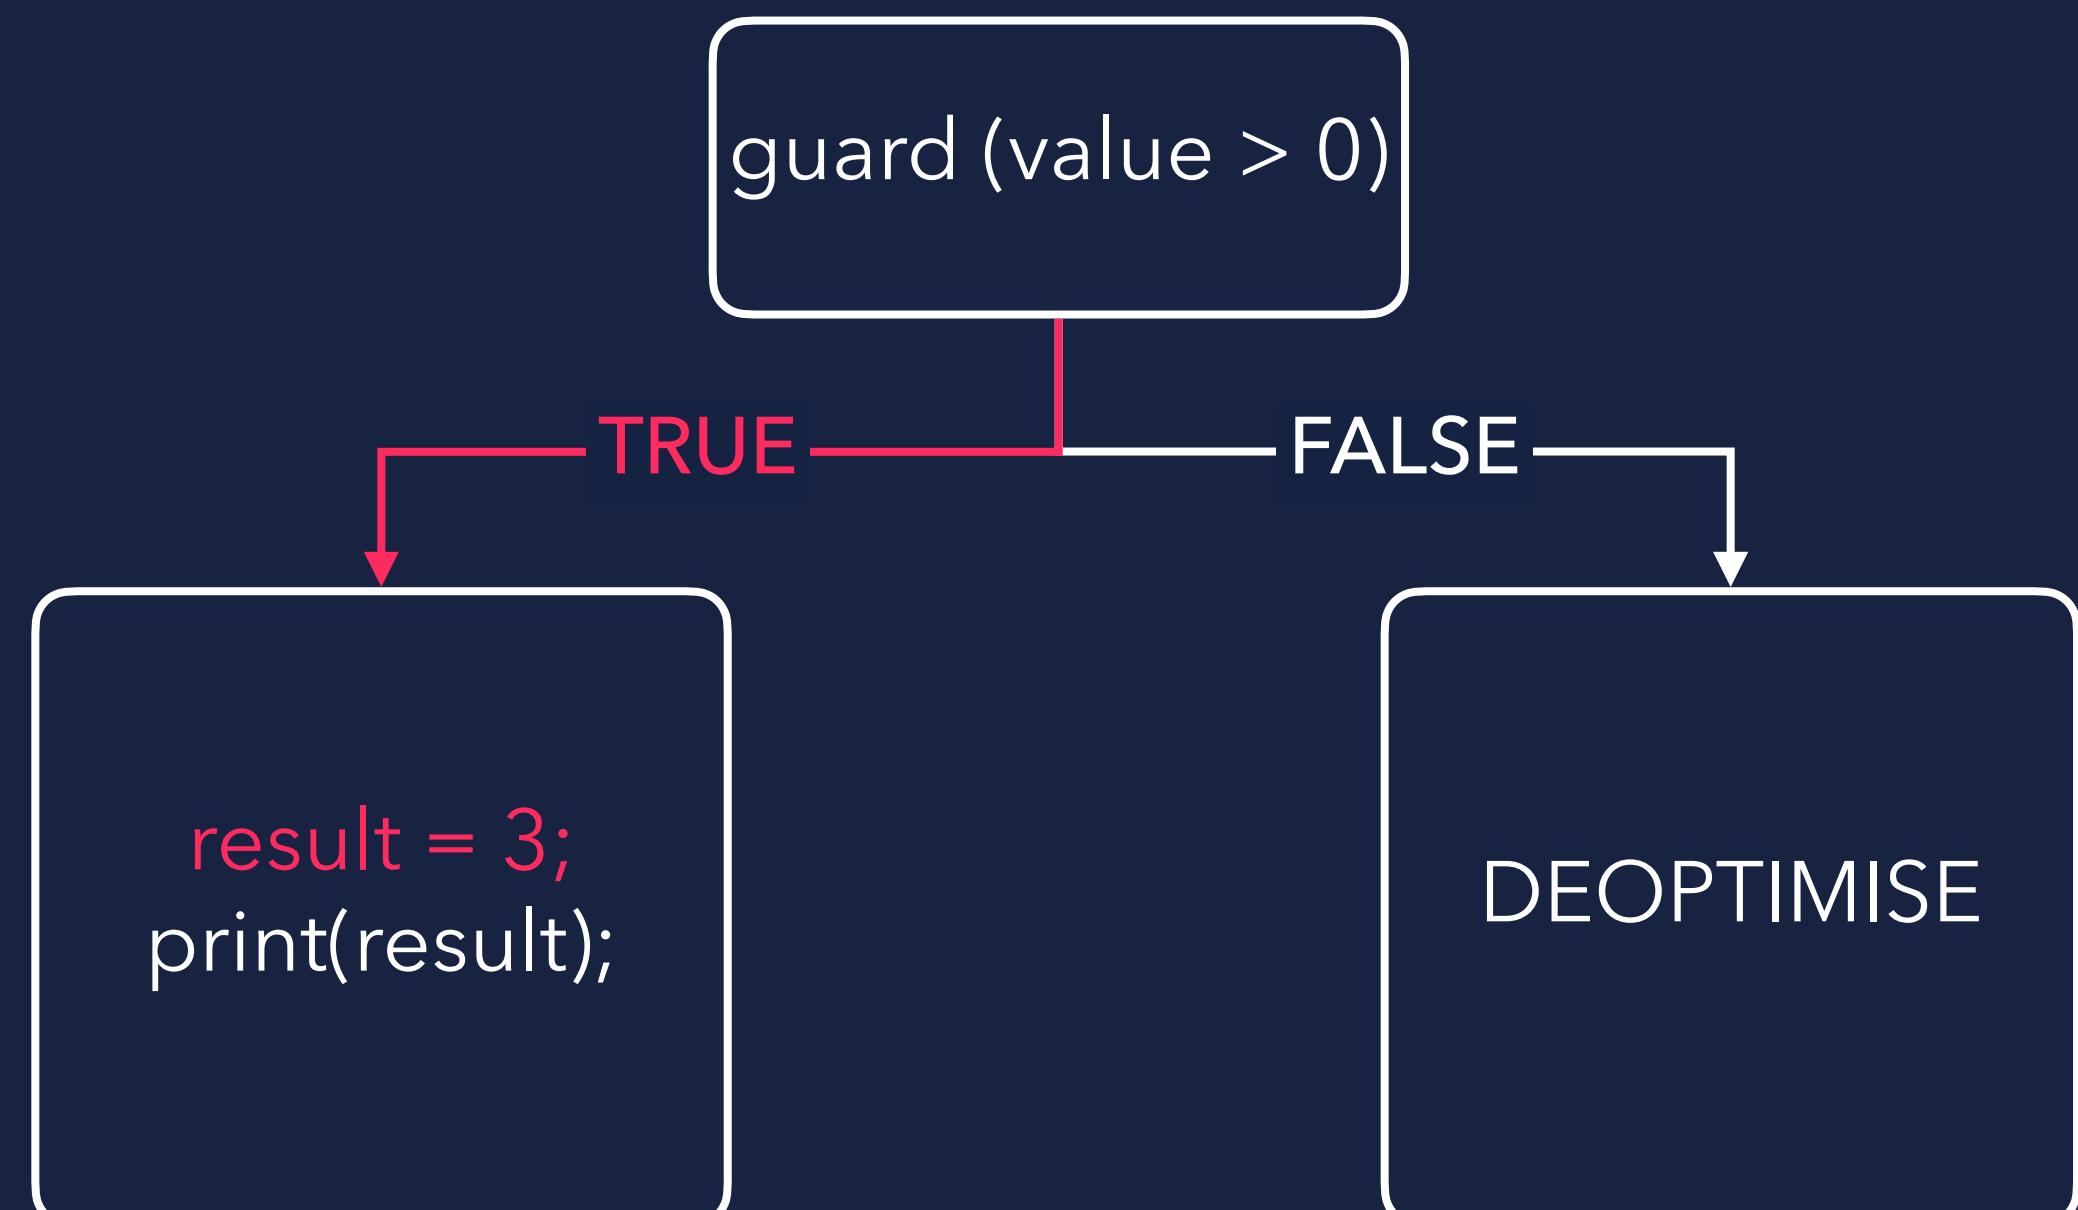
DEOPTIMISATION

CODE



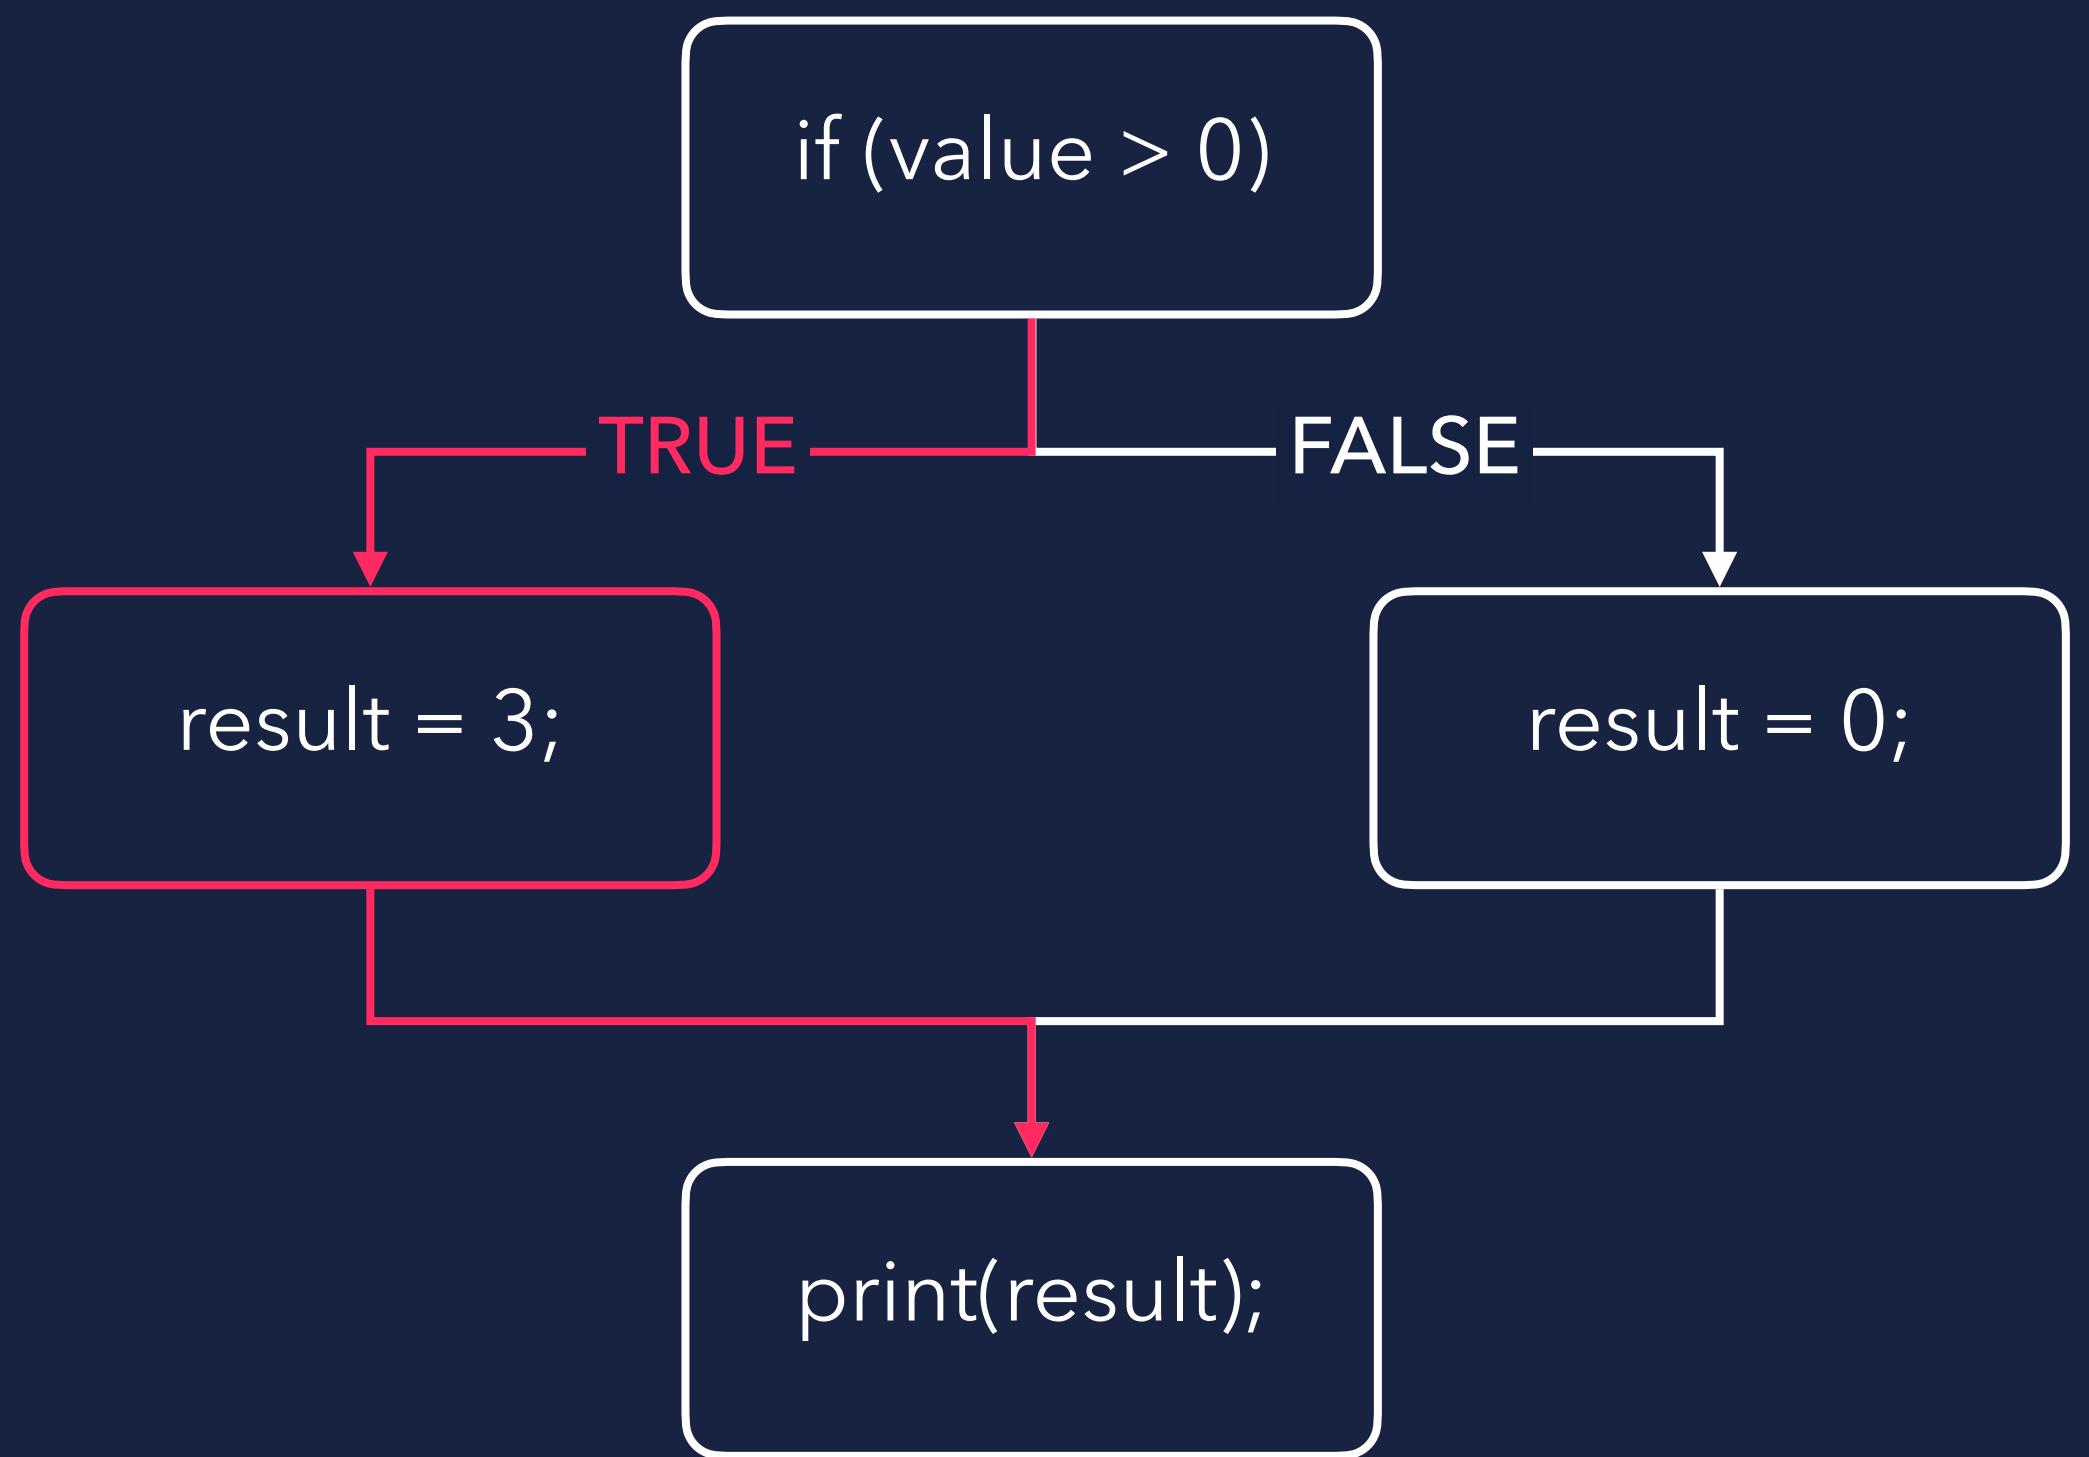
OPTIMIZED CODE

`value = 3`



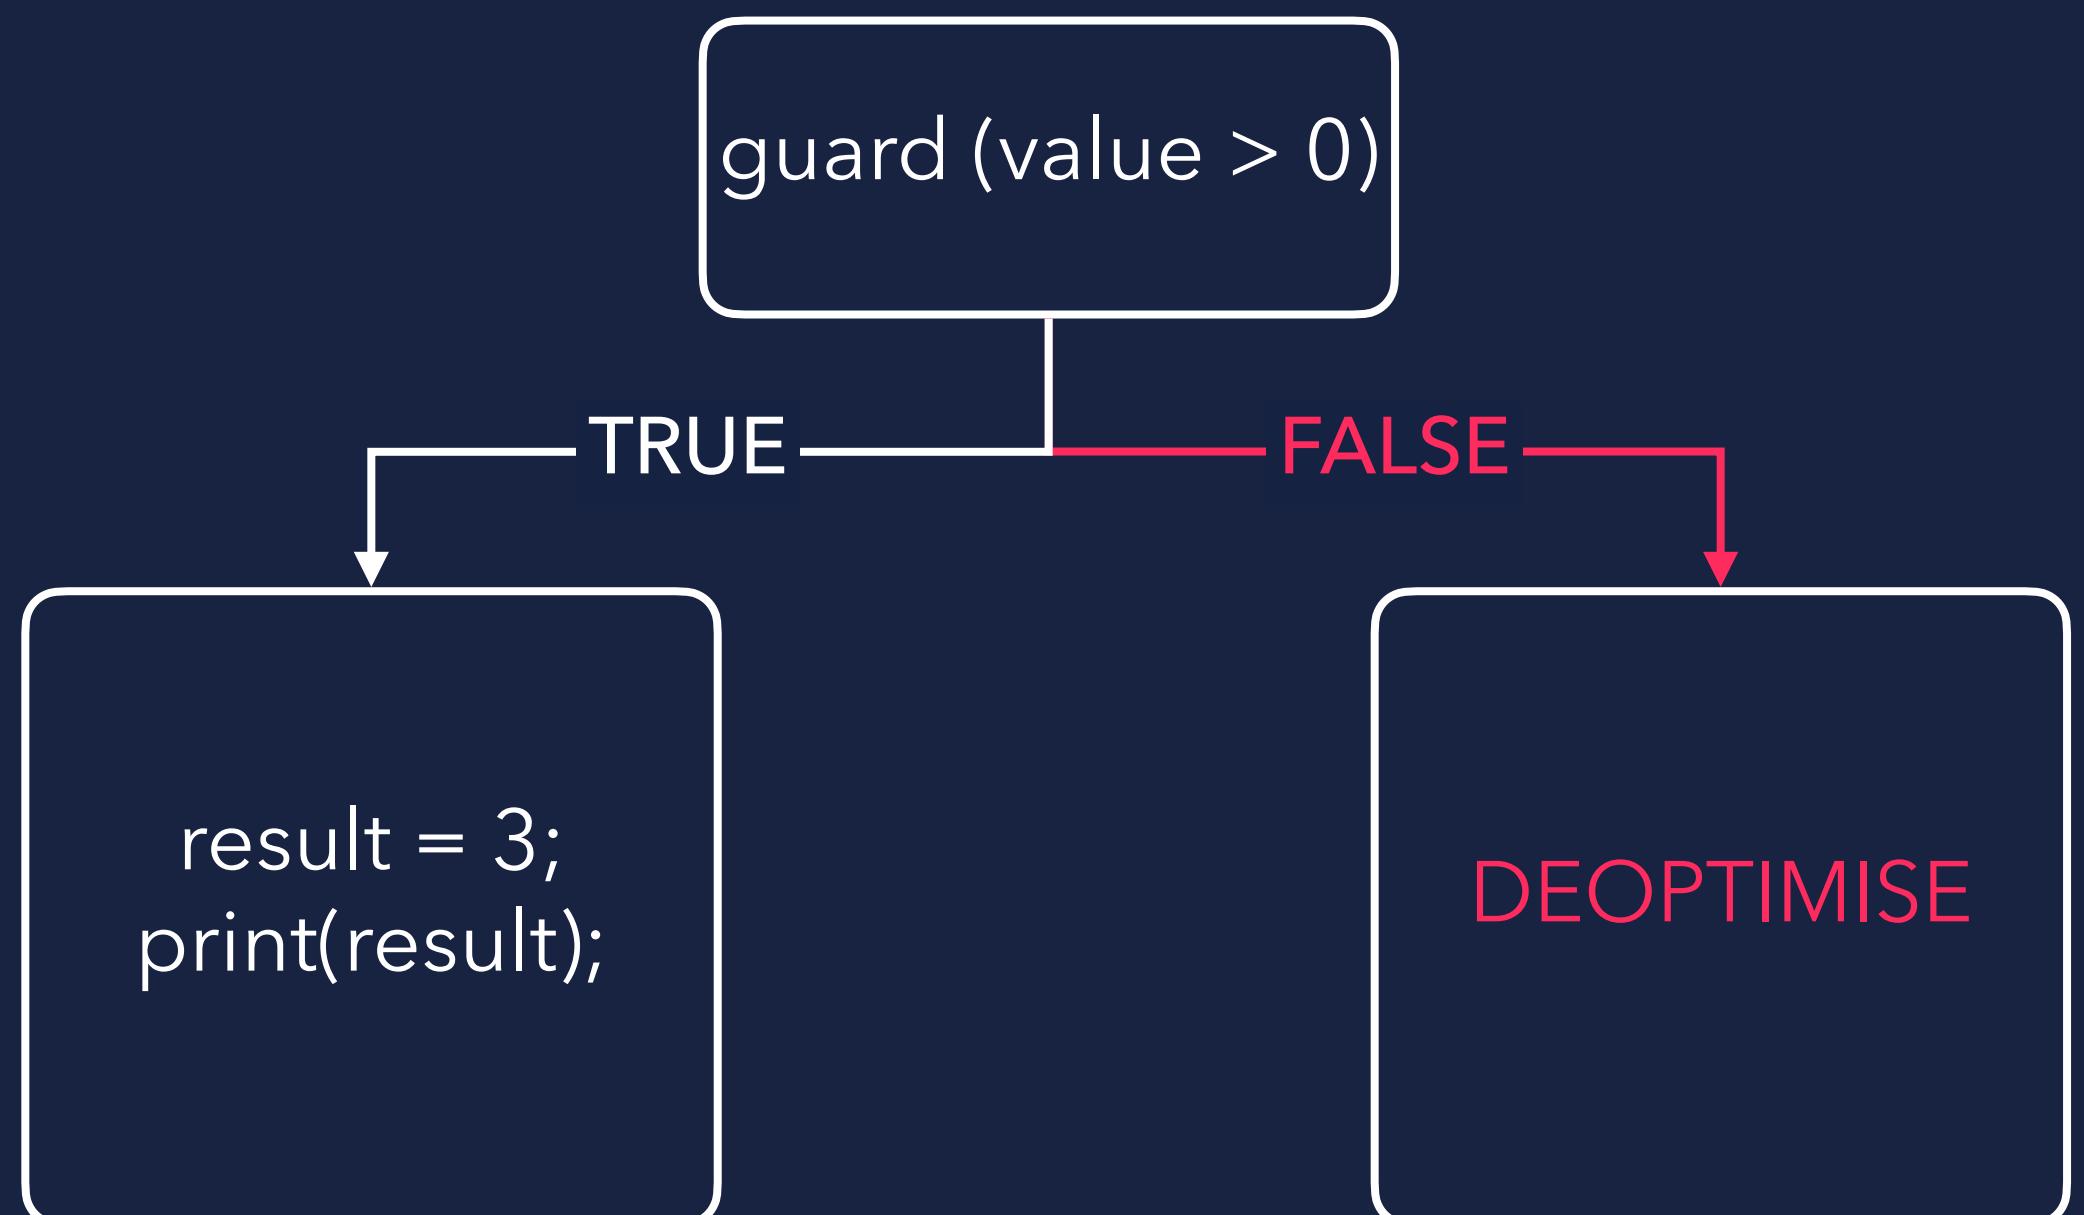
DEOPTIMISATION

CODE



OPTIMIZED CODE

value = 0



THAT'S

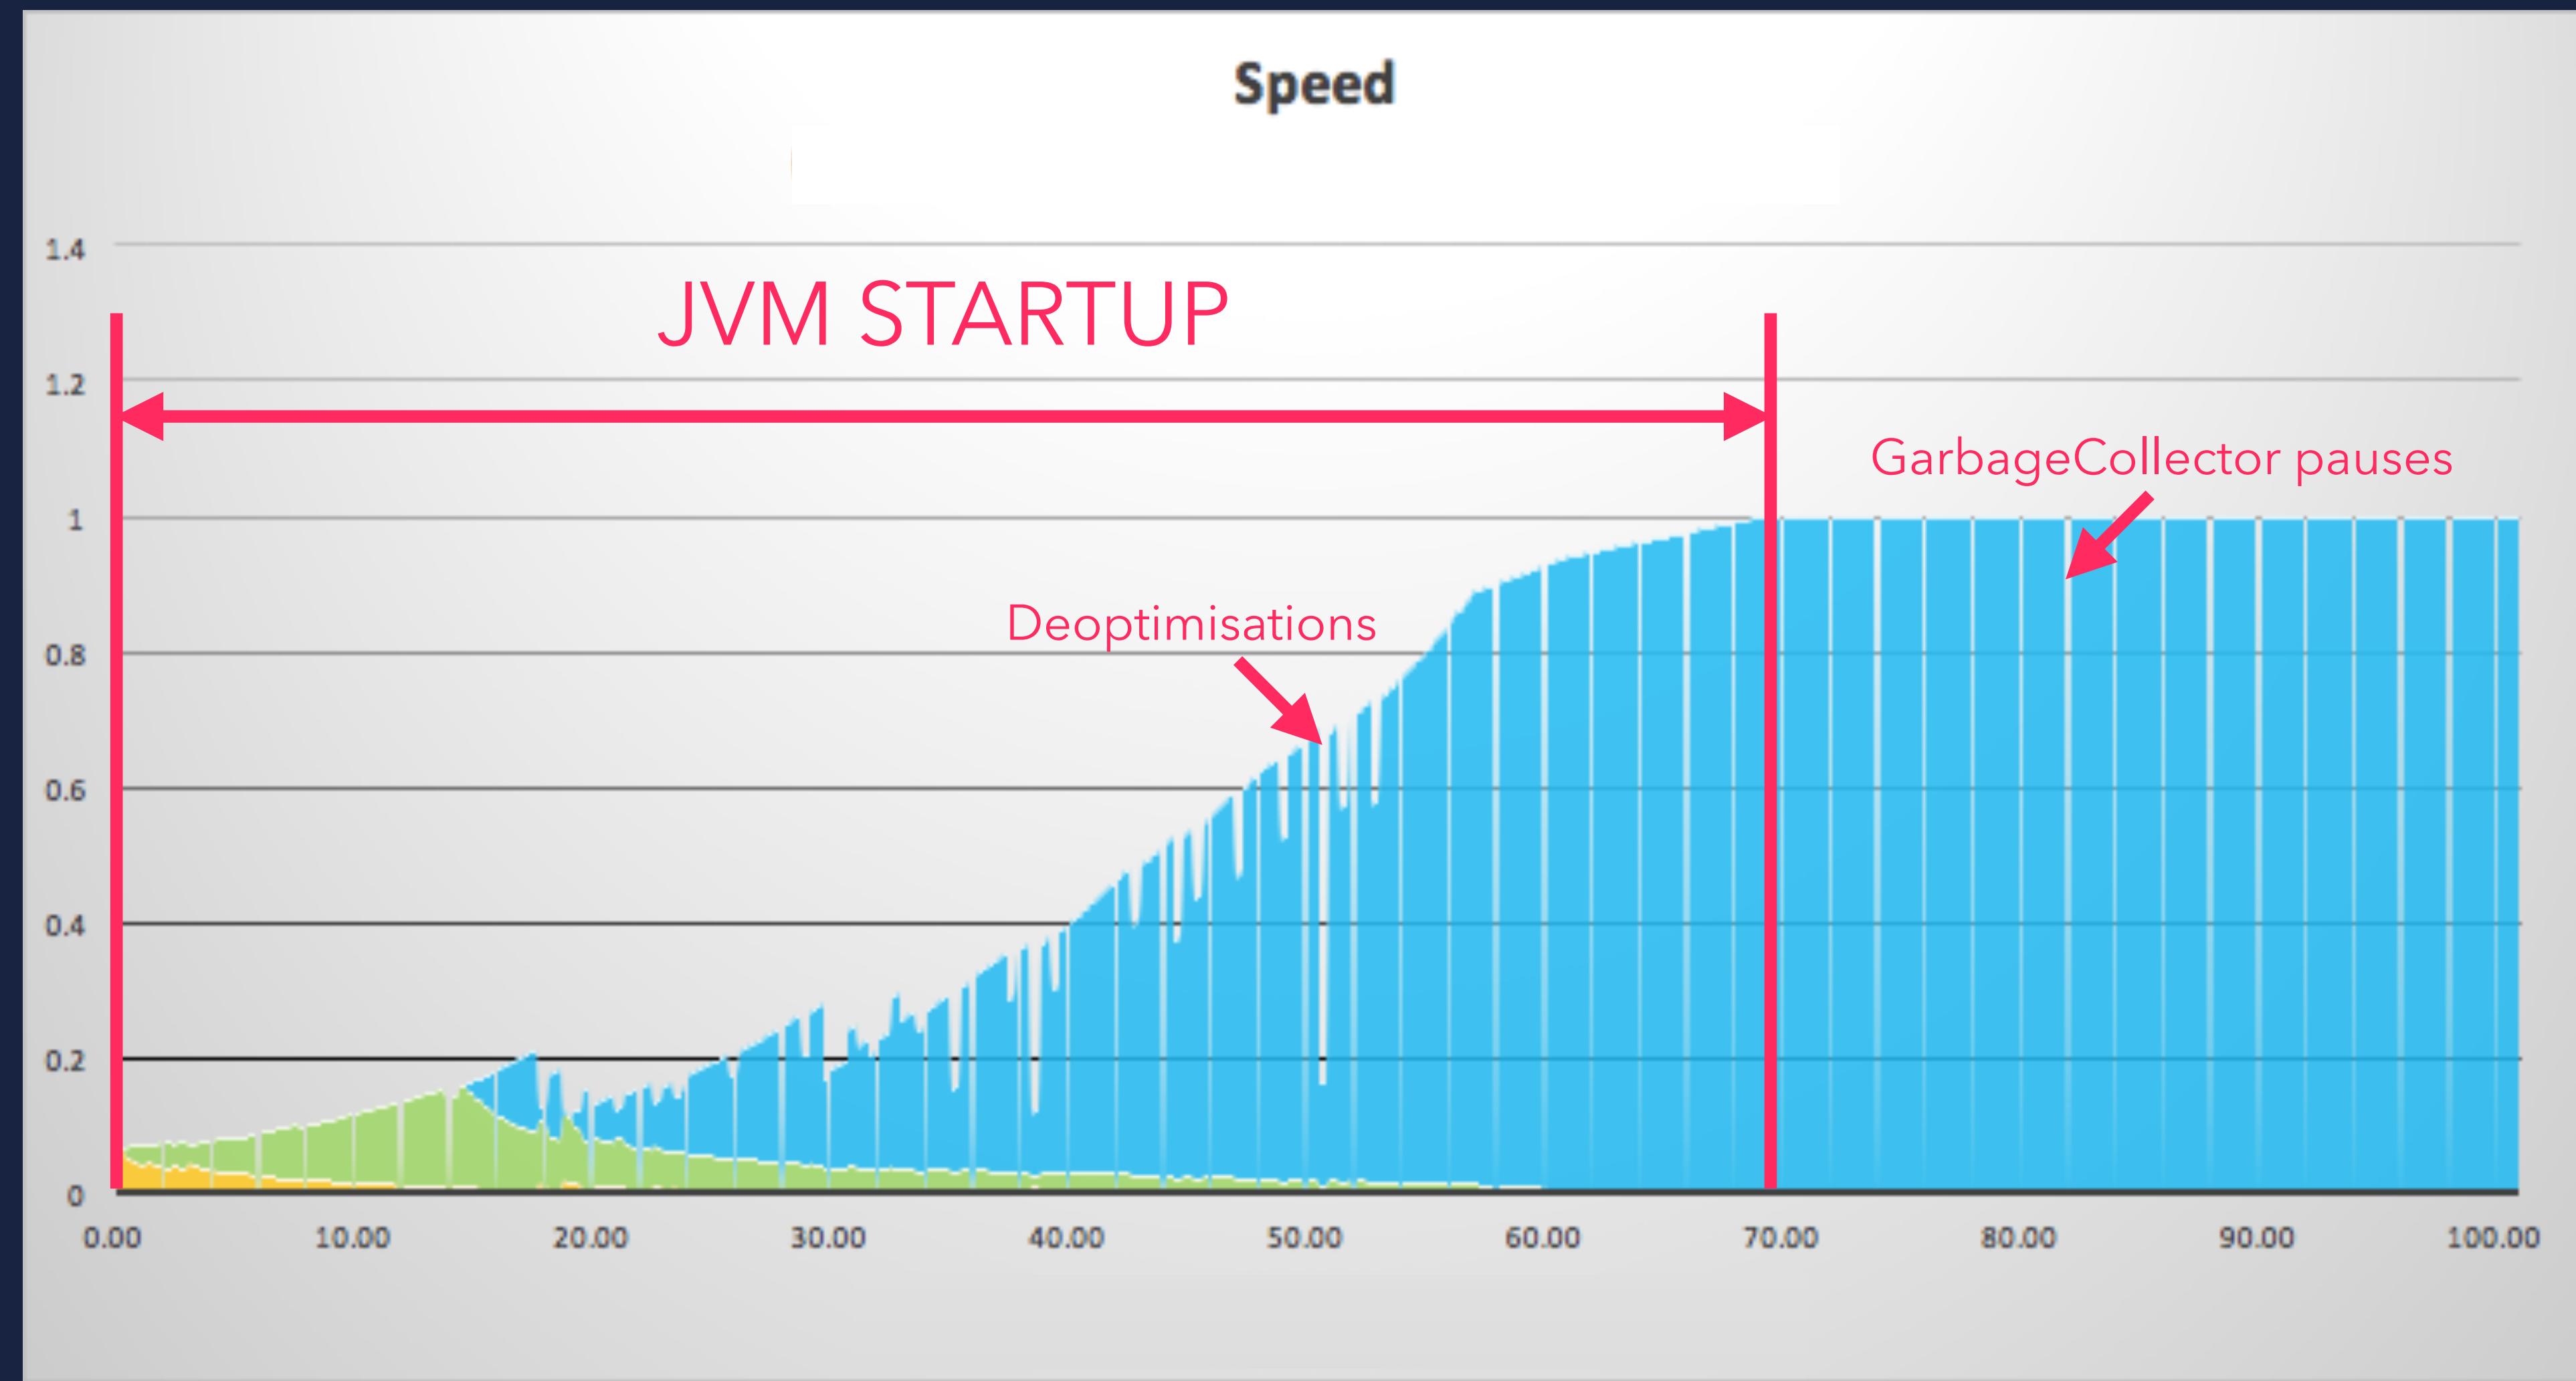
GREAT



...BUT

IT TAKES
TIME!

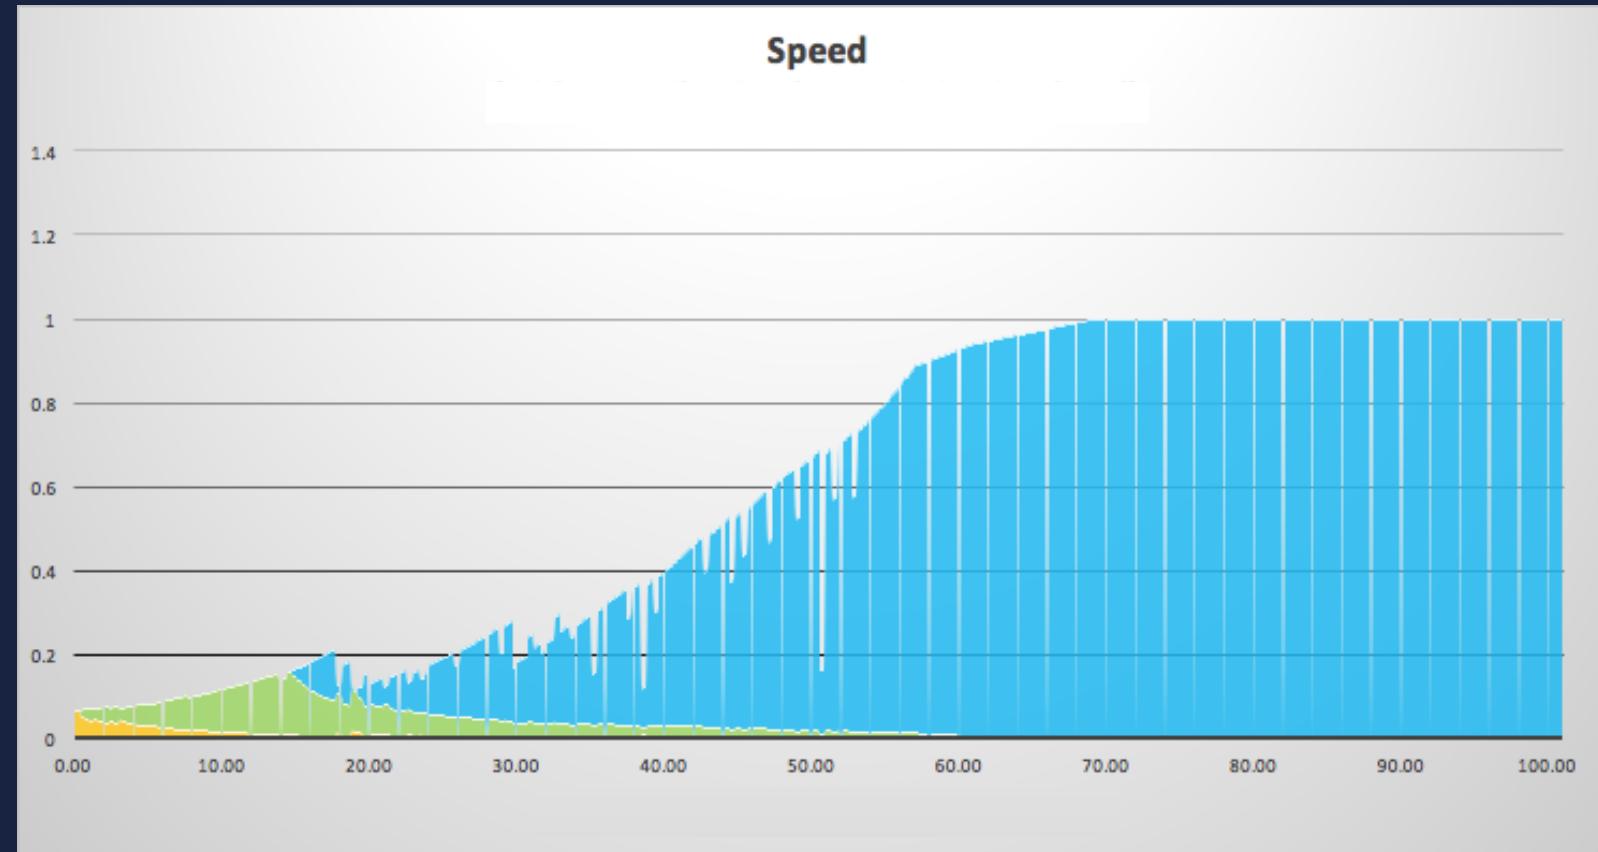
JVM PERFORMANCE GRAPH



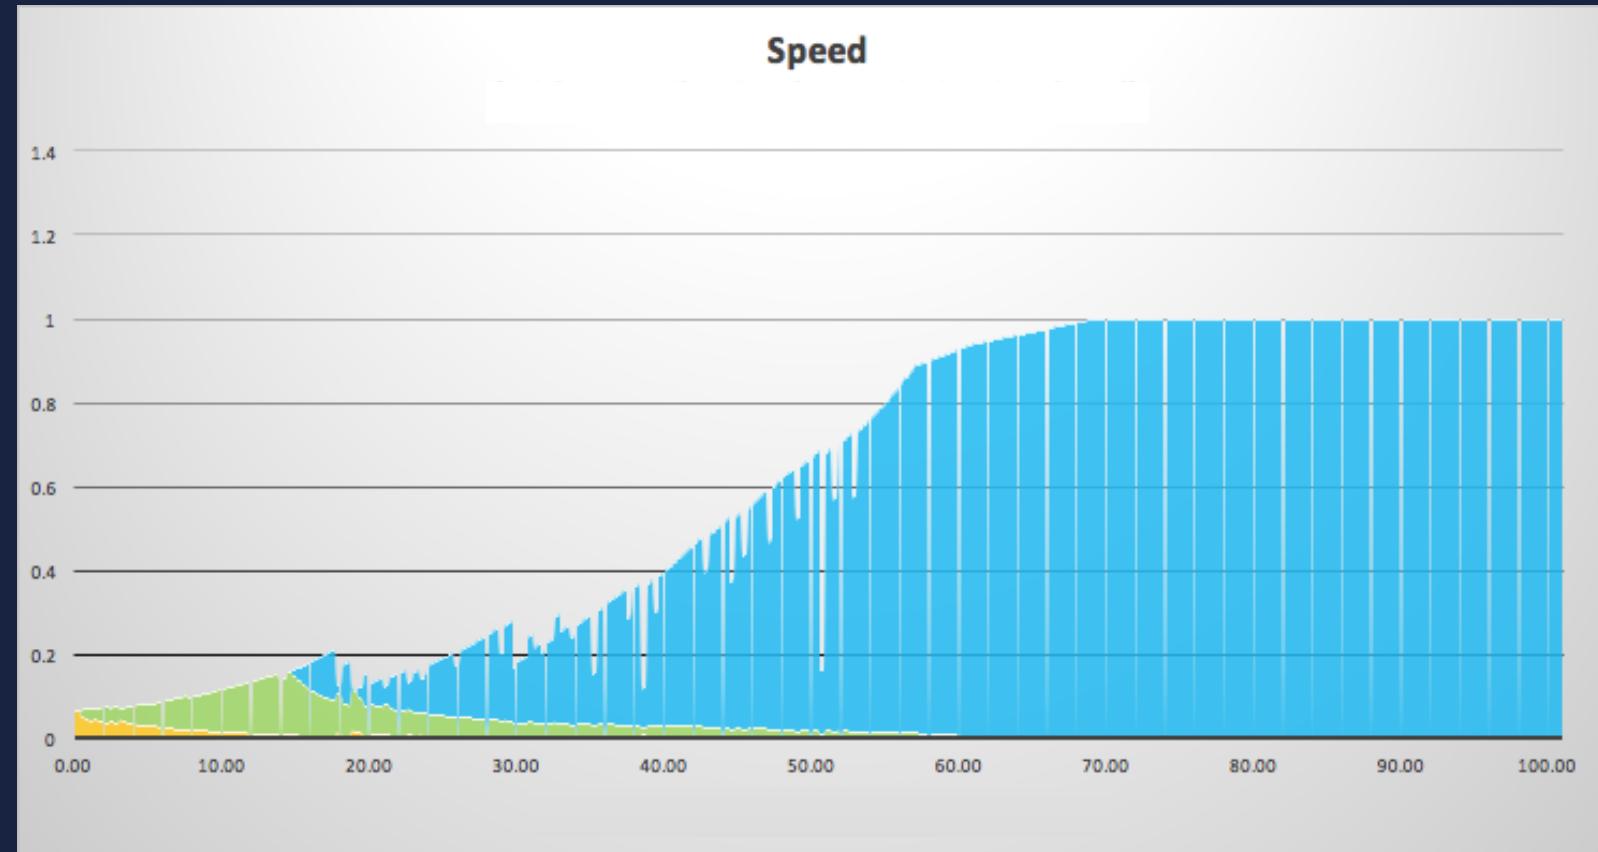
MICROSERVICE ENVIRONMENT

MICROSERVICE ENVIRONMENT

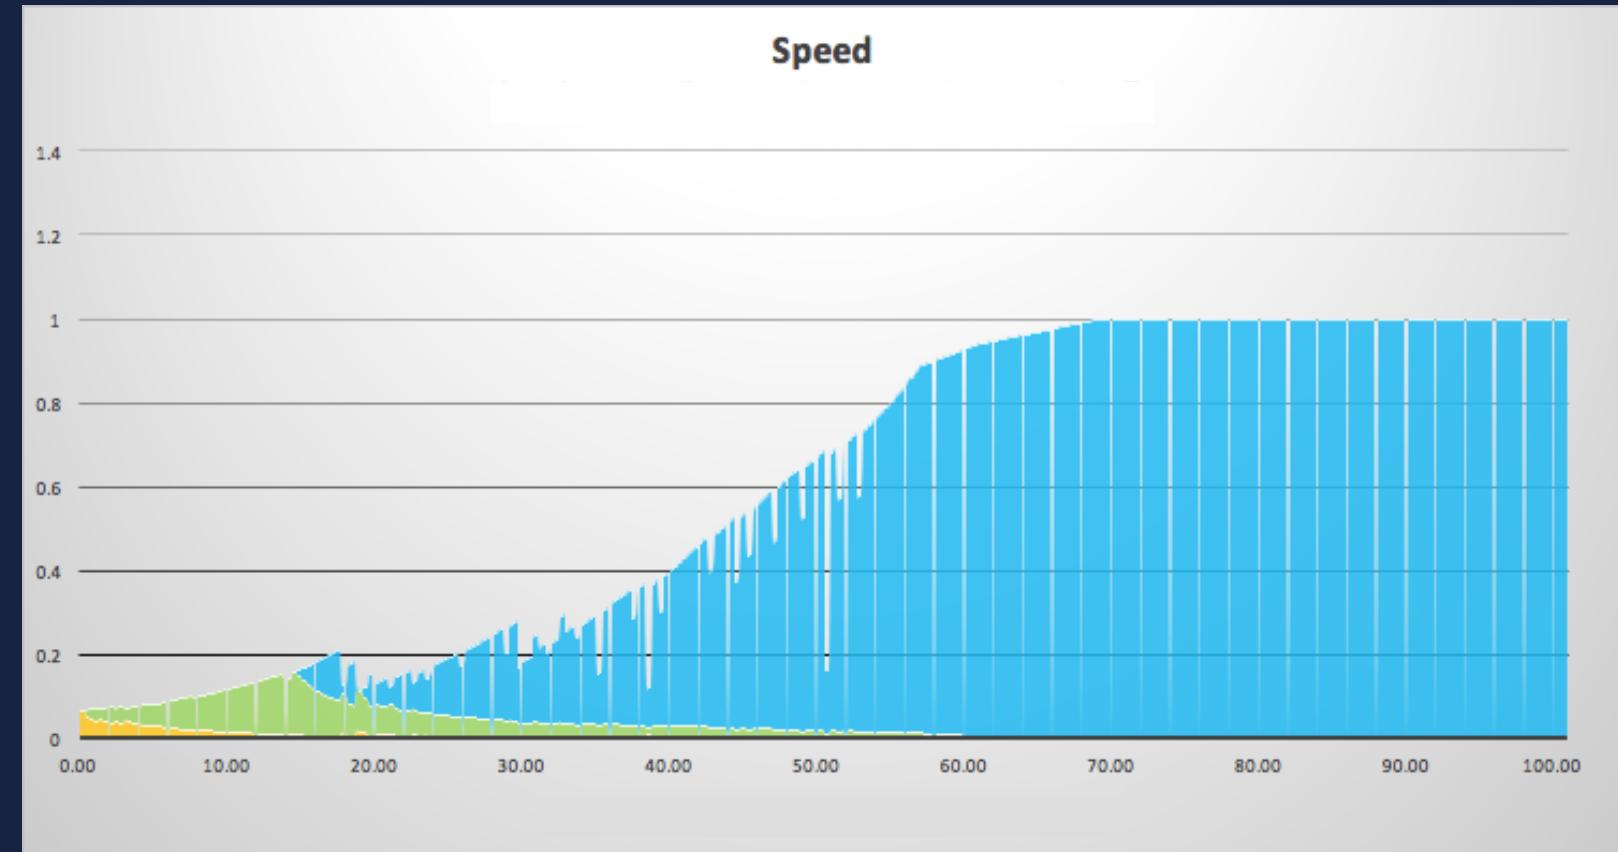
FIRST RUN



SECOND RUN



THIRD RUN



JVM STARTUP

JVM STARTUP

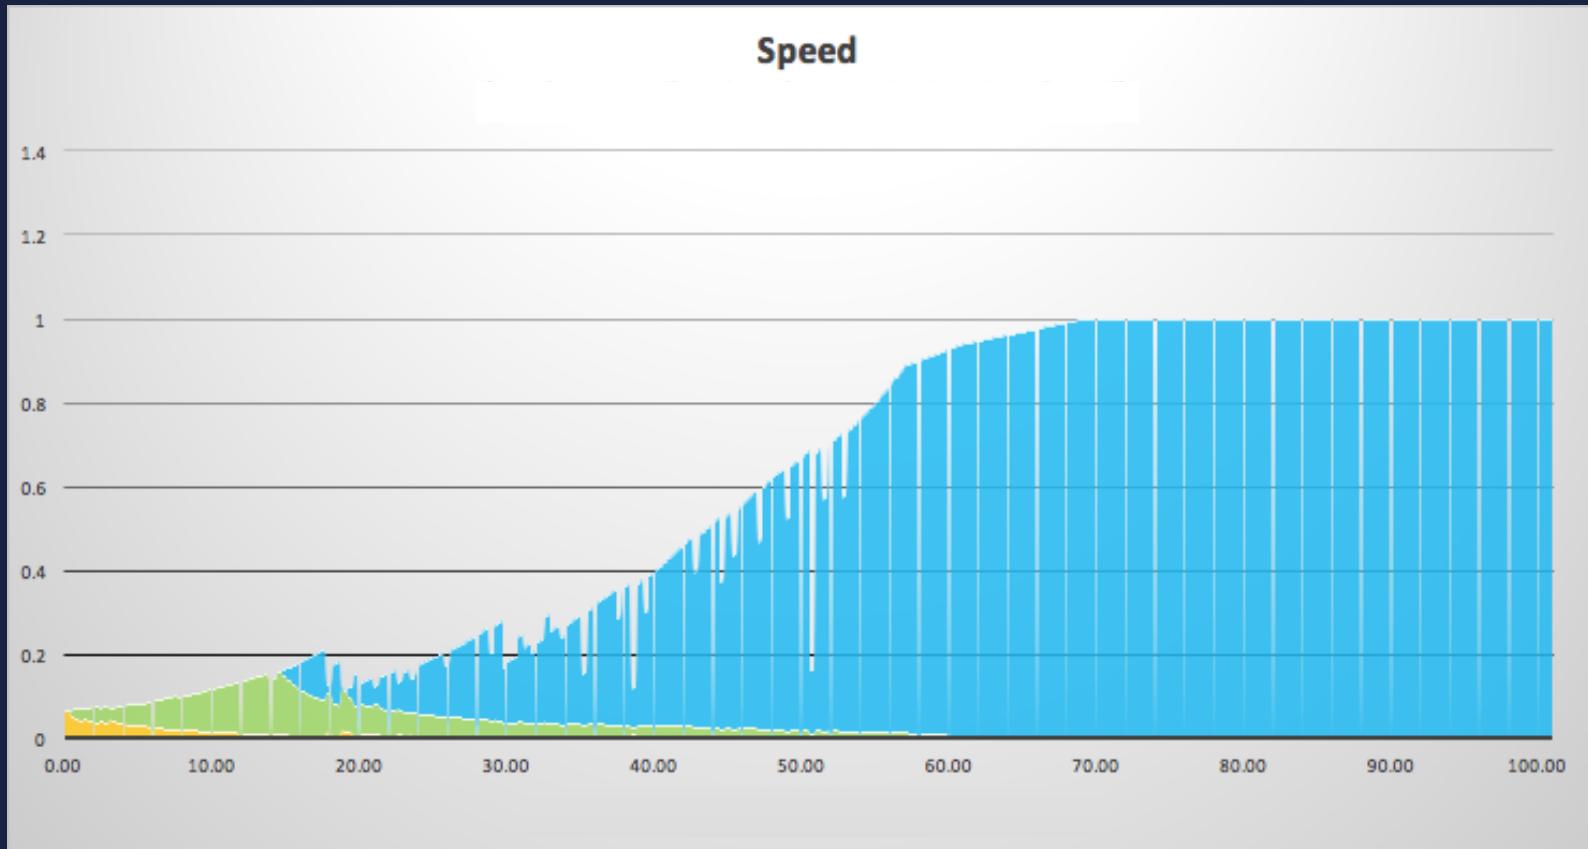
JVM STARTUP



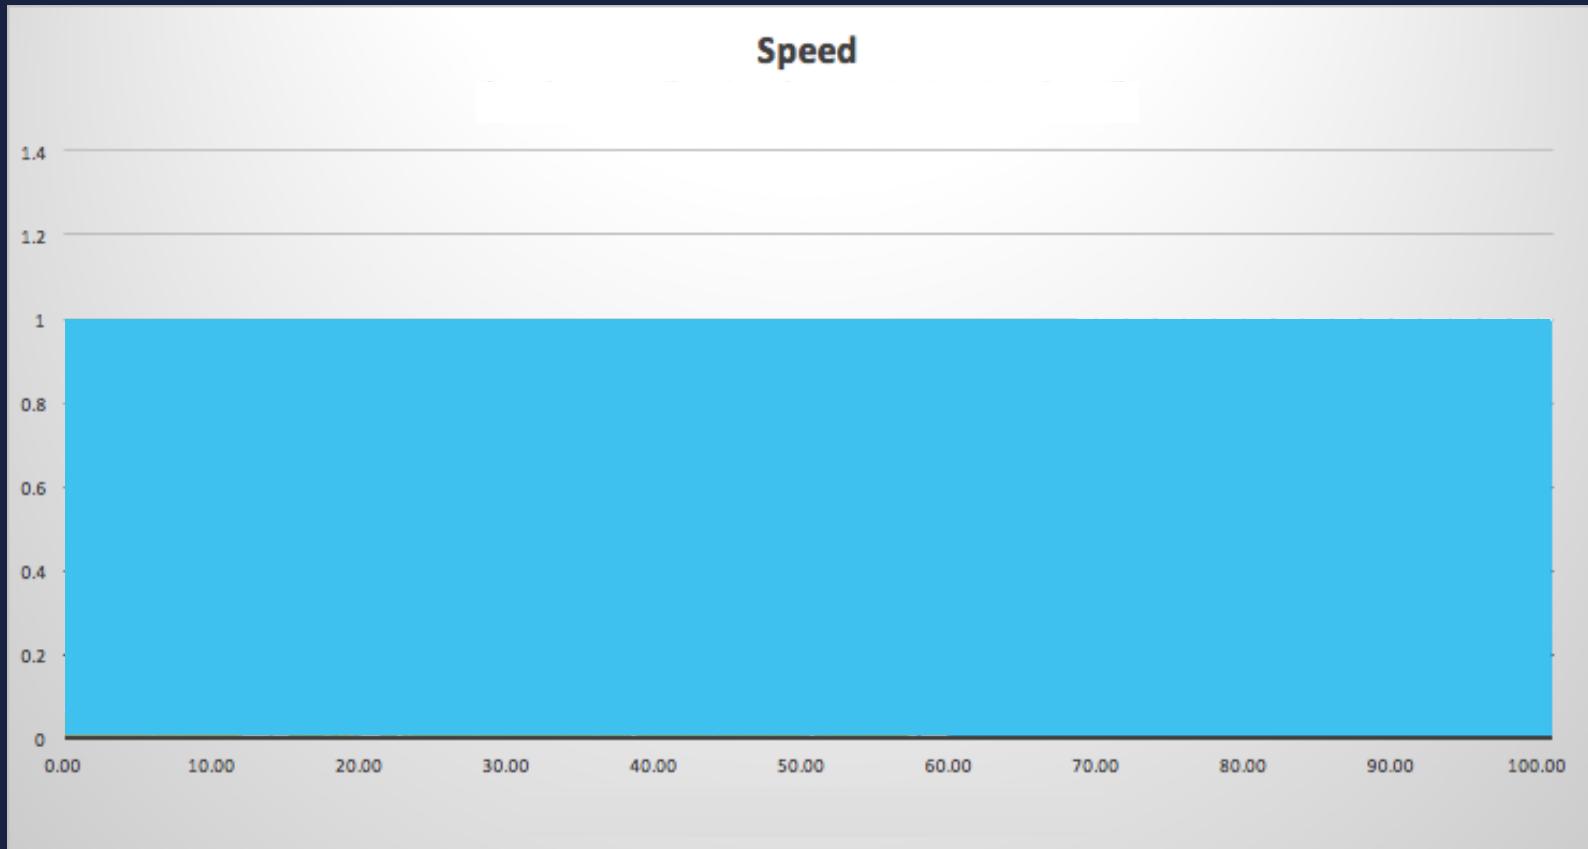
azul

WOULDN'T IT BE GREAT...?

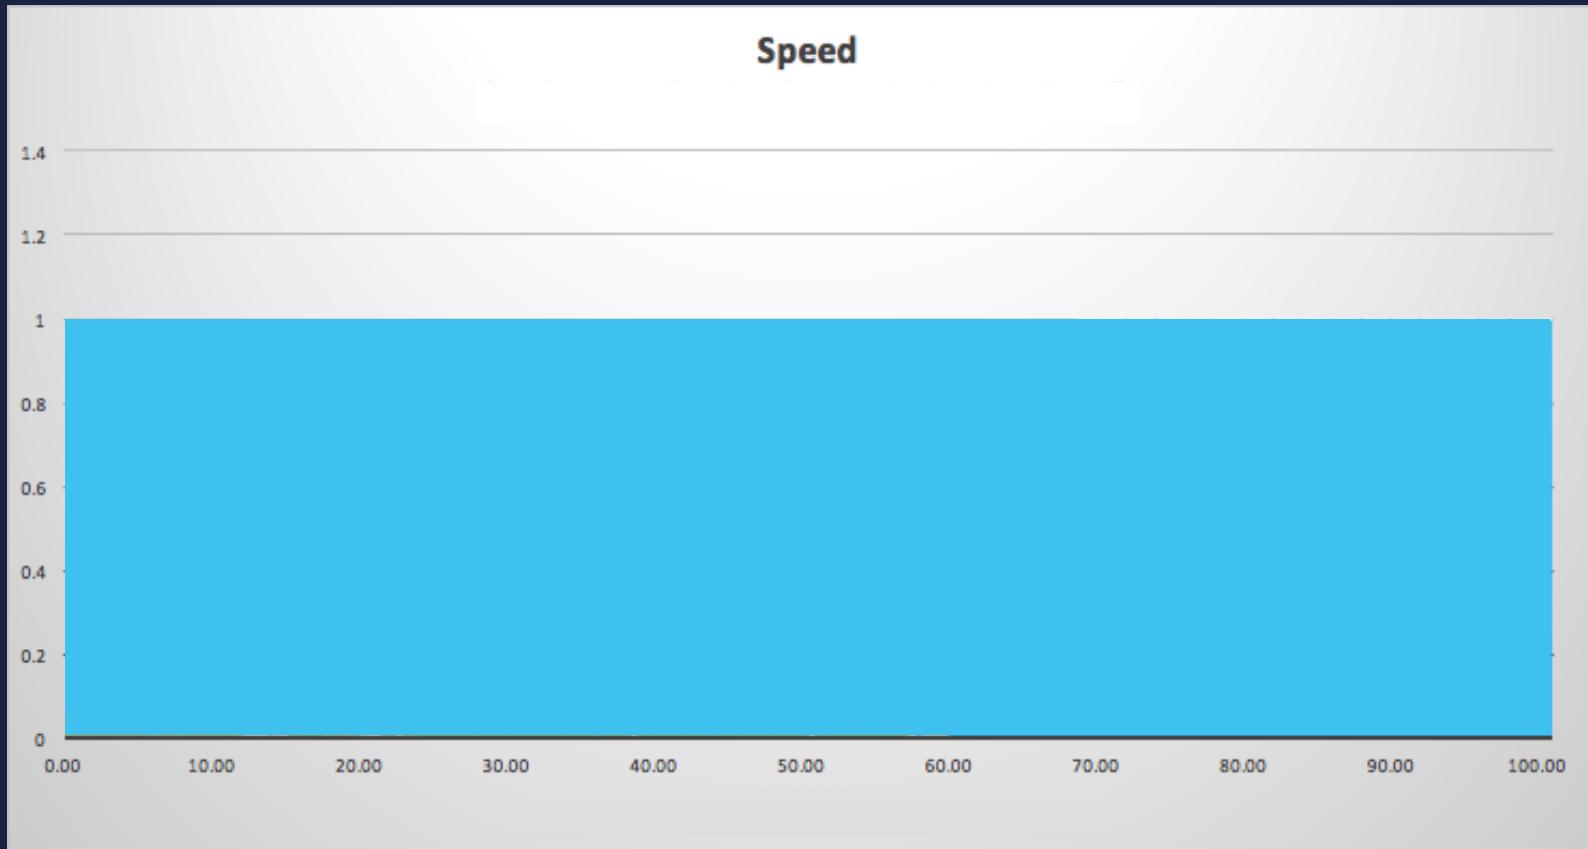
FIRST RUN



SECOND RUN



THIRD RUN



JVM STARTUP

NO STARTUP OVERHEAD

NO STARTUP OVERHEAD

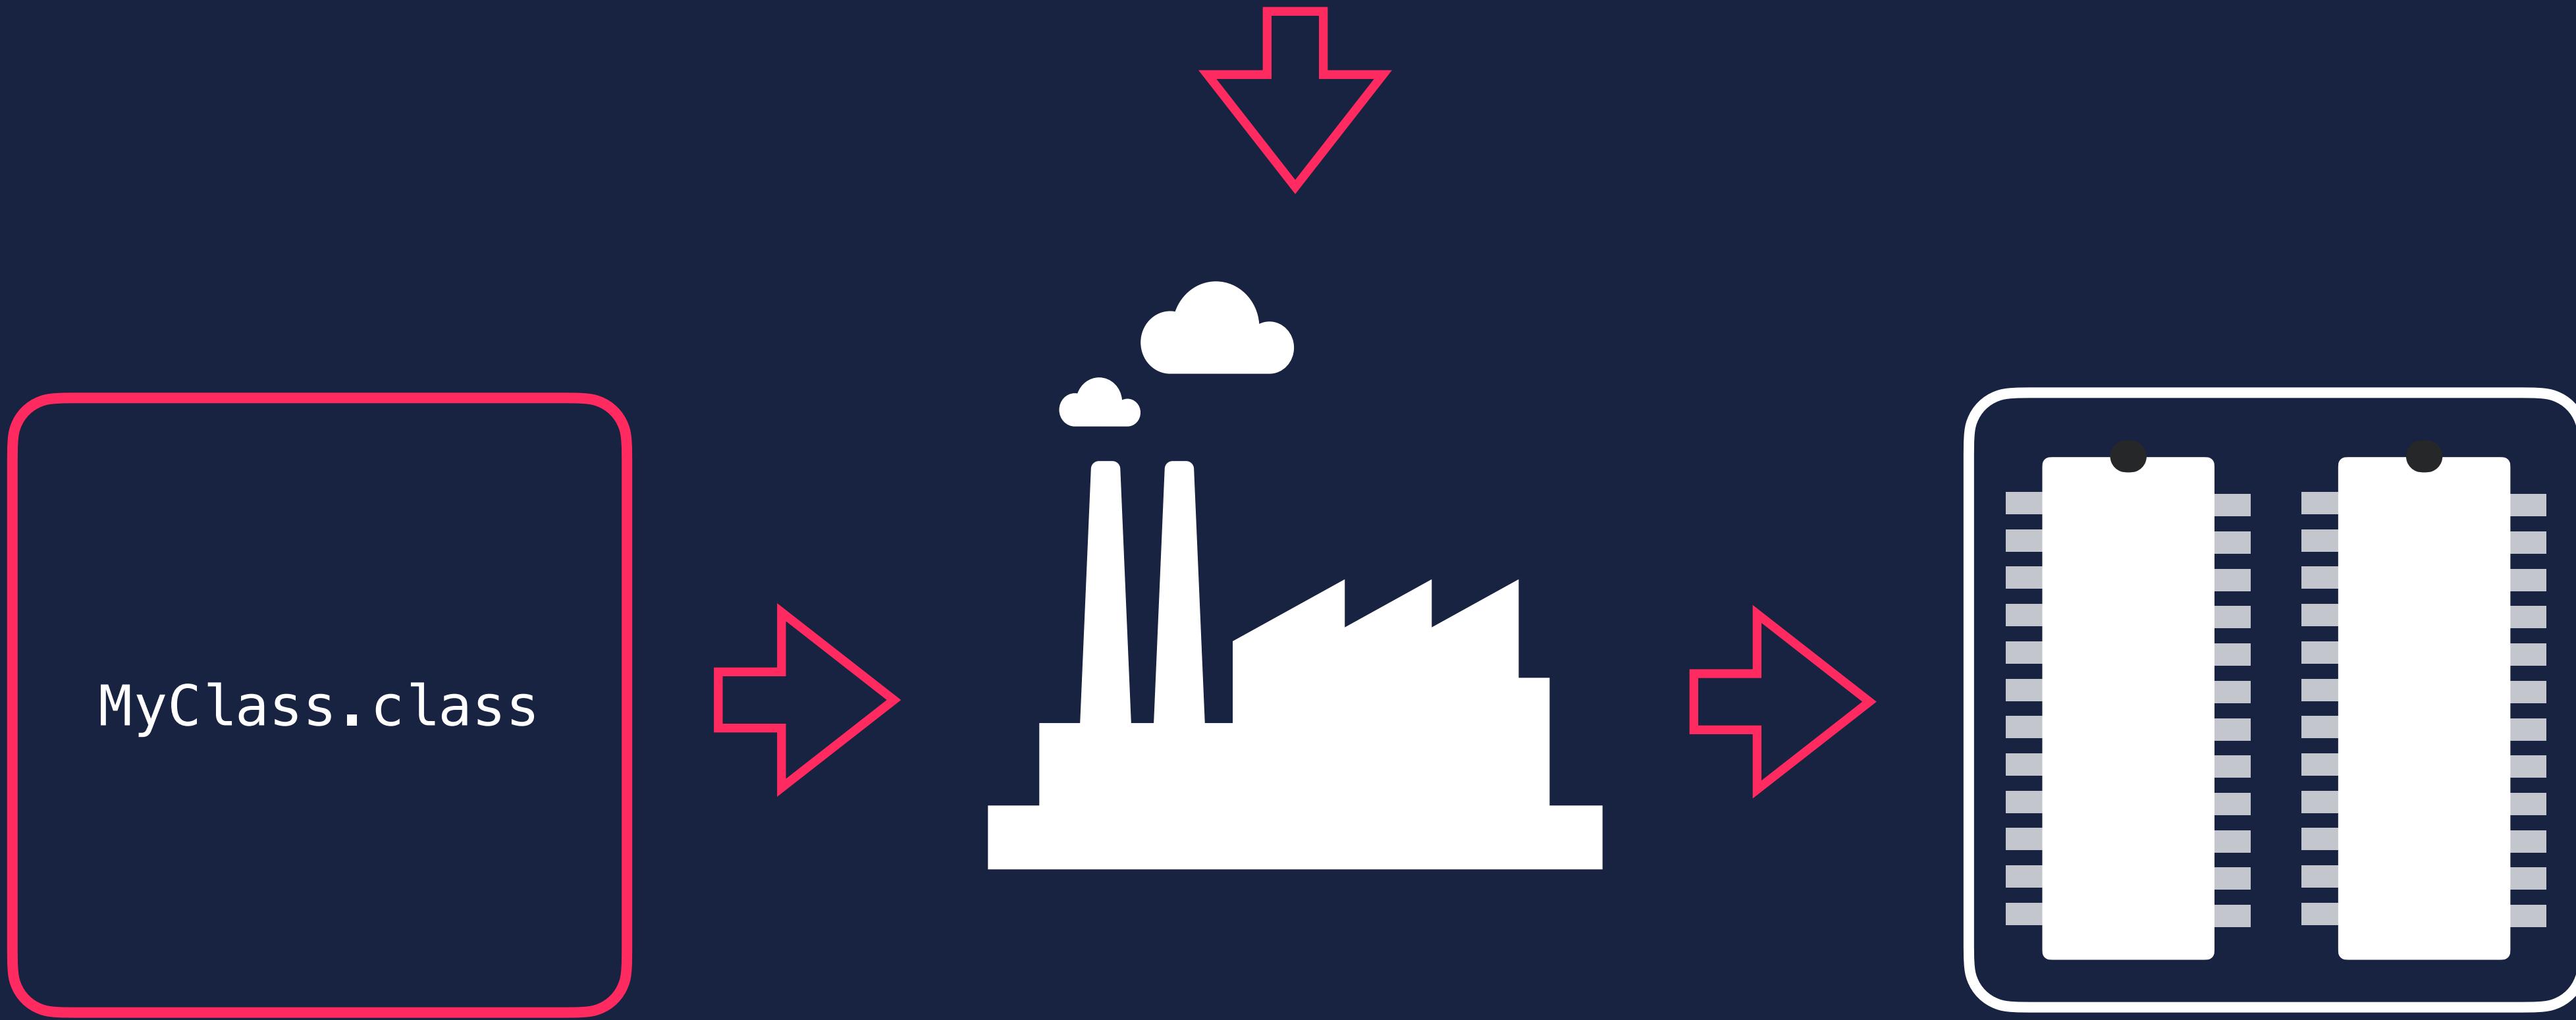
SOLUTIONS...?

CLASS DATA
SHARING

WHAT ABOUT CDS?

- Class Data Sharing (App/Dynamic)
- Dump internal class representation into file
- Shared on each JVM start (CDS)
- No optimization or hotspot detection
- Only reduces loading time of classes
- Start up could be up to 1-2 seconds faster

CDS



BYTE CODE

CLASS LOADER

JVM MEMORY

AHEAD OF TIME
COMPILATION

WHY NOT USE AOT?

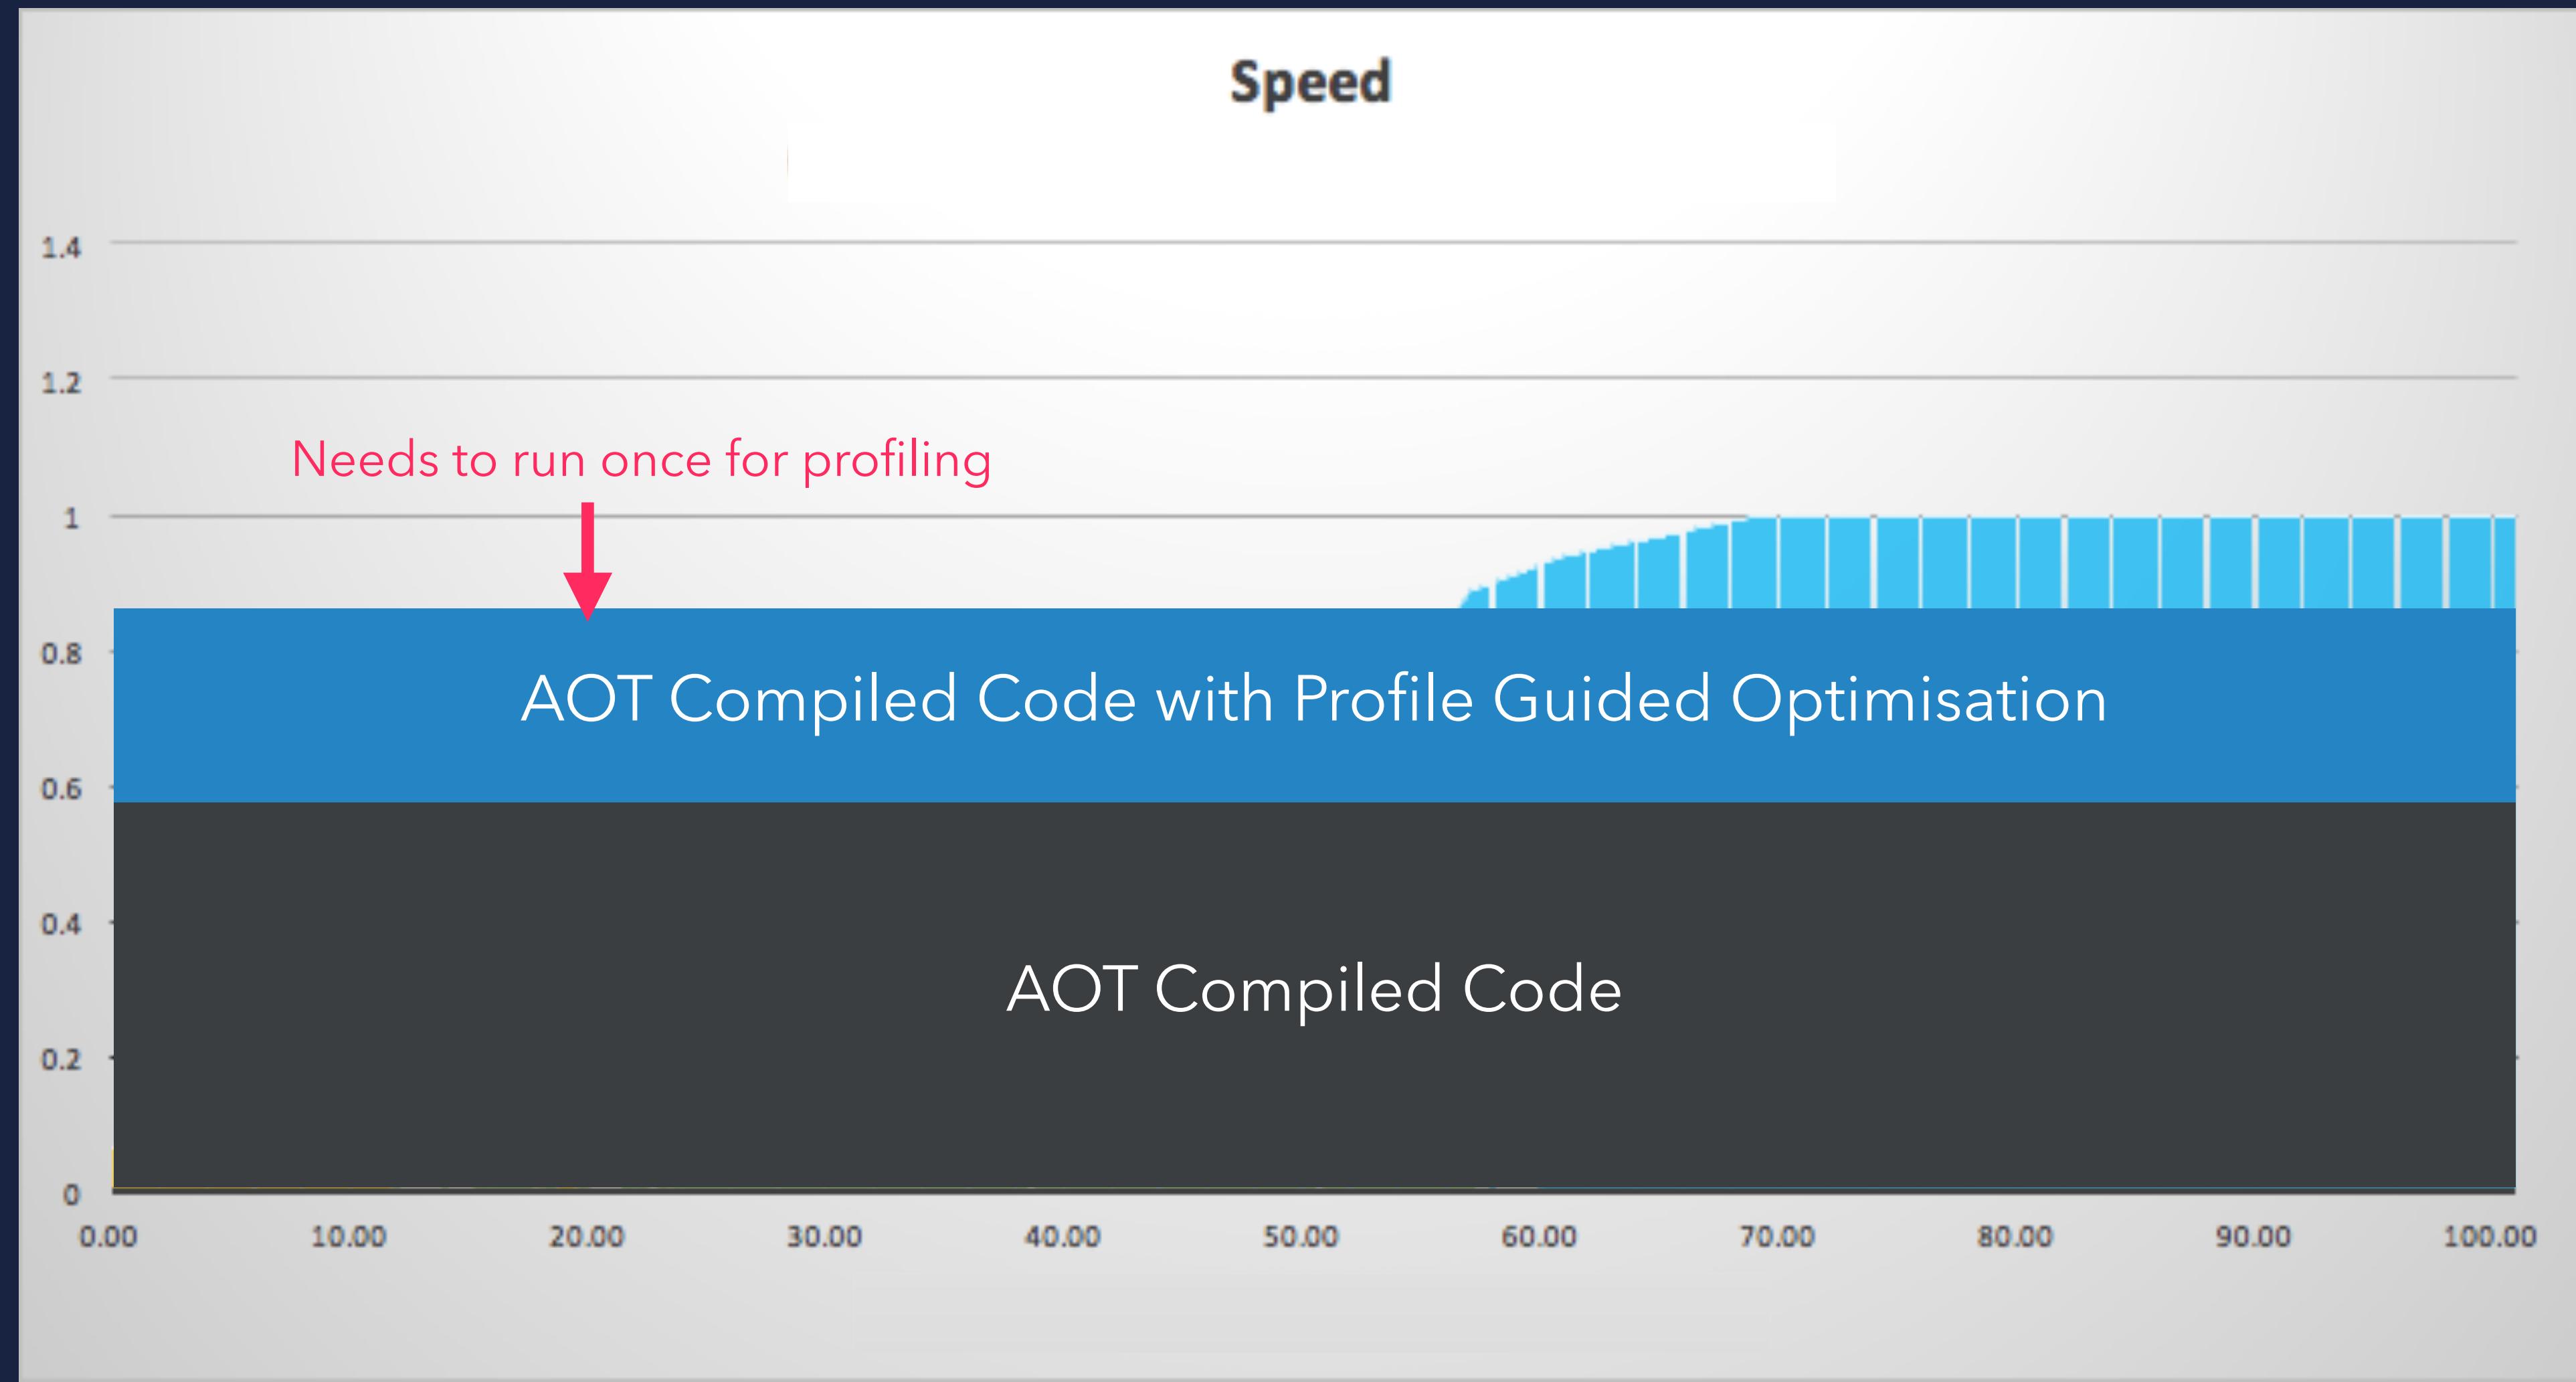
- No interpreting bytecodes
- No analysis of hotspots
- No runtime compilation of code
- Start at full speed, straight away
- GraalVM native image does that

PROBLEM SOLVED...?

NOT SO FAST...

- AOT is, by definition, static
- Code is compiled before it is run
- Compiler has no knowledge of how the code will actually run
- Profile Guided Optimisation (PGO) can partially help

JVM PERFORMANCE GRAPH



AOT VS JIT

AOT

- Limited use of method inlining
- No runtime bytecode generation
- Reflection is possible but complicated
- Unable to use speculative optimisations
- Overall performance will typically be lower
- Deployed env != Development env.

JIT

- Can use aggressive method inlining at runtime
- Can use runtime bytecode generation
- Reflection is simple
- Can use speculative optimisations
- Overall performance will typically be higher
- Deployed env. == Development env.

- Full speed from the start
- No overhead to compile code at runtime
- Small memory footprint

- Requires more time to start up
- Overhead to compile code at runtime
- Larger memory footprint

JIT DISADVANTAGES

- Requires more time to start up
- CPU overhead to compile code at runtime
- Larger memory footprint

AOT

JIT

Low Memory
Footprint

Startup Speed

Peak Throughput

Reduced
Max Latency

Small Packaging

A DIFFERENT
APPROACH



CRIU

CHECKPOINT RESTORE IN USERSPACE

CHECKPOINT | RESTORE IN USERSPACE



- Linux project
- Part of kernel >= 3.11 (2013)
- Freeze a running container/application
- Checkpoint its state to disk
- Restore the container/application from the saved data.
- Used by/integrated in OpenVZ, LXC/LXD, Docker, Podman and others

CHECKPOINT | RESTORE IN USERSPACE



- Mainly implemented in userspace (not kernel)
- Main goal is migration of containers for microservices
- Heavily relies on `/proc` file system
- It can checkpoint:
 - Processes and threads
 - Application memory, memory mapped files and shared memory
 - Open files, pipes and FIFOs
 - Sockets
 - Interprocess communication channels
 - Timers and signals
- Can rebuild TCP connection from one side without need to exchange packets

CRIU CHALLENGES



- Restart from saved state on another machine
(open files, shared memory etc.)
- Start multiple instances of same state on same machine
(PID will be restored which will lead to problems)
- A Java Virtual Machine would assume it was continuing its tasks
(very difficult to use effectively, e.g. running applications might have open files etc.)

CRAAC

Coordinated Restore at Checkpoint

CRaC

- CRIU comes bundled with the JDK
- Heap is cleaned, compacted (JVM is in a safe state)
- Throws CheckpointException (in case of open files/sockets)
- Comes with a simple API
- Creates checkpoints using code or jcmd
- Uses JVM safepoint mechanism (to stop threads etc.)

openjdk.org/projects/crac

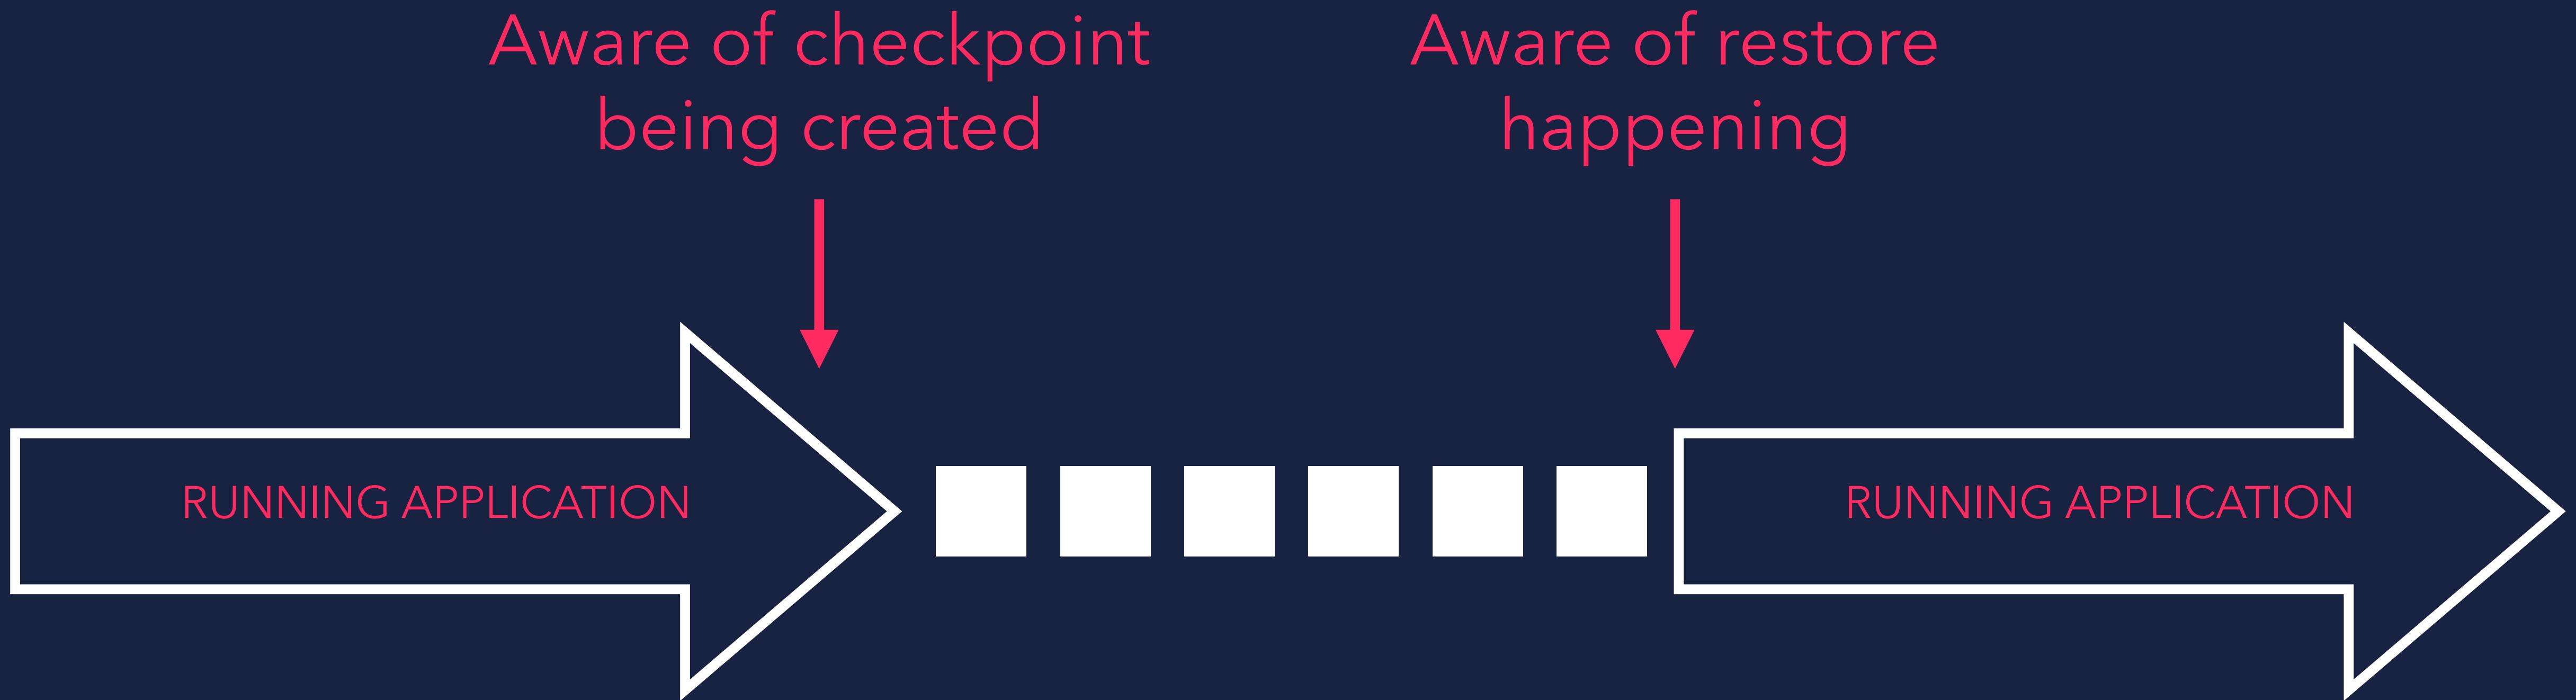
Lead by Anton Kozlov (Azul)

azul



CRaC

CRaC has more restrictions (e.g. no open files, sockets etc.)



CRaC

START

```
>java -XX:CRaCCheckpointTo=PATH -jar app.jar
```

RESTORE

```
>java -XX:CRaCRestoreFrom=PATH
```

CRAC API

CRaC API

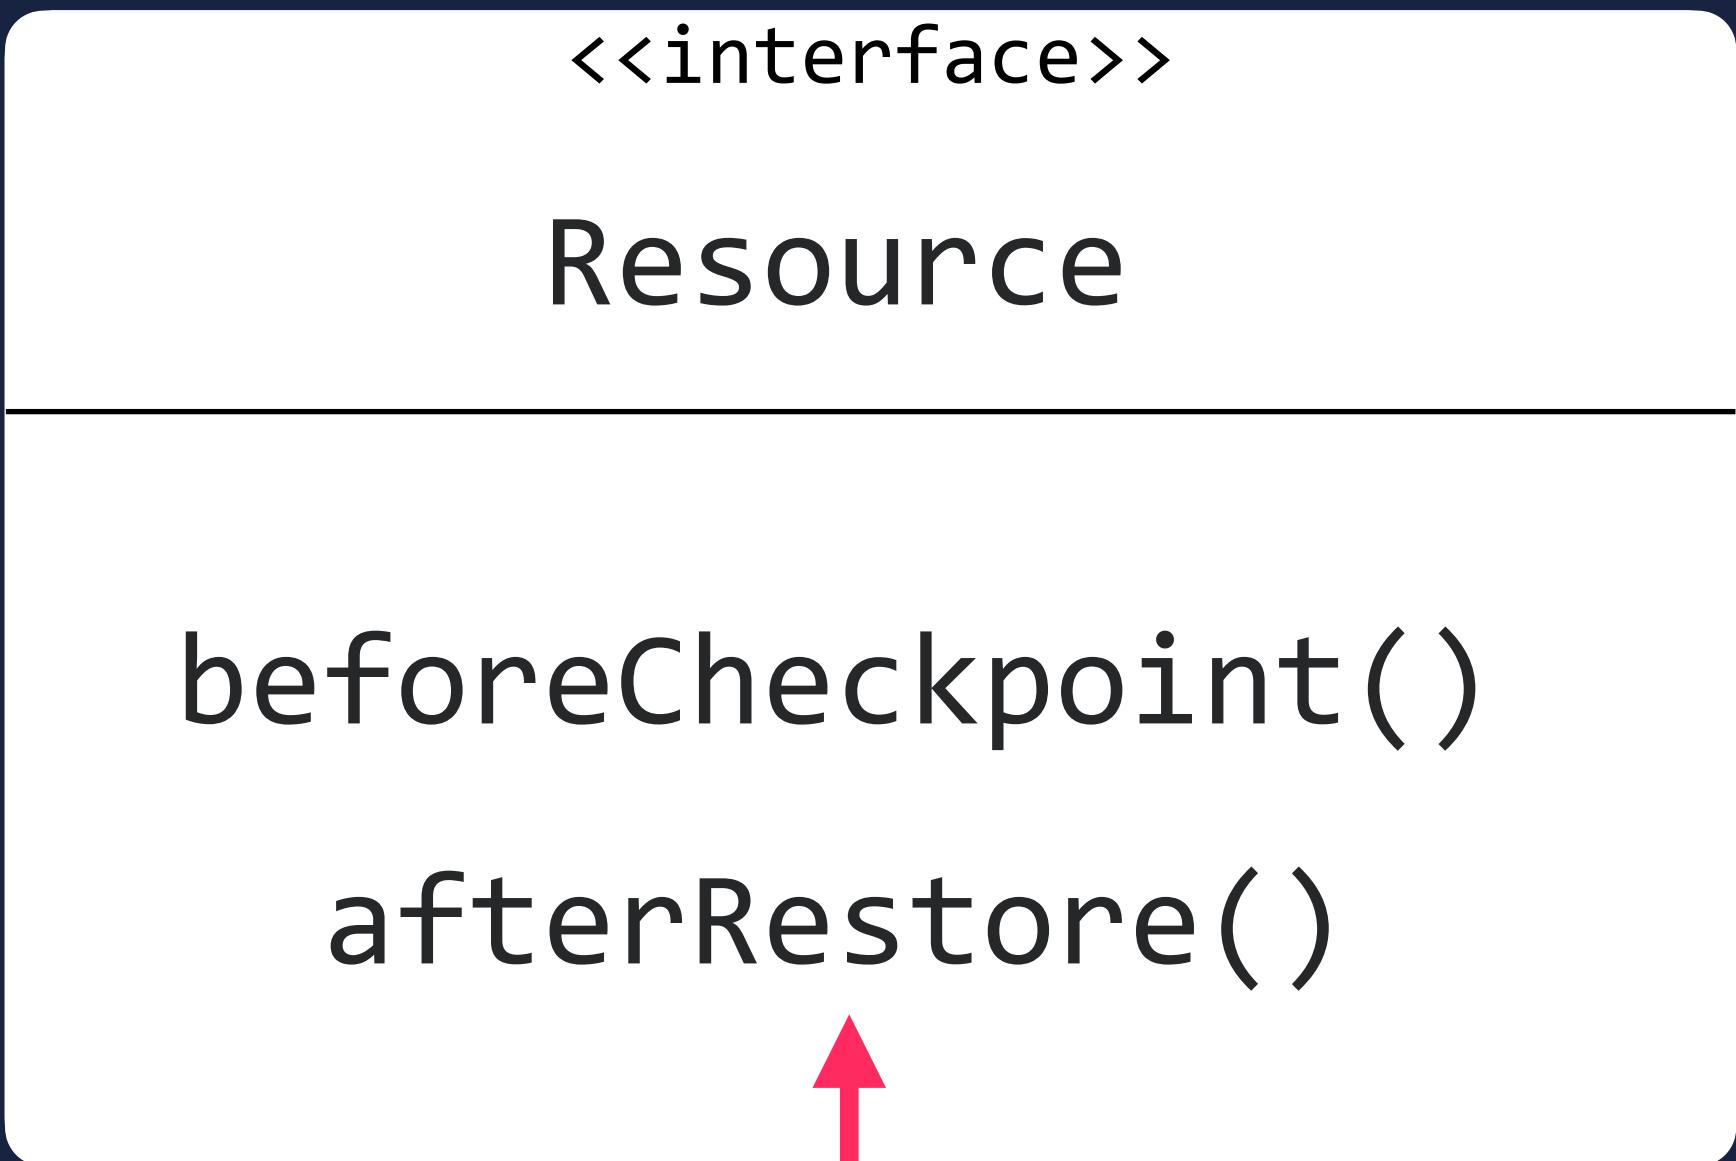
- CRaC uses Resources that can be notified about a Checkpoint and Restore
- Classes in application code implement the Resource interface
- The application receives callbacks during checkpointing and restoring
- Makes it possible to close/restore resources (e.g. open files, sockets)

```
<<interface>>
Resource
-----
beforeCheckpoint()
afterRestore()
```

CRaC API

- Resource objects need to be registered with a Context so that they can receive notifications
- There is a global Context accessible via the static `getGlobalContext()` method of the Core class

CRaC API



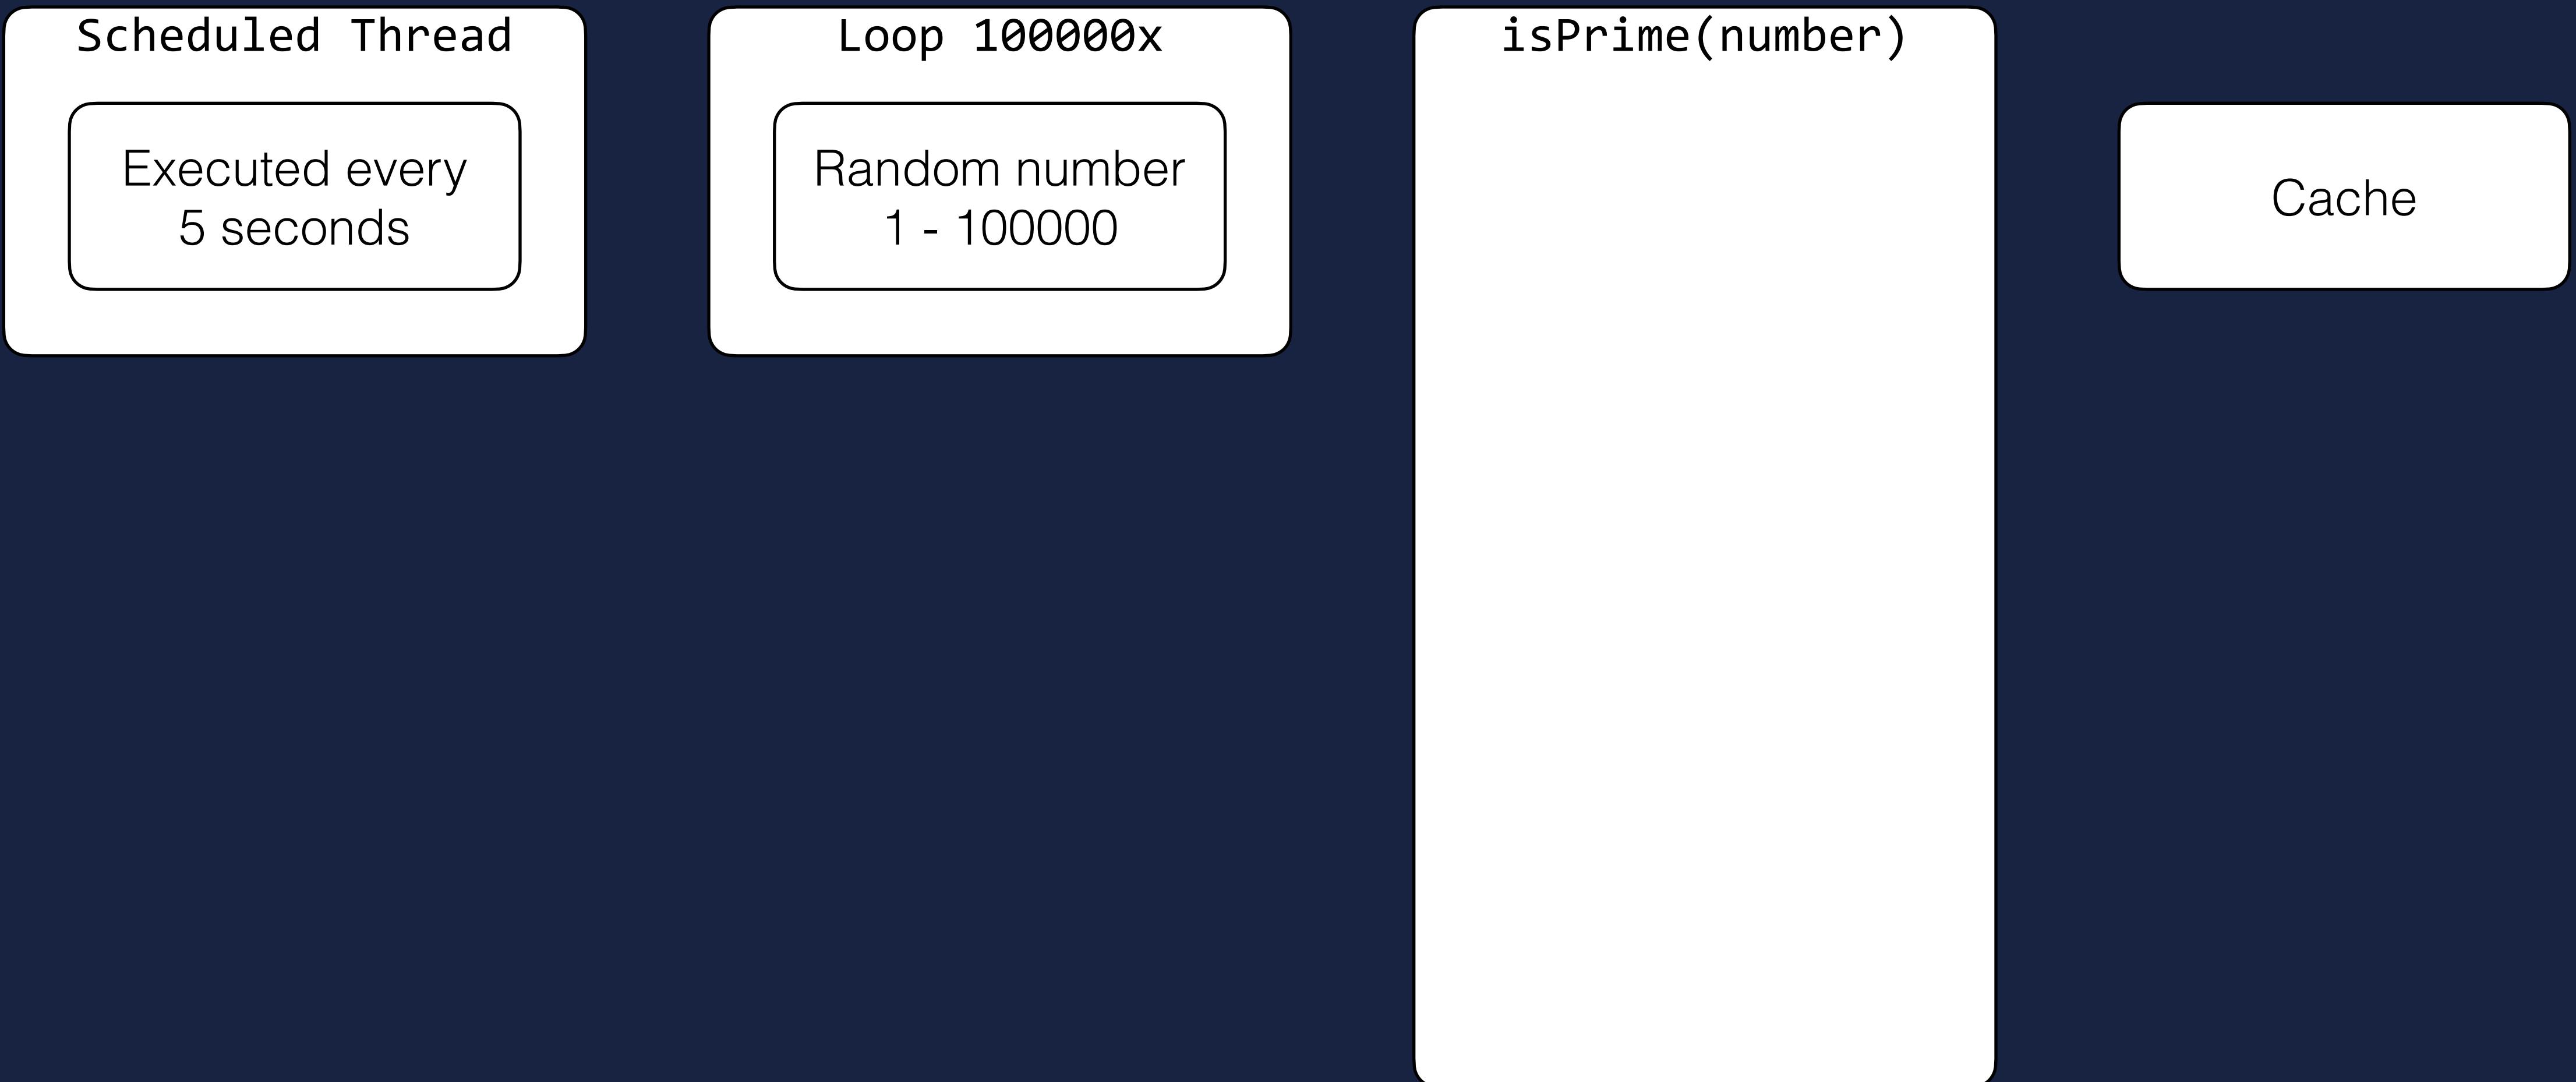
CRaCAPI

- The global `Context` maintains a list of `Resource` objects
- The `beforeCheckpoint()` methods are called in the reverse order the `Resource` objects were added
- The `afterRestore()` methods are called in the order the `Resource` objects were added
- Checkpoint and restore can be done programmatically by using `Core.checkpointRestore()`
- Checkpoint can also be created by `jcmd`

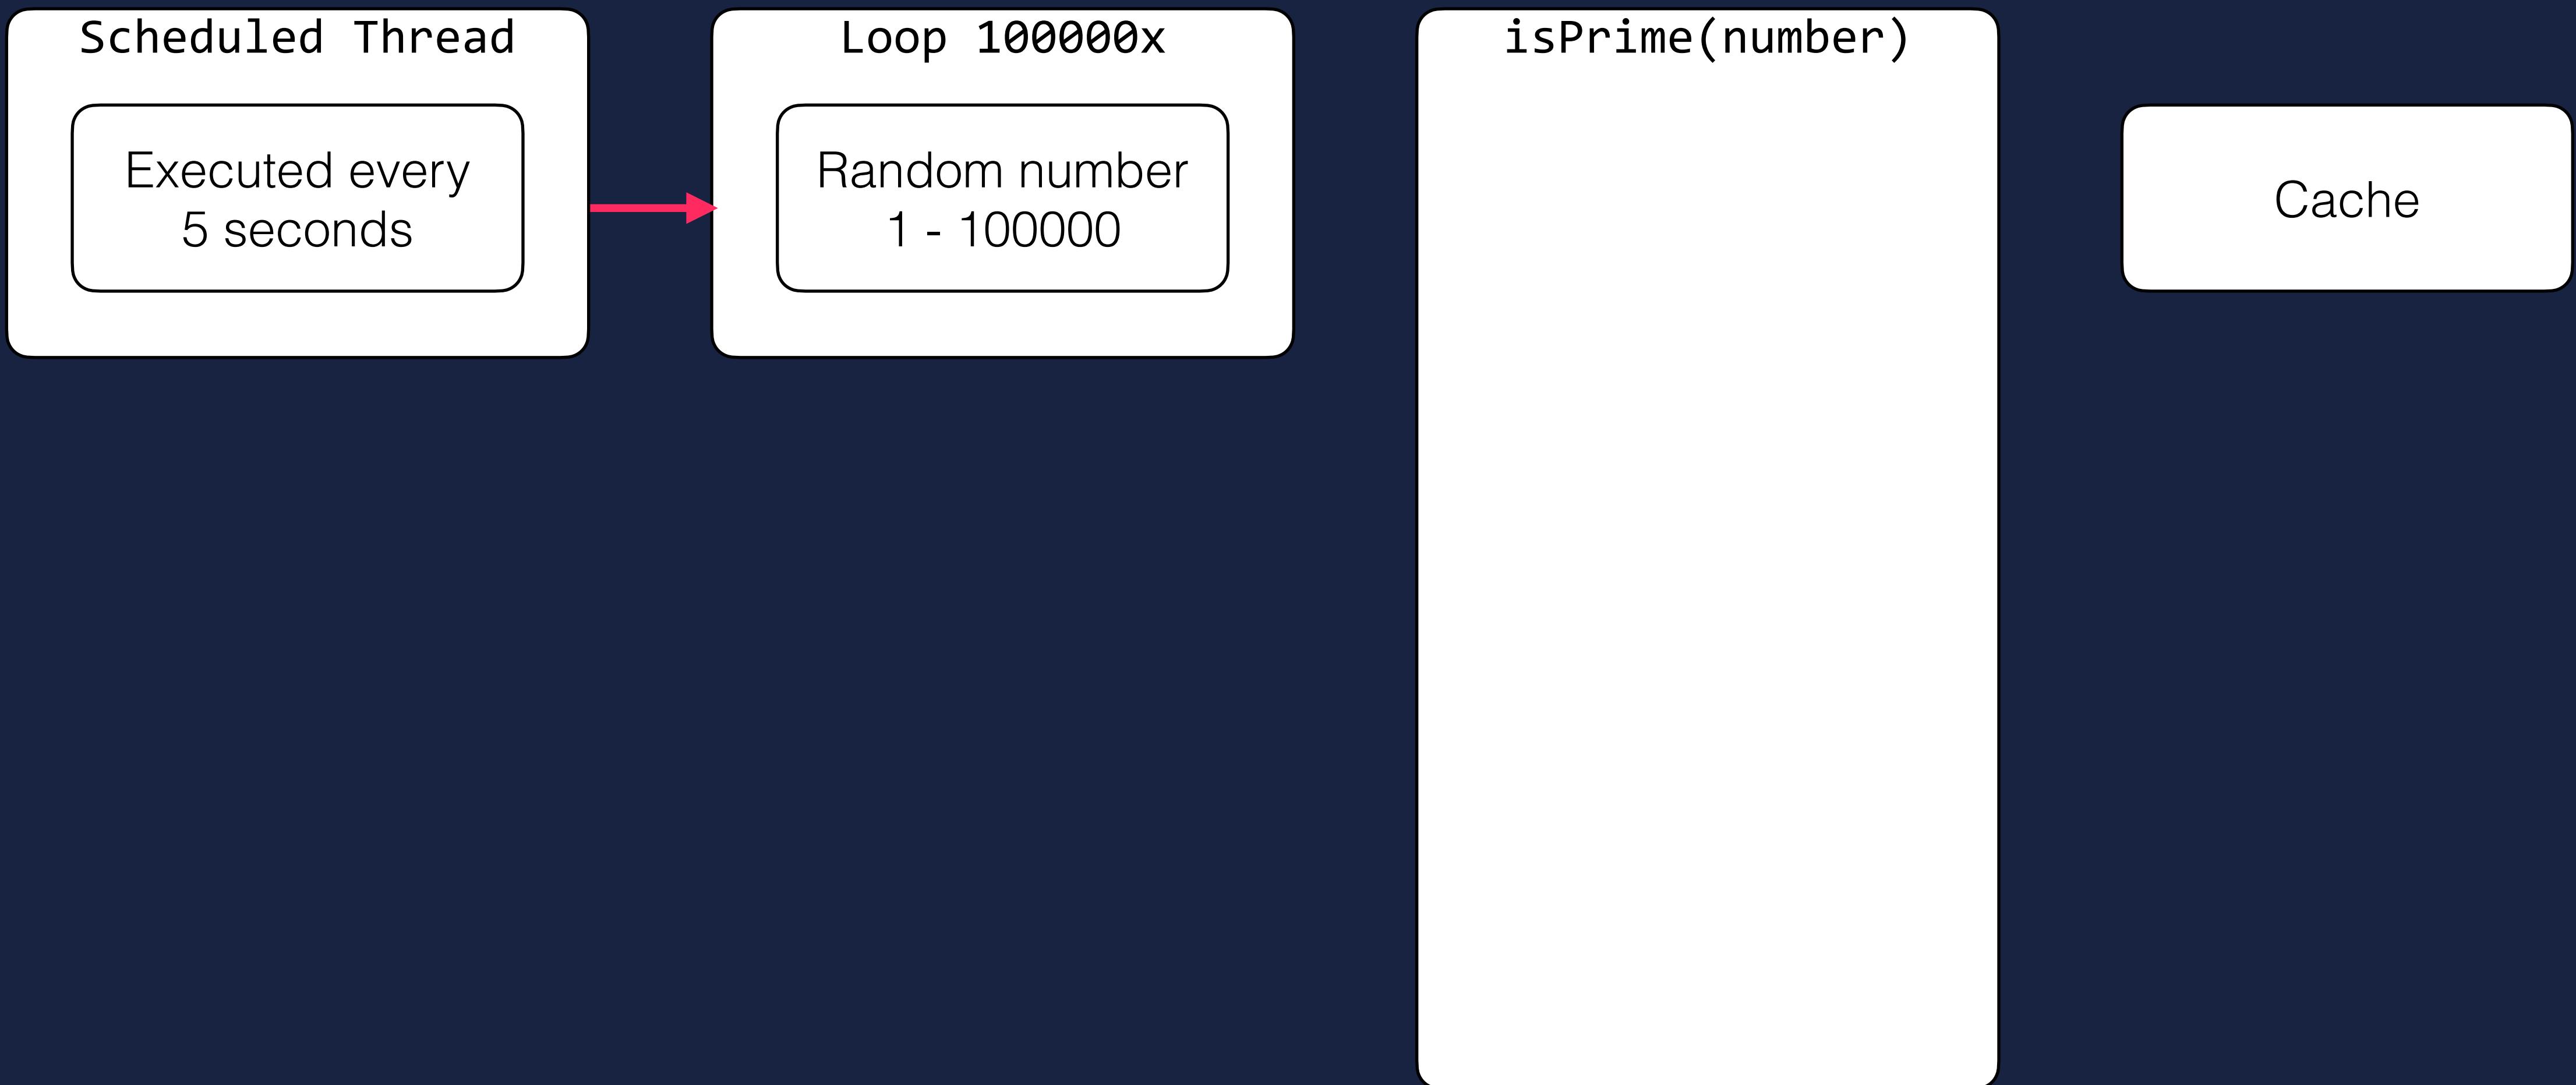
SIMPLE

EXAMPLE

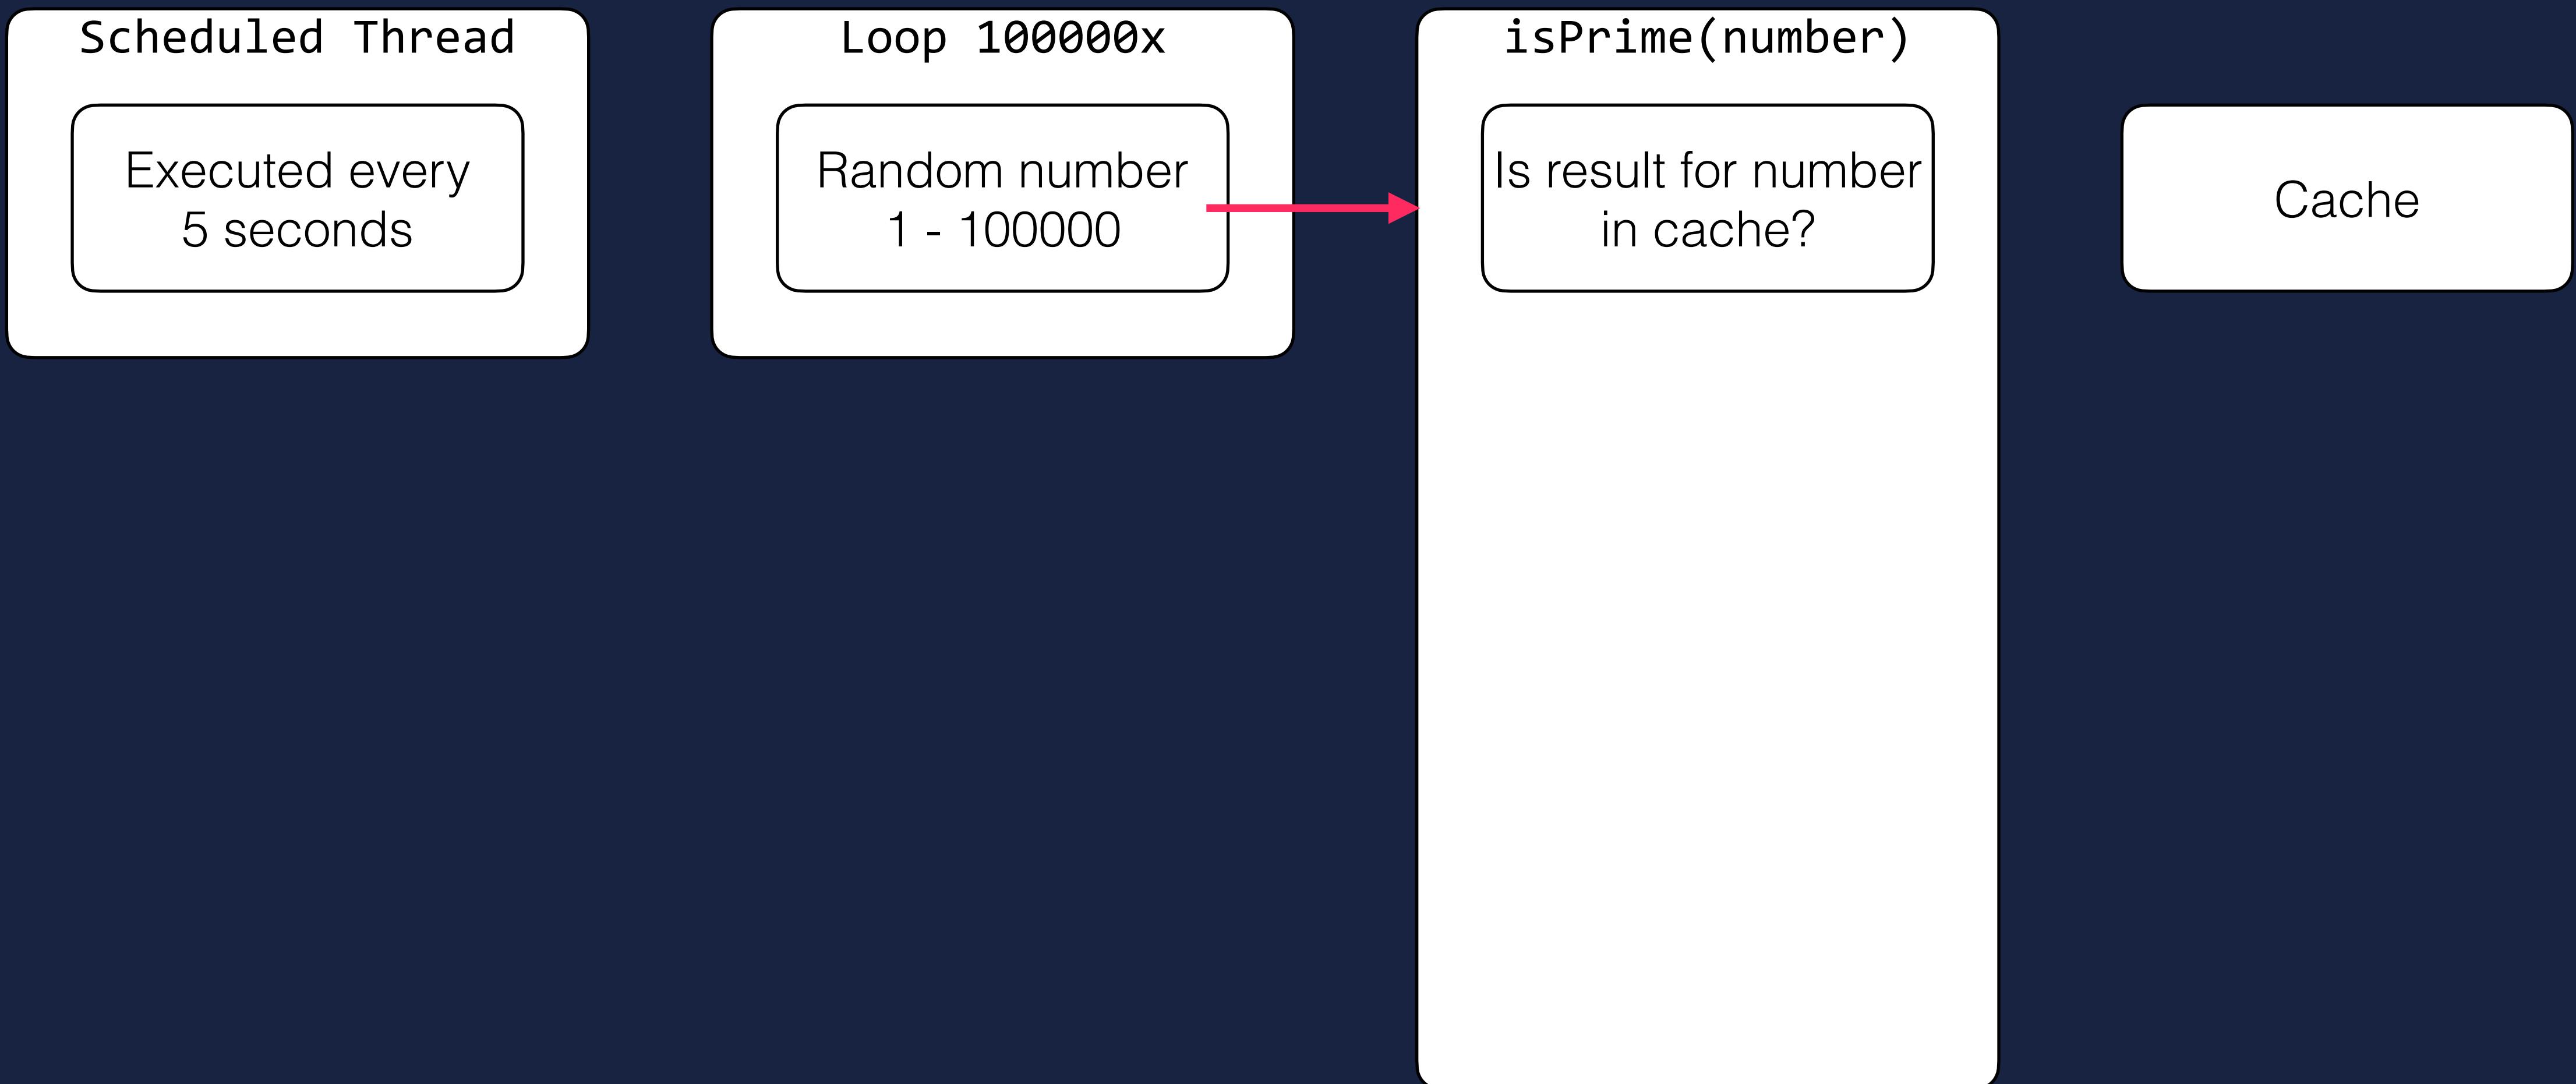
SIMPLE EXAMPLE



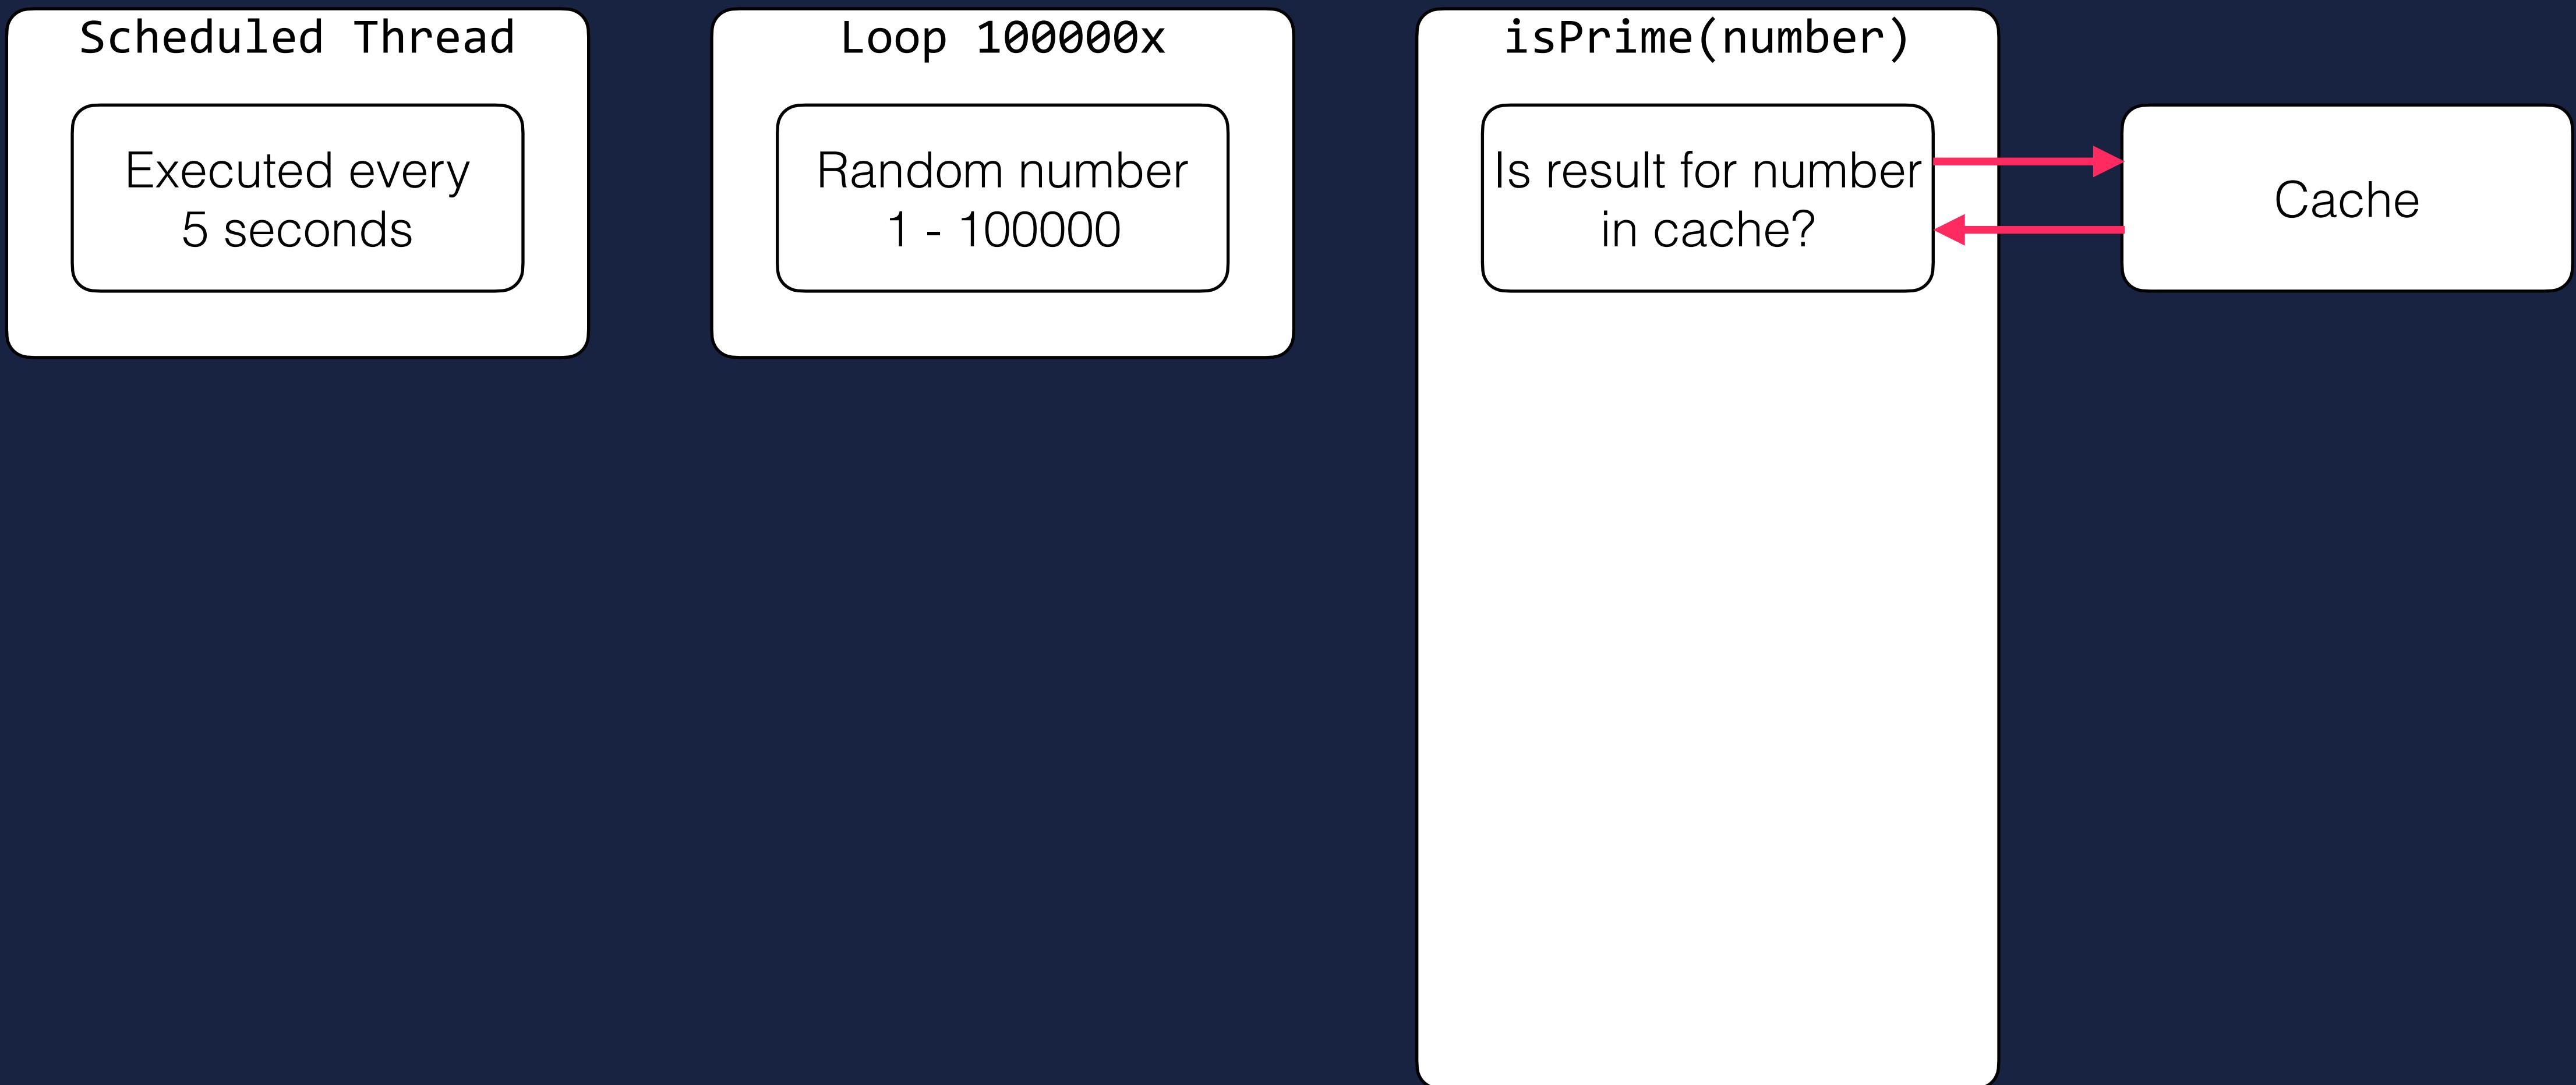
SIMPLE EXAMPLE



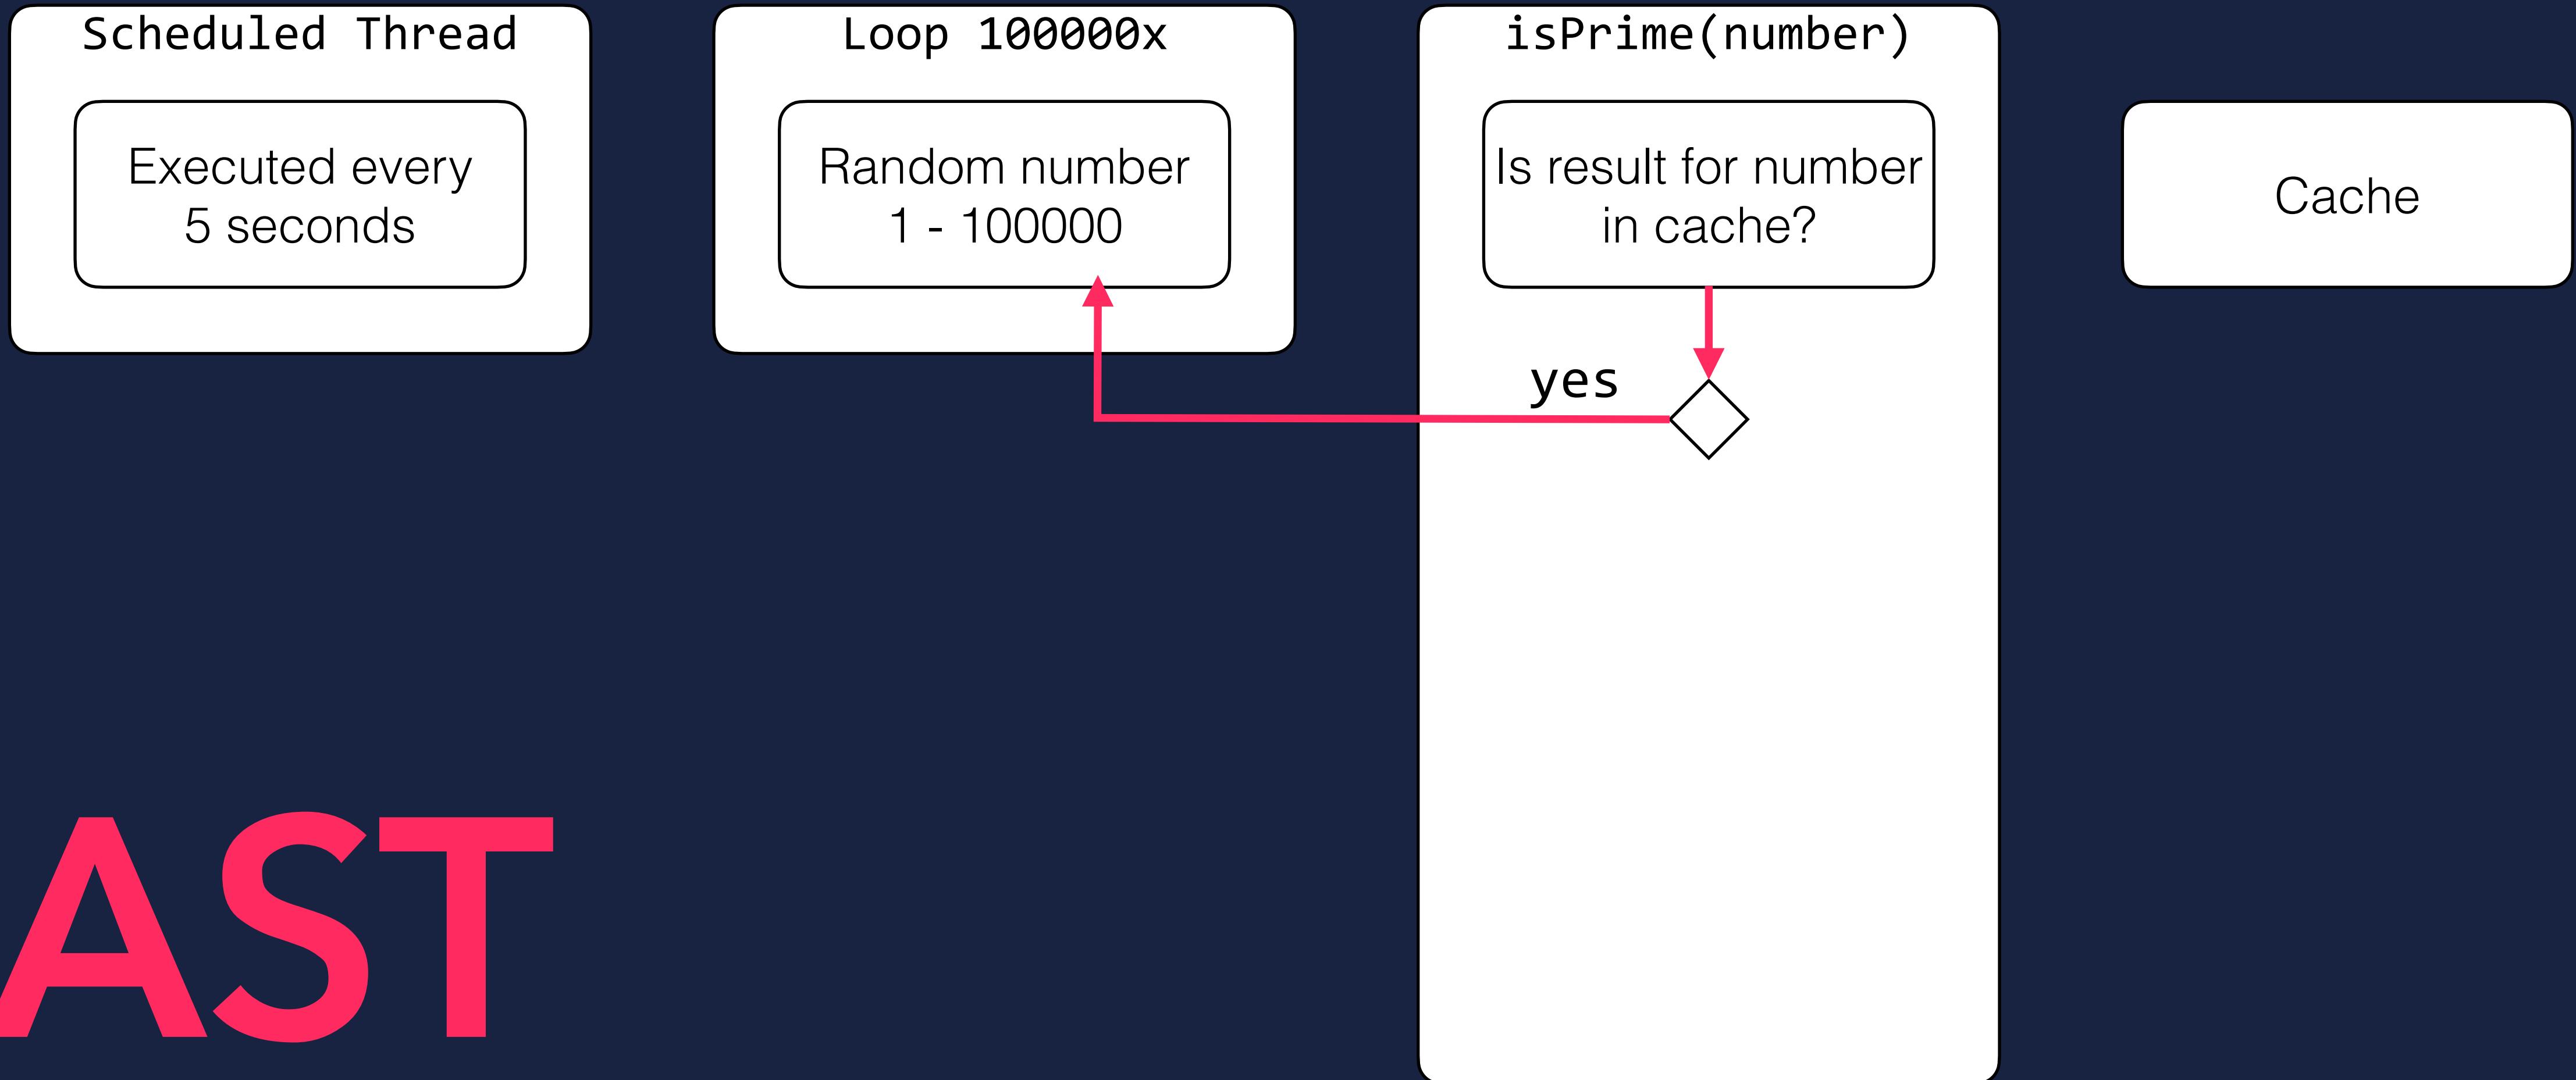
SIMPLE EXAMPLE



SIMPLE EXAMPLE

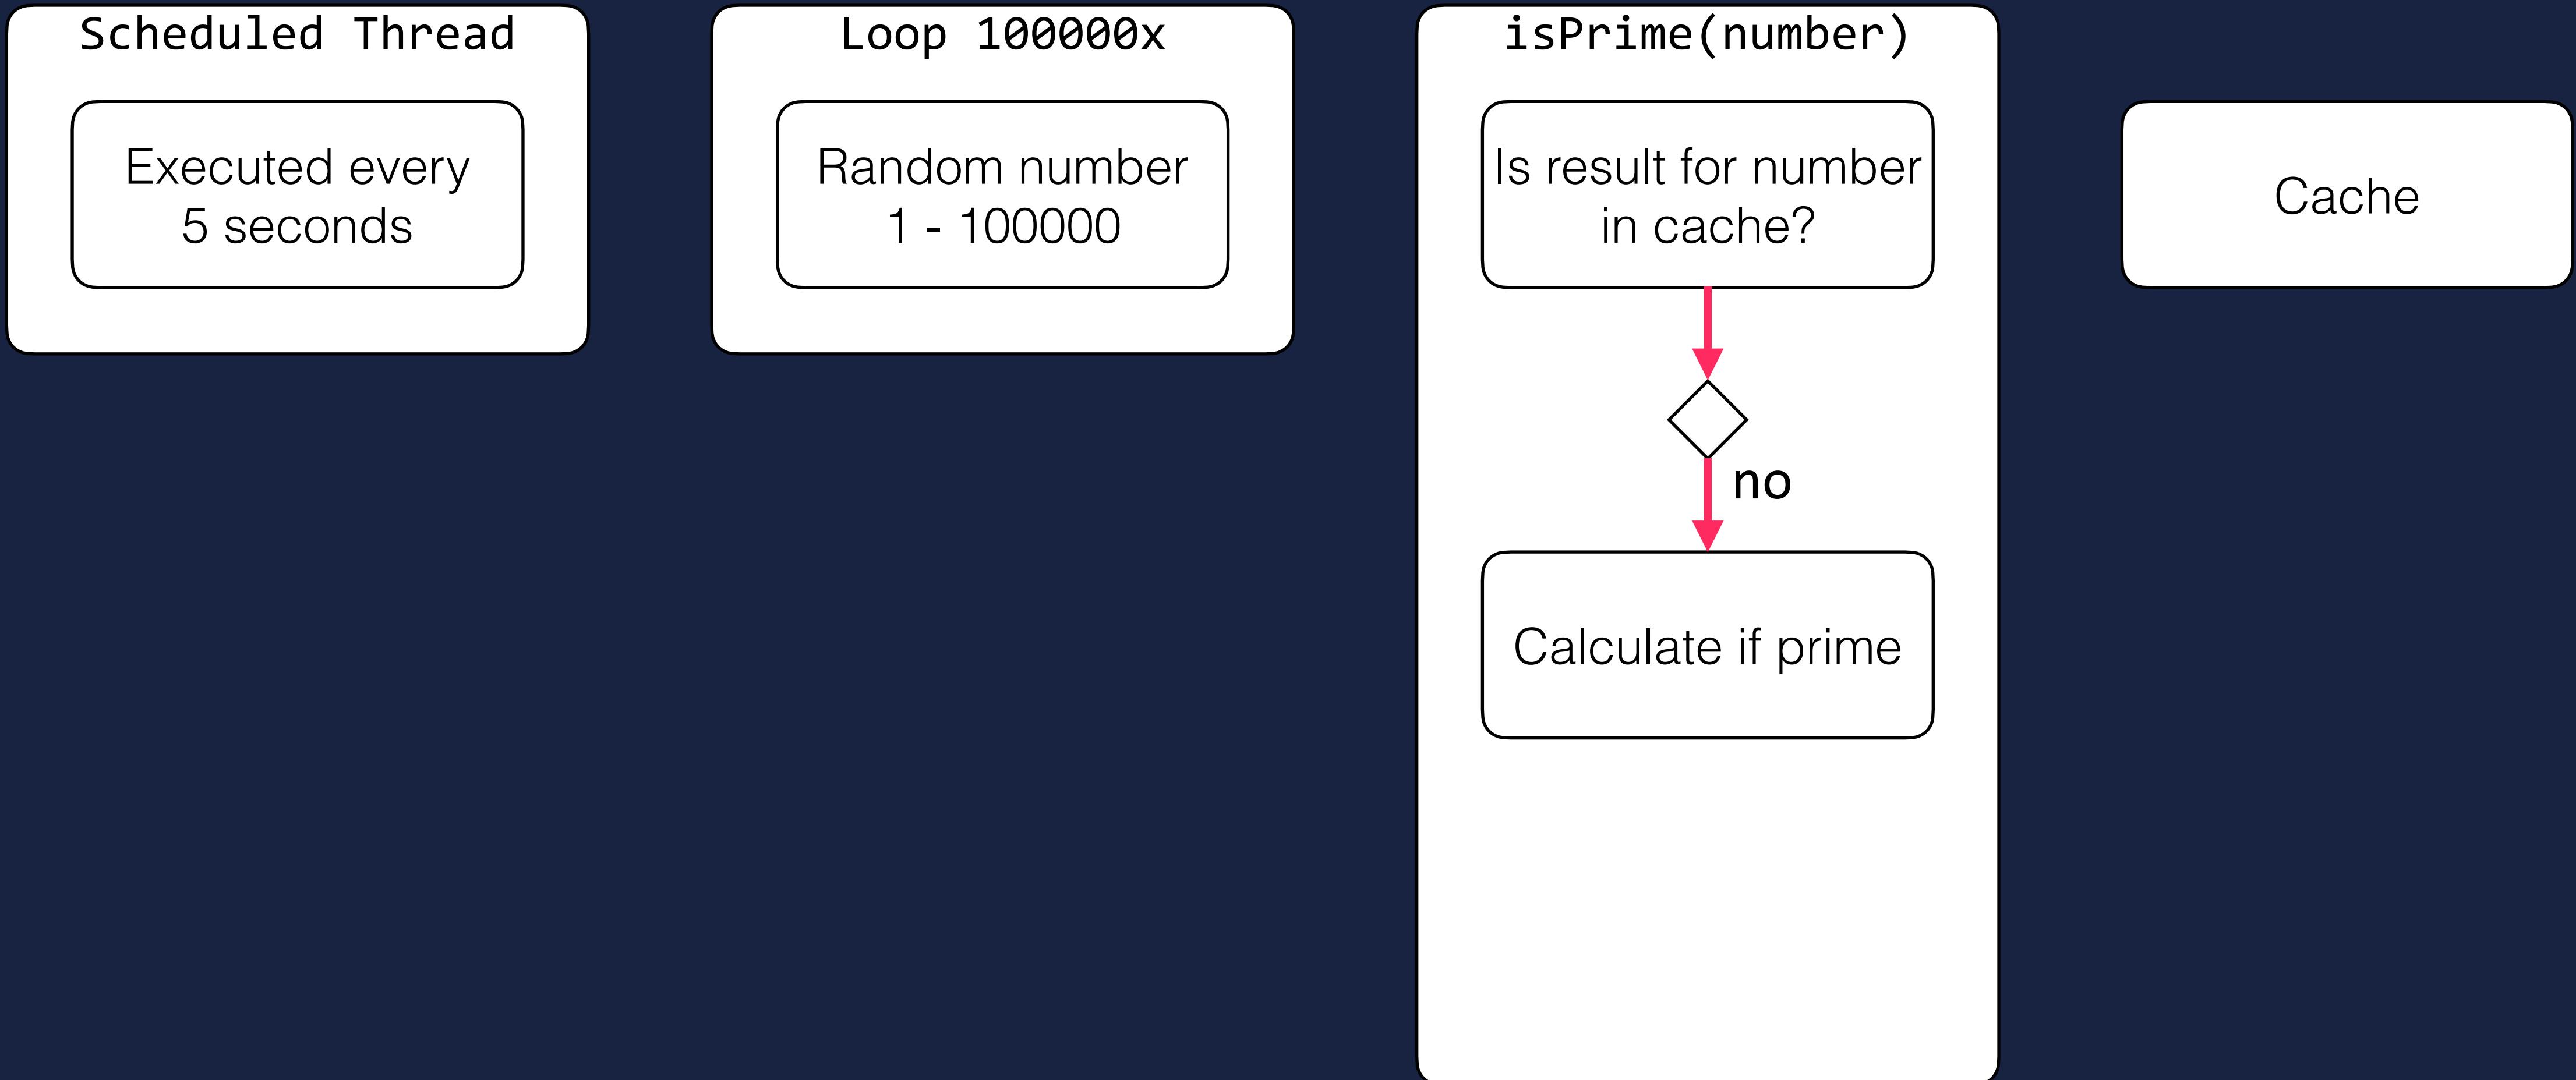


SIMPLE EXAMPLE

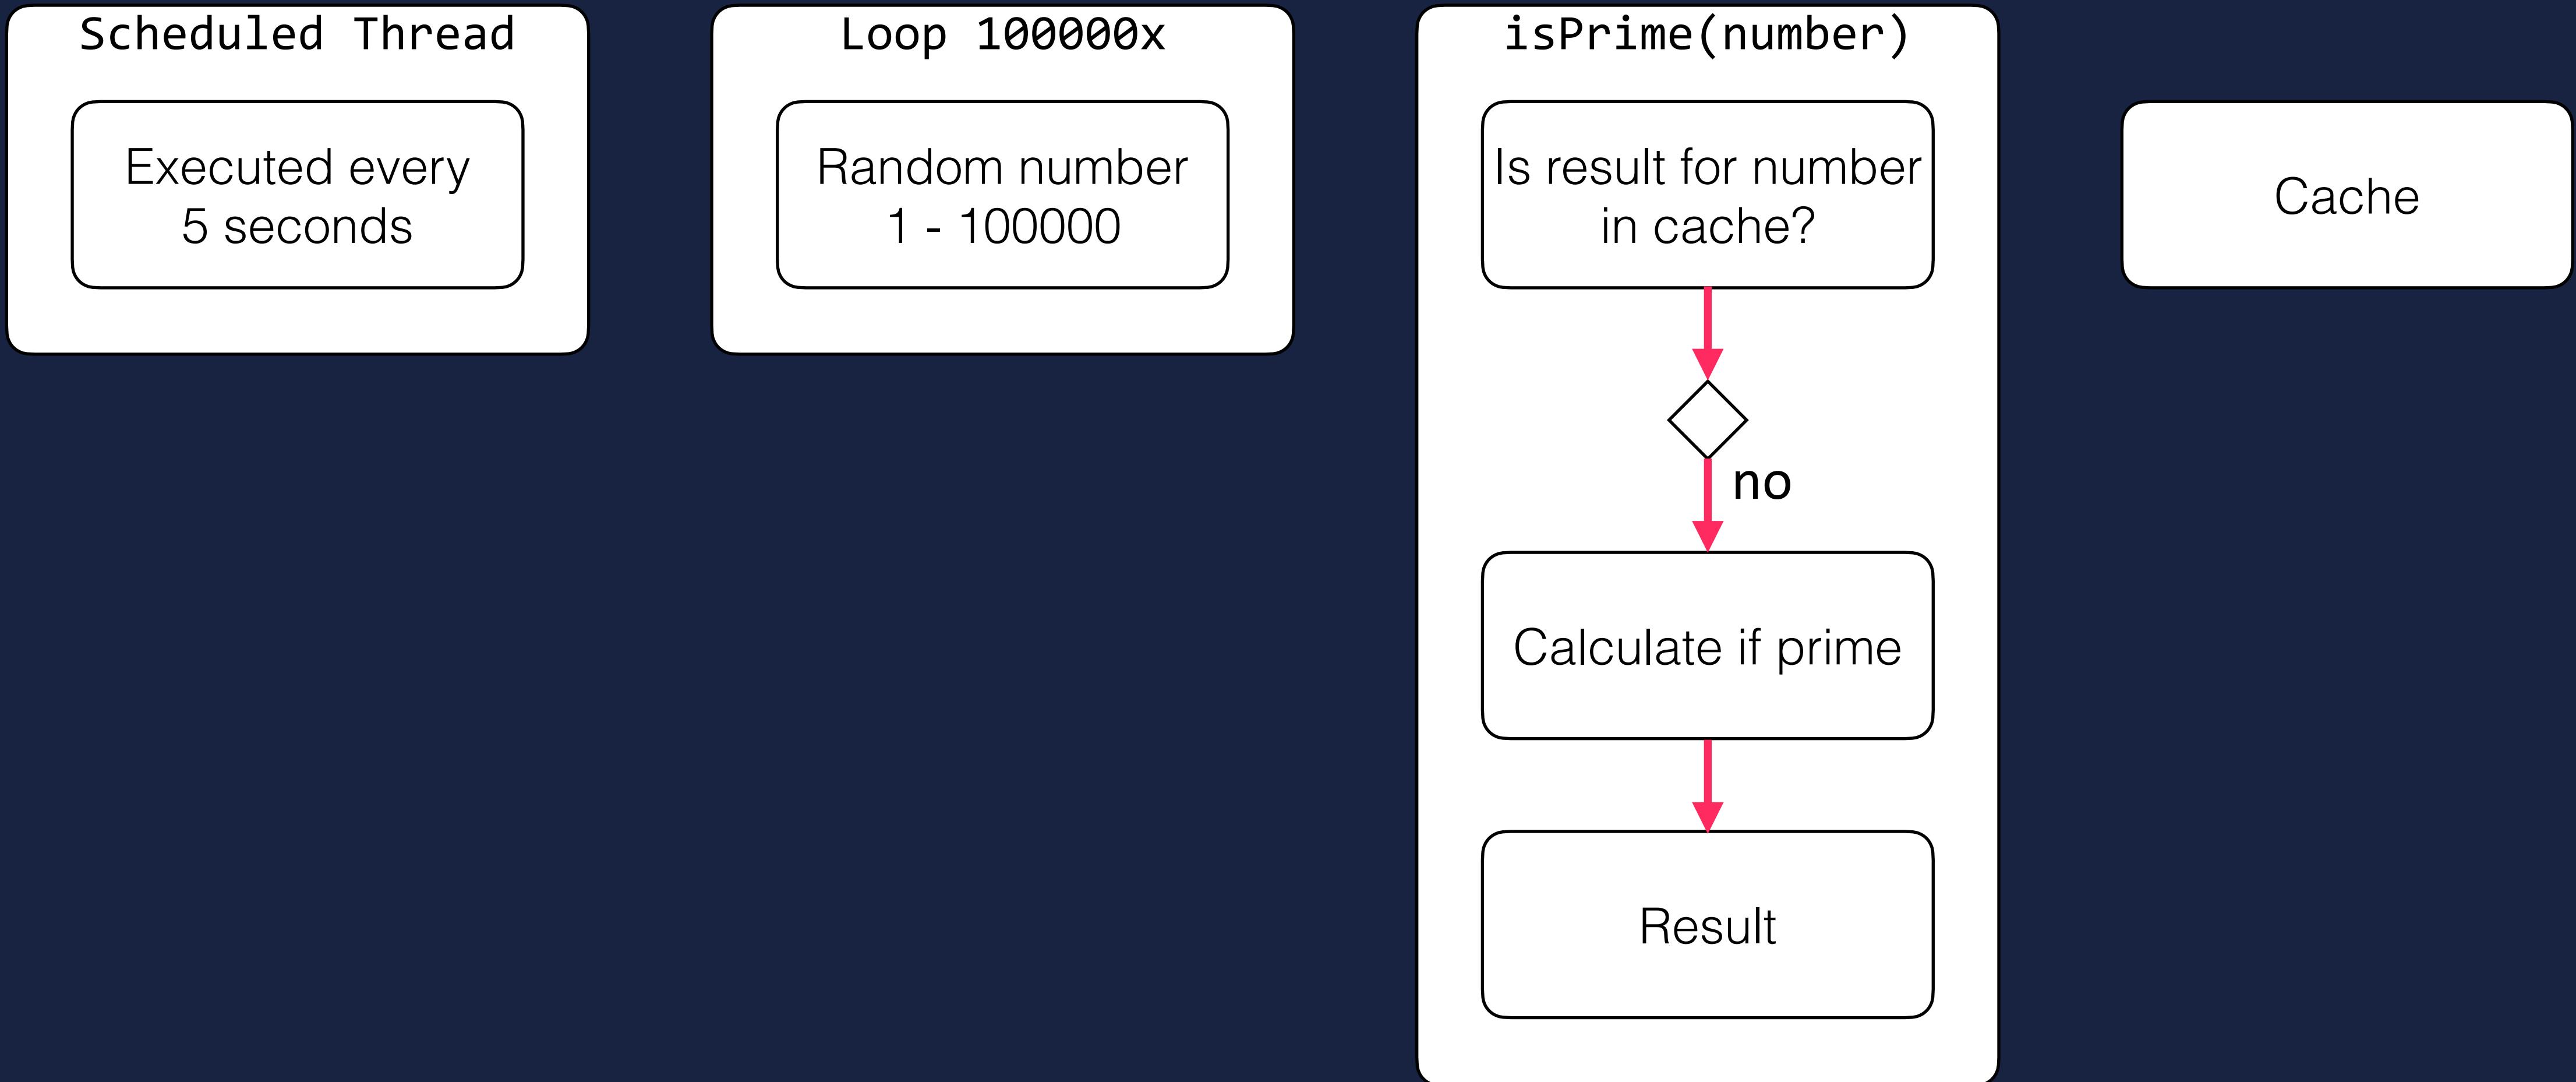


FAST

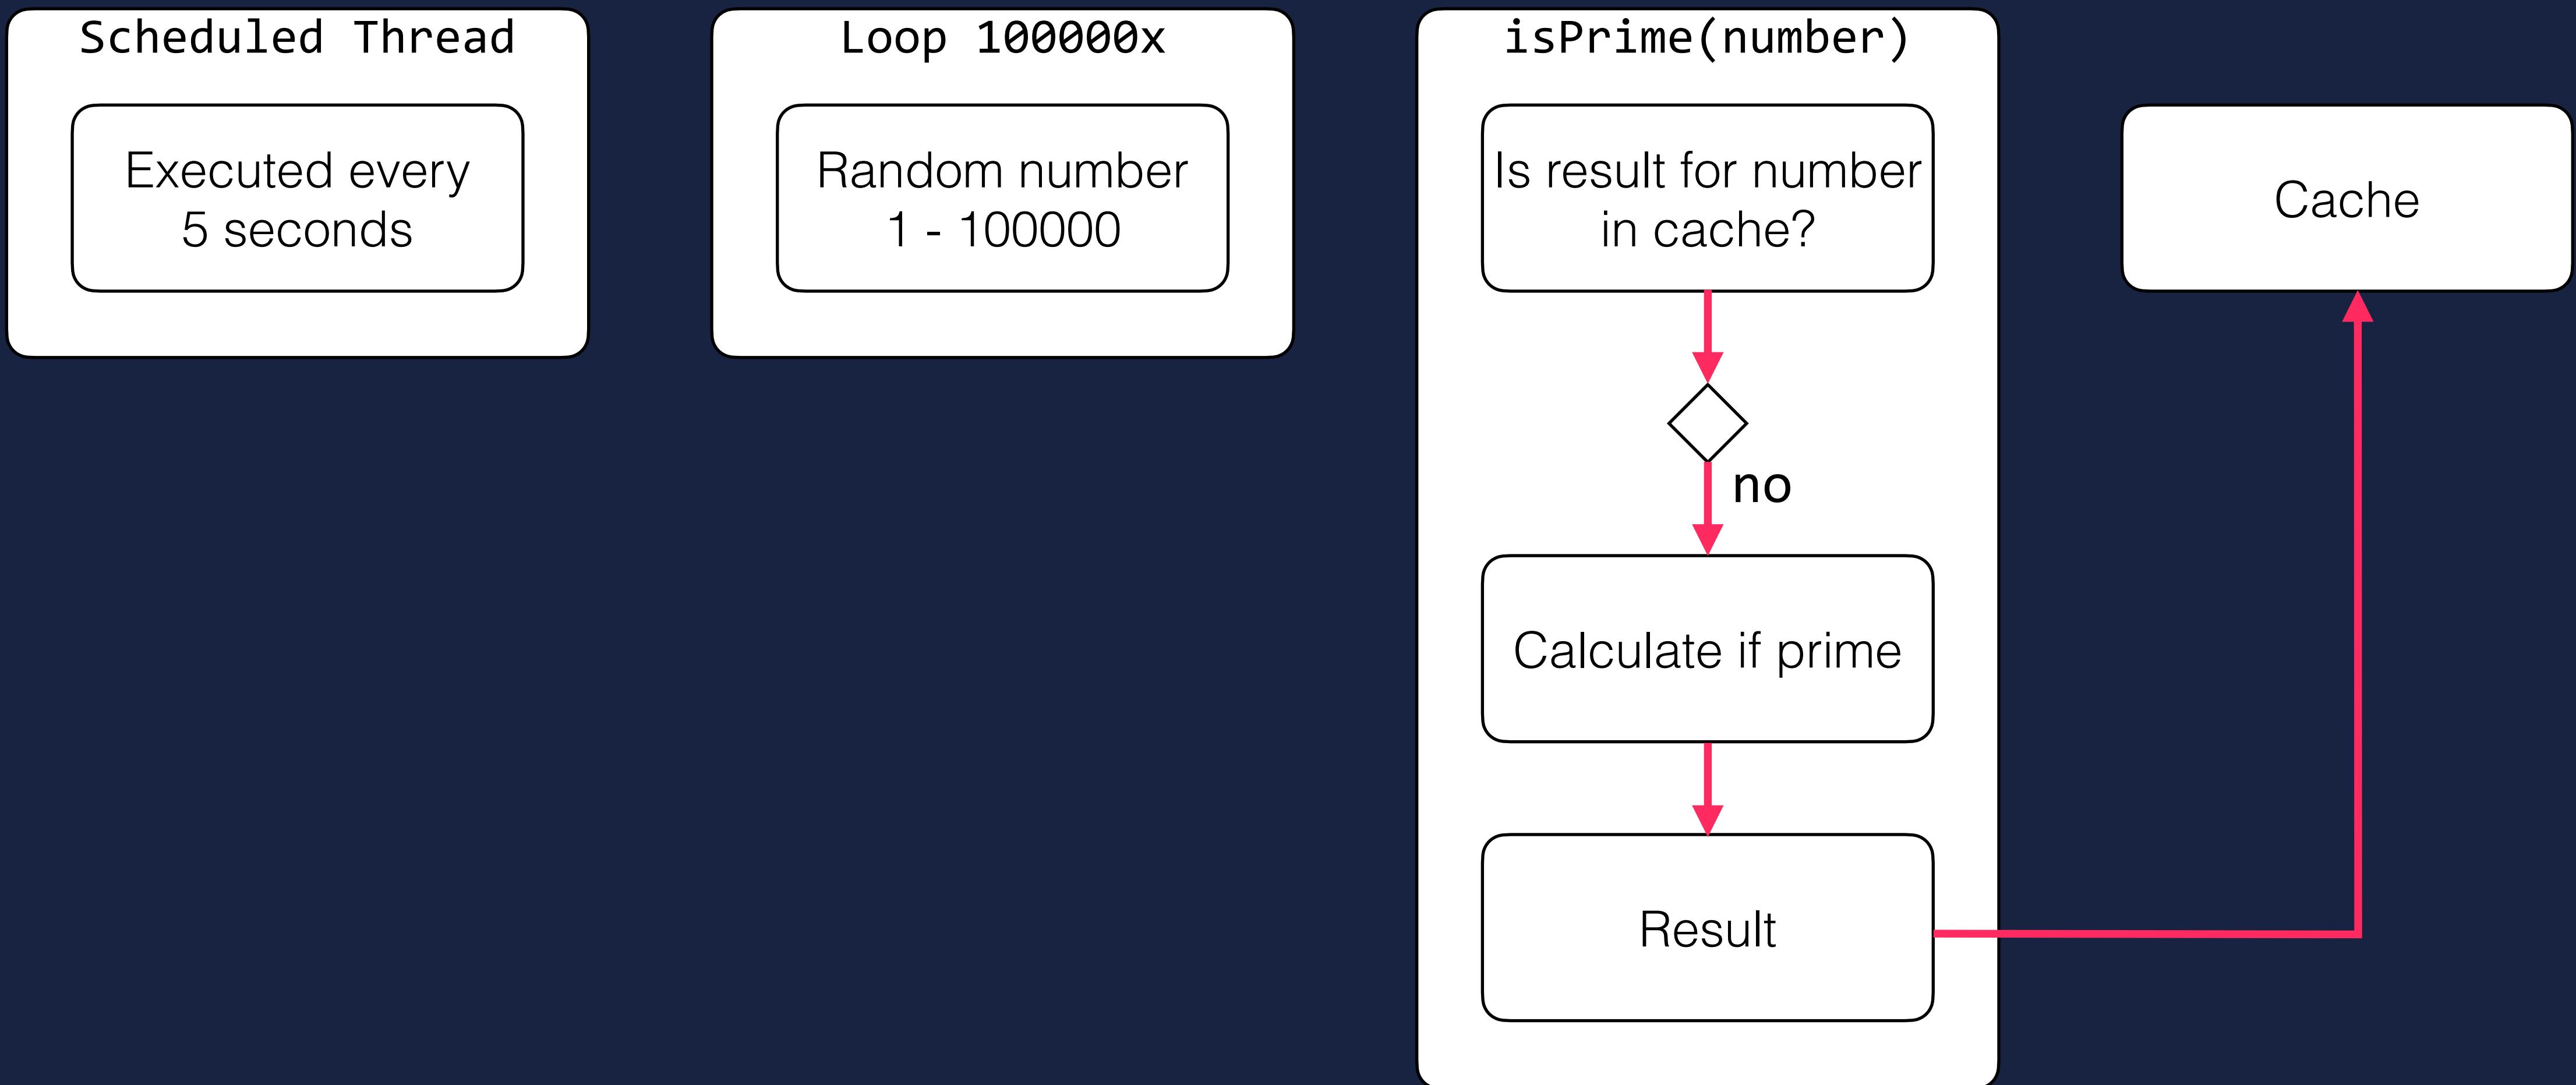
SIMPLE EXAMPLE



SIMPLE EXAMPLE

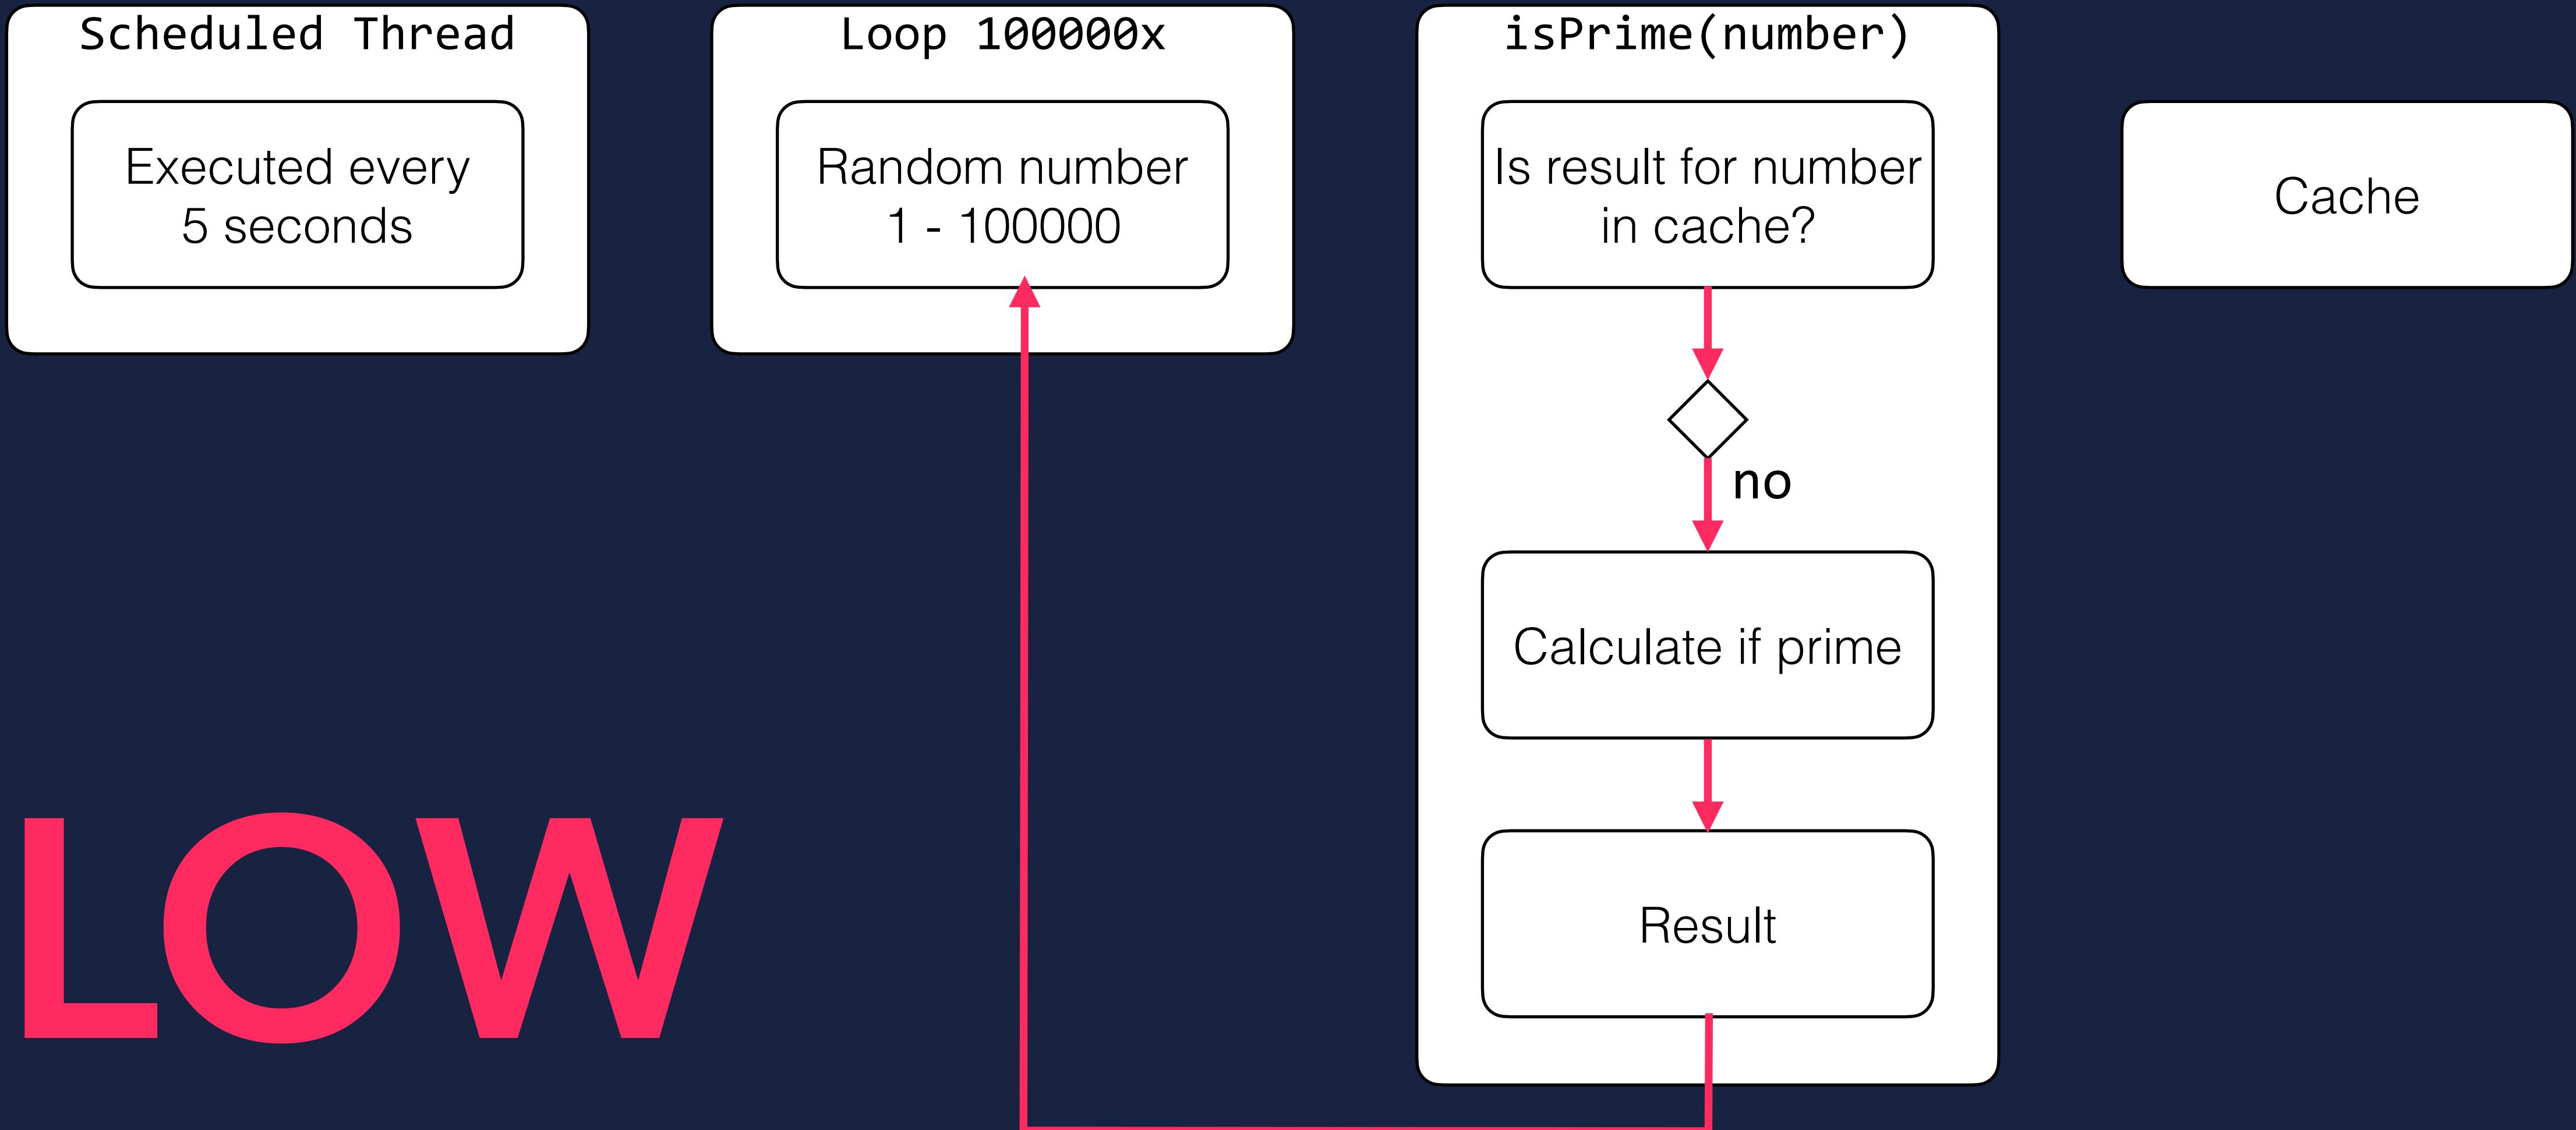


SIMPLE EXAMPLE

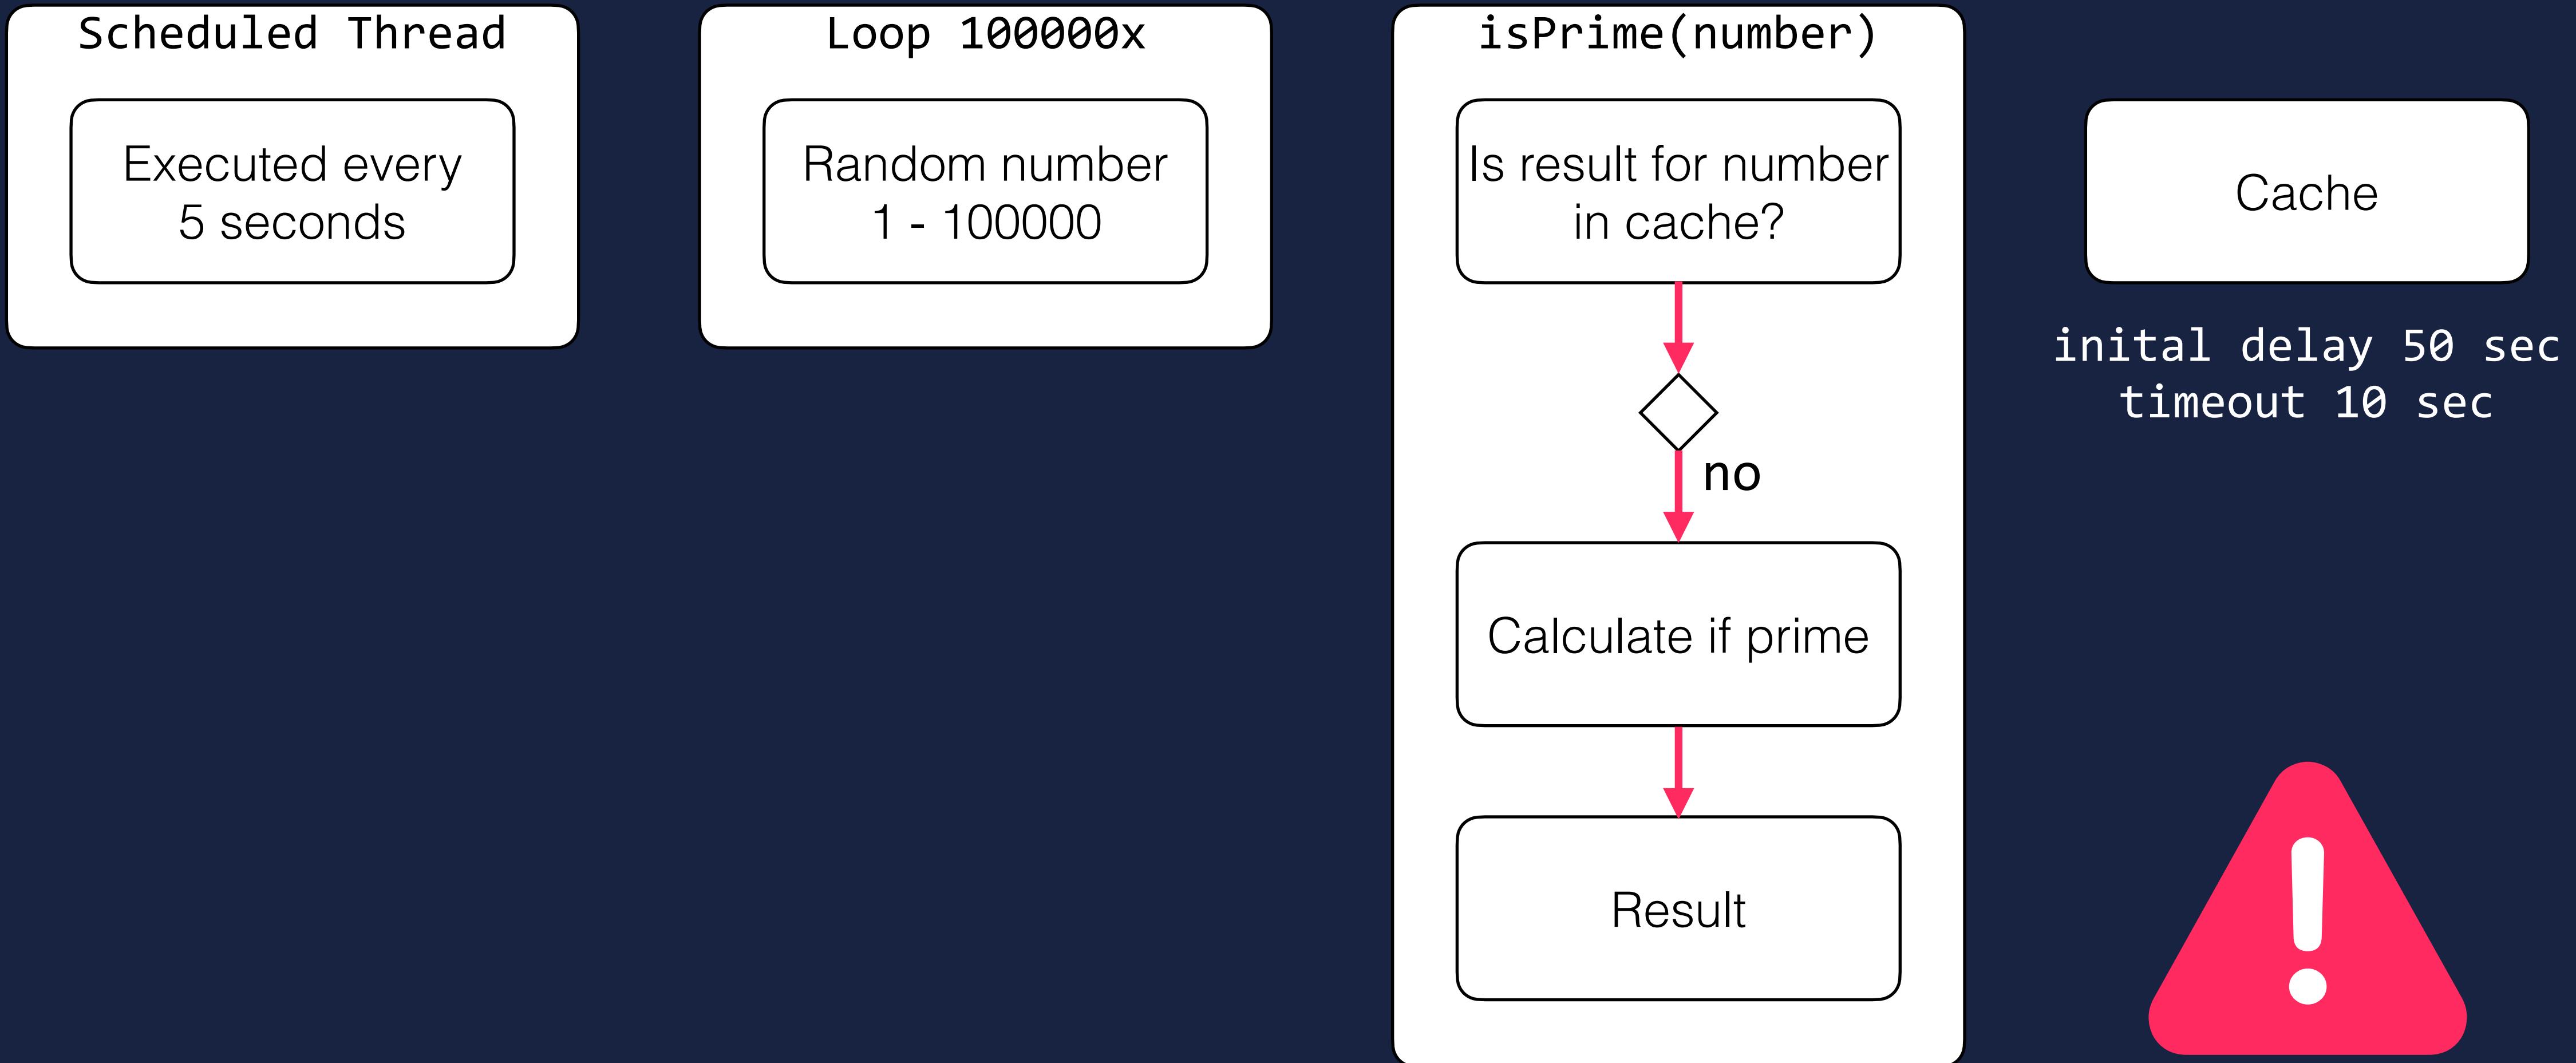


SIMPLE EXAMPLE

SLOW



SIMPLE EXAMPLE

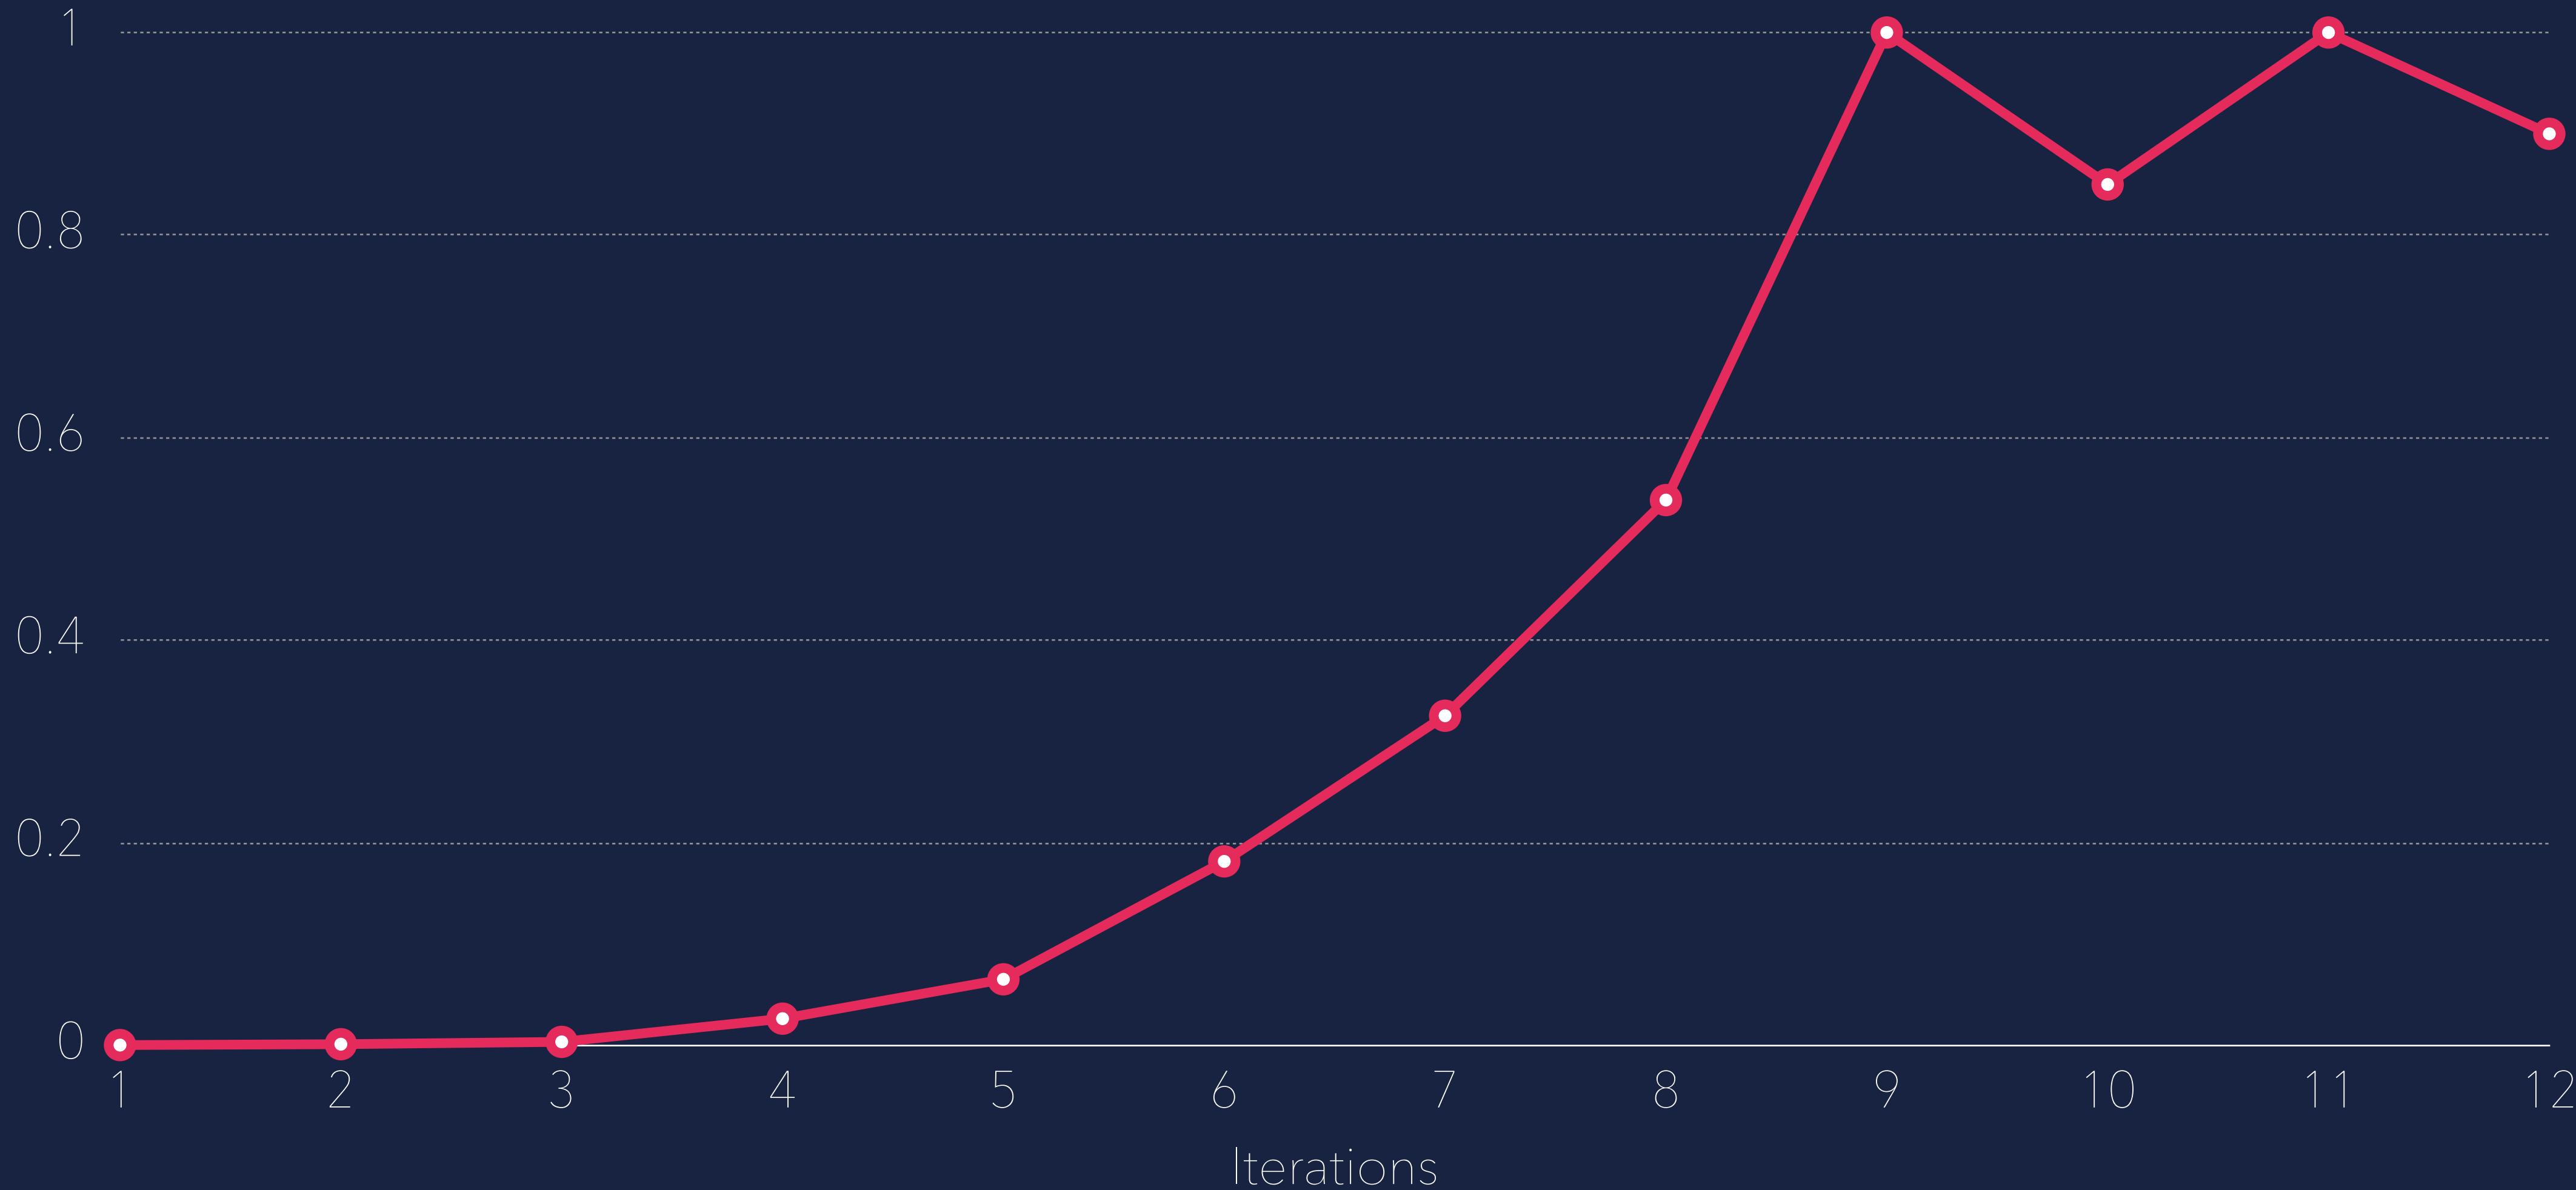


SIMPLE EXAMPLE

- First run is slow because the cache is empty
- Time to check 100 000 random numbers for prime should decrease with every run

SIMPLE EXAMPLE

Application Performance



DEMO

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
```

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
```



Defines where to store files
for the checkpoint

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
```

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
```

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
```

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
```



Slow because cache is not filled yet

SHELL 2

```
>
```

SIMPLEX EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
```

SHELL 2

```
>
```

SIMPLEX EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
```

SHELL 2

```
>
```

SIMPLEX EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
```

SHELL 2

```
>
```

SIMPLEX EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
```

SHELL 2

```
>
```

SIMPLEX EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
```

SHELL 2

```
>
```

SIMPLEX EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
```

SHELL 2

```
>
```

SIMPLEX EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
```

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
```

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
```

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
```

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
```



Fast because cache is filled

SHELL 2

```
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
```

SHELL 2

```
>jcmd crac4.jar JDK.checkpoint
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
```

SHELL 2

```
>jcmd crac4.jar JDK.checkpoint
```

CREATE CHECKPOINT

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
beforeCheckpoint() called in Main
beforeCheckpoint() called in GenericCache
CR: Checkpoint ...
Killed
>
```

SHELL 2

```
>jccmd crac4.jar JDK.checkpoint
4240:
Command executed successfully
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
beforeCheckpoint() called in Main
beforeCheckpoint() called in GenericCache
CR: Checkpoint ...
Killed
>
```

SHELL 2

```
>jccmd crac4.jar JDK.checkpoint
4240:
Command executed successfully
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
beforeCheckpoint() called in Main
beforeCheckpoint() called in GenericCache
CR: Checkpoint ...
Killed
>java -XX:CRaCRestoreFrom=/crac-files
```

SHELL 2

```
>jcmd crac4.jar JDK.checkpoint
4240:
Command executed successfully
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
beforeCheckpoint() called in Main
beforeCheckpoint() called in GenericCache
CR: Checkpoint ...
Killed
>java -XX:CRaCRestoreFrom=/crac-files
```

RESTORE

SHELL 2

```
>jcmd crac4.jar JDK.checkpoint
4240:
Command executed successfully
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
beforeCheckpoint() called in Main
beforeCheckpoint() called in GenericCache
CR: Checkpoint ...
Killed
>java -XX:CRaCRestoreFrom=/crac-files
afterRestore() called in GenericCache
afterRestore() called in Main
12. Run: 23 ms (99999 elements cached, 100.0%)
```

SHELL 2

```
>jcmd crac4.jar JDK.checkpoint
4240:
Command executed successfully
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
beforeCheckpoint() called in Main
beforeCheckpoint() called in GenericCache
CR: Checkpoint ...
Killed
>java -XX:CRaCRestoreFrom=/crac-files
afterRestore() called in GenericCache
afterRestore() called in Main
12. Run: 23 ms (99999 elements cached, 100.0%)
```



Continues to run...fast

SHELL 2

```
>jcmsg crac4.jar JDK.checkpoint
4240:
Command executed successfully
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
beforeCheckpoint() called in Main
beforeCheckpoint() called in GenericCache
CR: Checkpoint ...
Killed
>java -XX:CRaCRestoreFrom=/crac-files
afterRestore() called in GenericCache
afterRestore() called in Main
12. Run: 23 ms (99999 elements cached, 100.0%)
13. Run: 21 ms (99995 elements cached, 100.0%)
```

SHELL 2

```
>jcmd crac4.jar JDK.checkpoint
4240:
Command executed successfully
>
```

SIMPLE EXAMPLE

SHELL 1

```
>java -XX:CRaCCheckpointTo=/crac-files -jar crac4.jar
Running on CRaC (PID 4240)
First run will take up to 30 seconds...
Register Resource: GenericCache
Register Resource: Main
1. Run: 27863 ms (63303 elements cached, 63.3%)
2. Run: 10611 ms (86580 elements cached, 86.6%)
3. Run: 3965 ms (95161 elements cached, 95.2%)
4. Run: 1428 ms (98245 elements cached, 98.2%)
5. Run: 533 ms (99340 elements cached, 99.3%)
6. Run: 207 ms (99750 elements cached, 99.8%)
7. Run: 87 ms (99898 elements cached, 99.9%)
8. Run: 42 ms (99960 elements cached, 100.0%)
9. Run: 25 ms (99986 elements cached, 100.0%)
10. Run: 20 ms (99992 elements cached, 100.0%)
11. Run: 17 ms (99999 elements cached, 100.0%)
beforeCheckpoint() called in Main
beforeCheckpoint() called in GenericCache
CR: Checkpoint ...
Killed
>java -XX:CRaCRestoreFrom=/crac-files
afterRestore() called in GenericCache
afterRestore() called in Main
12. Run: 23 ms (99999 elements cached, 100.0%)
13. Run: 21 ms (99995 elements cached, 100.0%)
14. Run: 24 ms (99997 elements cached, 100.0%)
```

SHELL 2

```
>jcmd crac4.jar JDK.checkpoint
4240:
Command executed successfully
>
```

CREATING A CHECKPOINT

FROM THE COMMAND LINE:

```
>jcmd YOUR_AWESOME.jar JDK.checkpoint
```

CREATING A CHECKPOINT

FROM CODE:

```
Core.checkpointRestore();
```

WHEN TO CHECKPOINT ?

- Run your app with your typical workload
- Use the parameter `-XX:+PrintCompilation`
- Observe the moment the compilations are ramped down
- Create the checkpoint

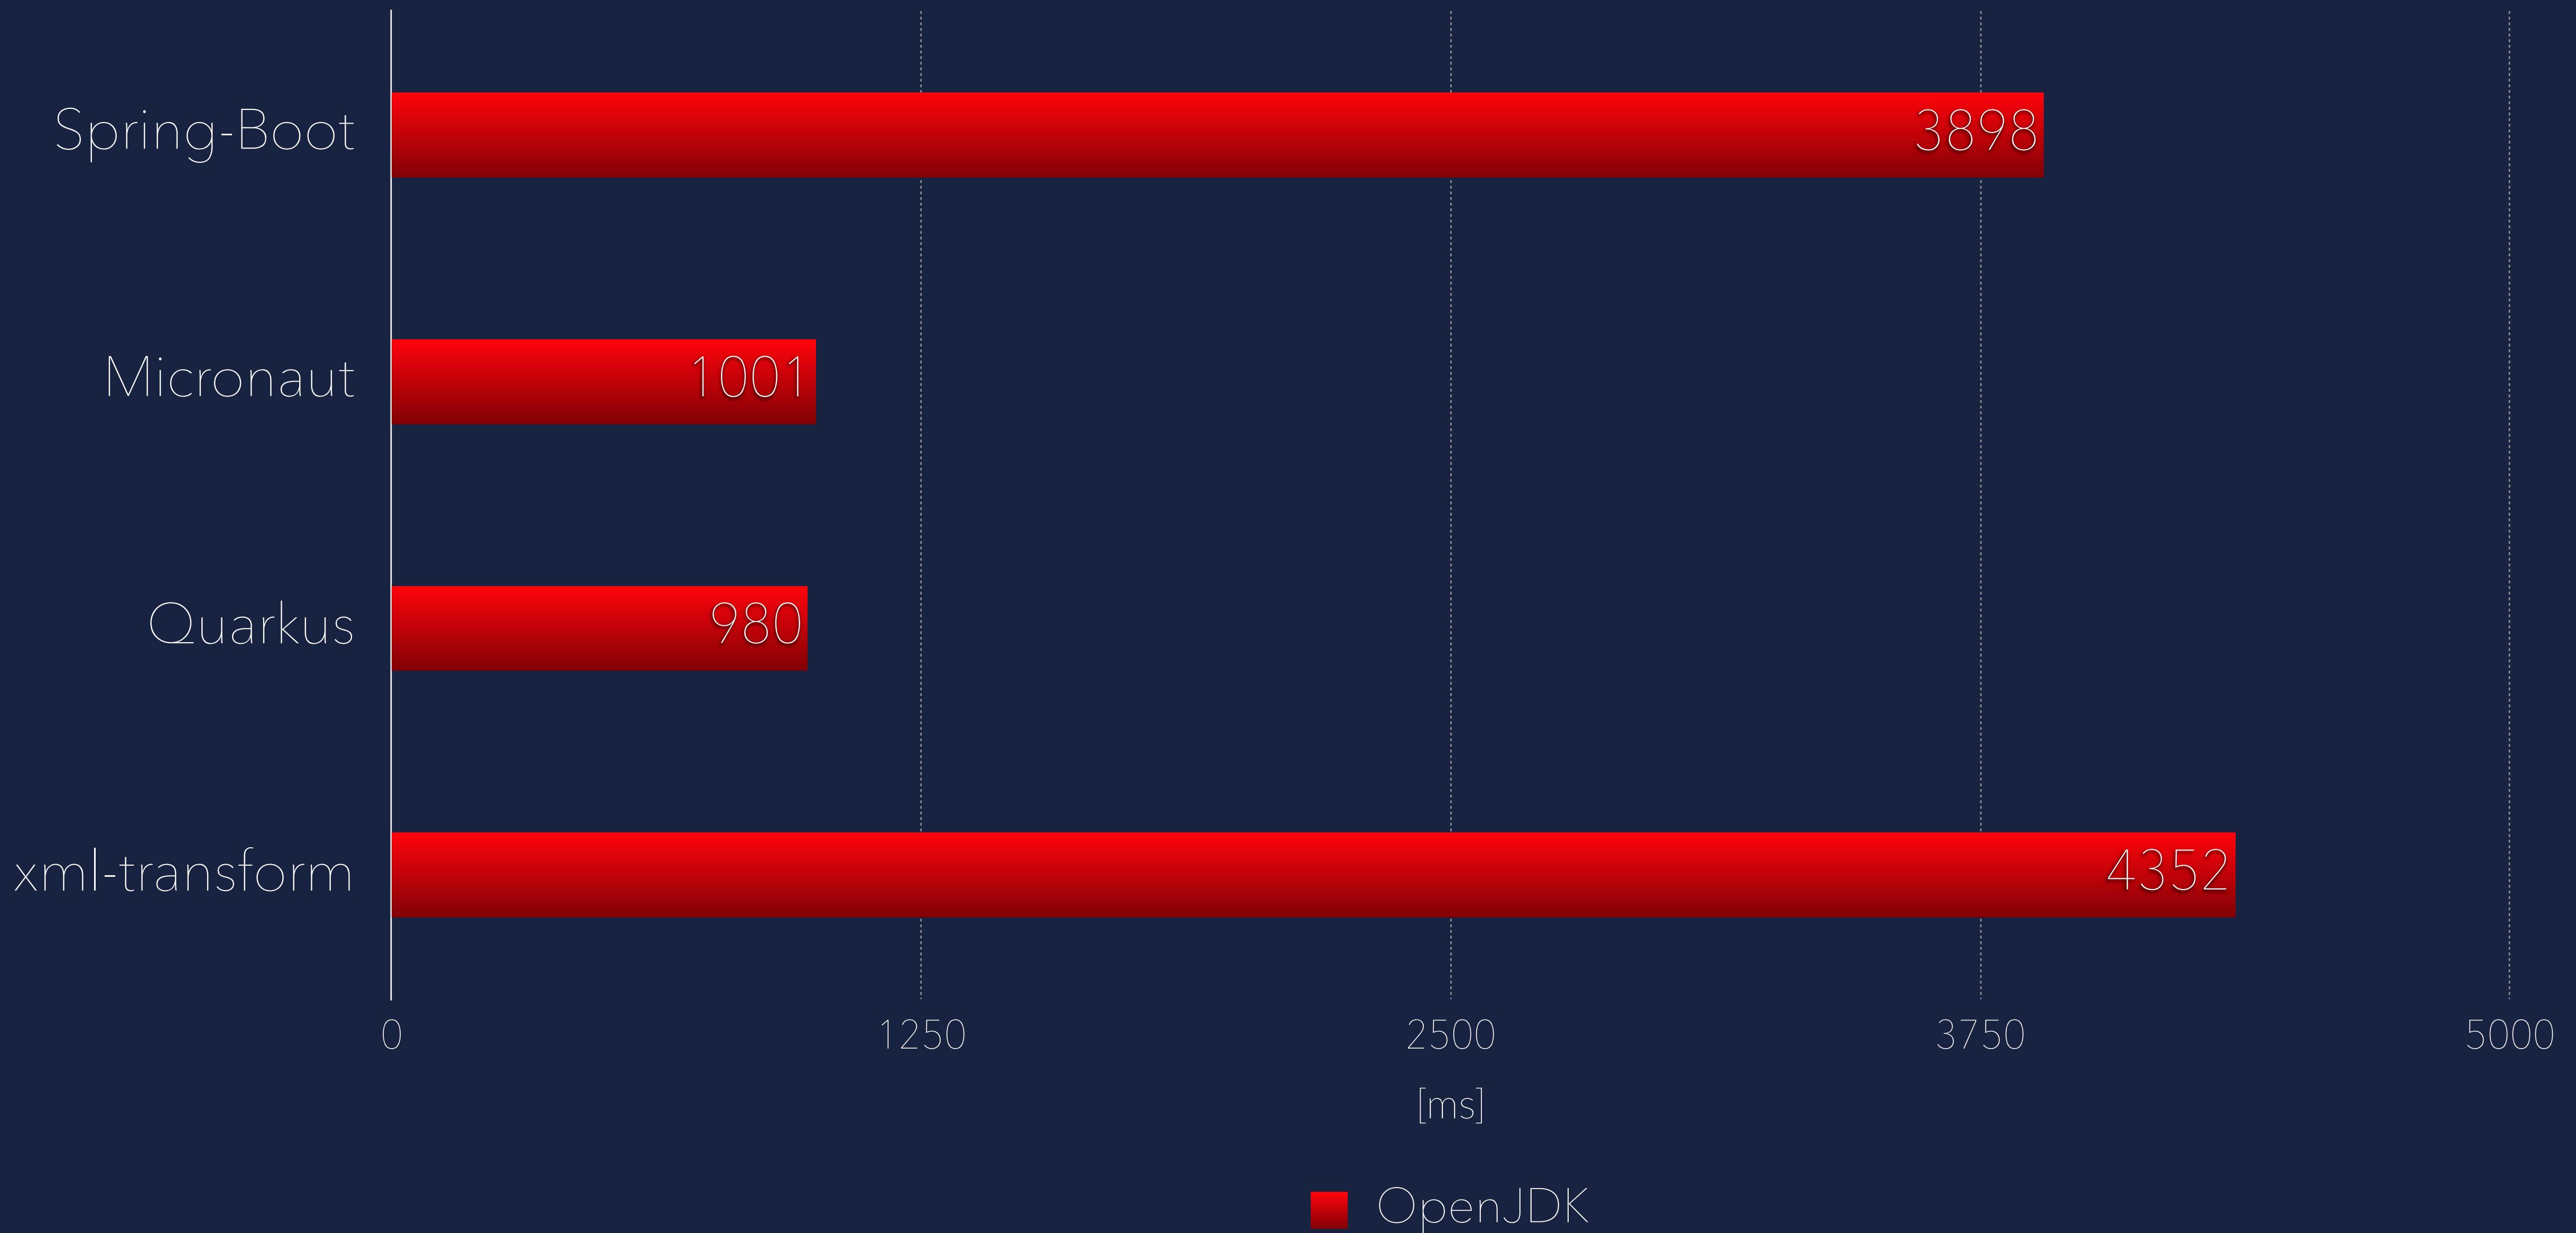
github.com/HanSolo/crac4



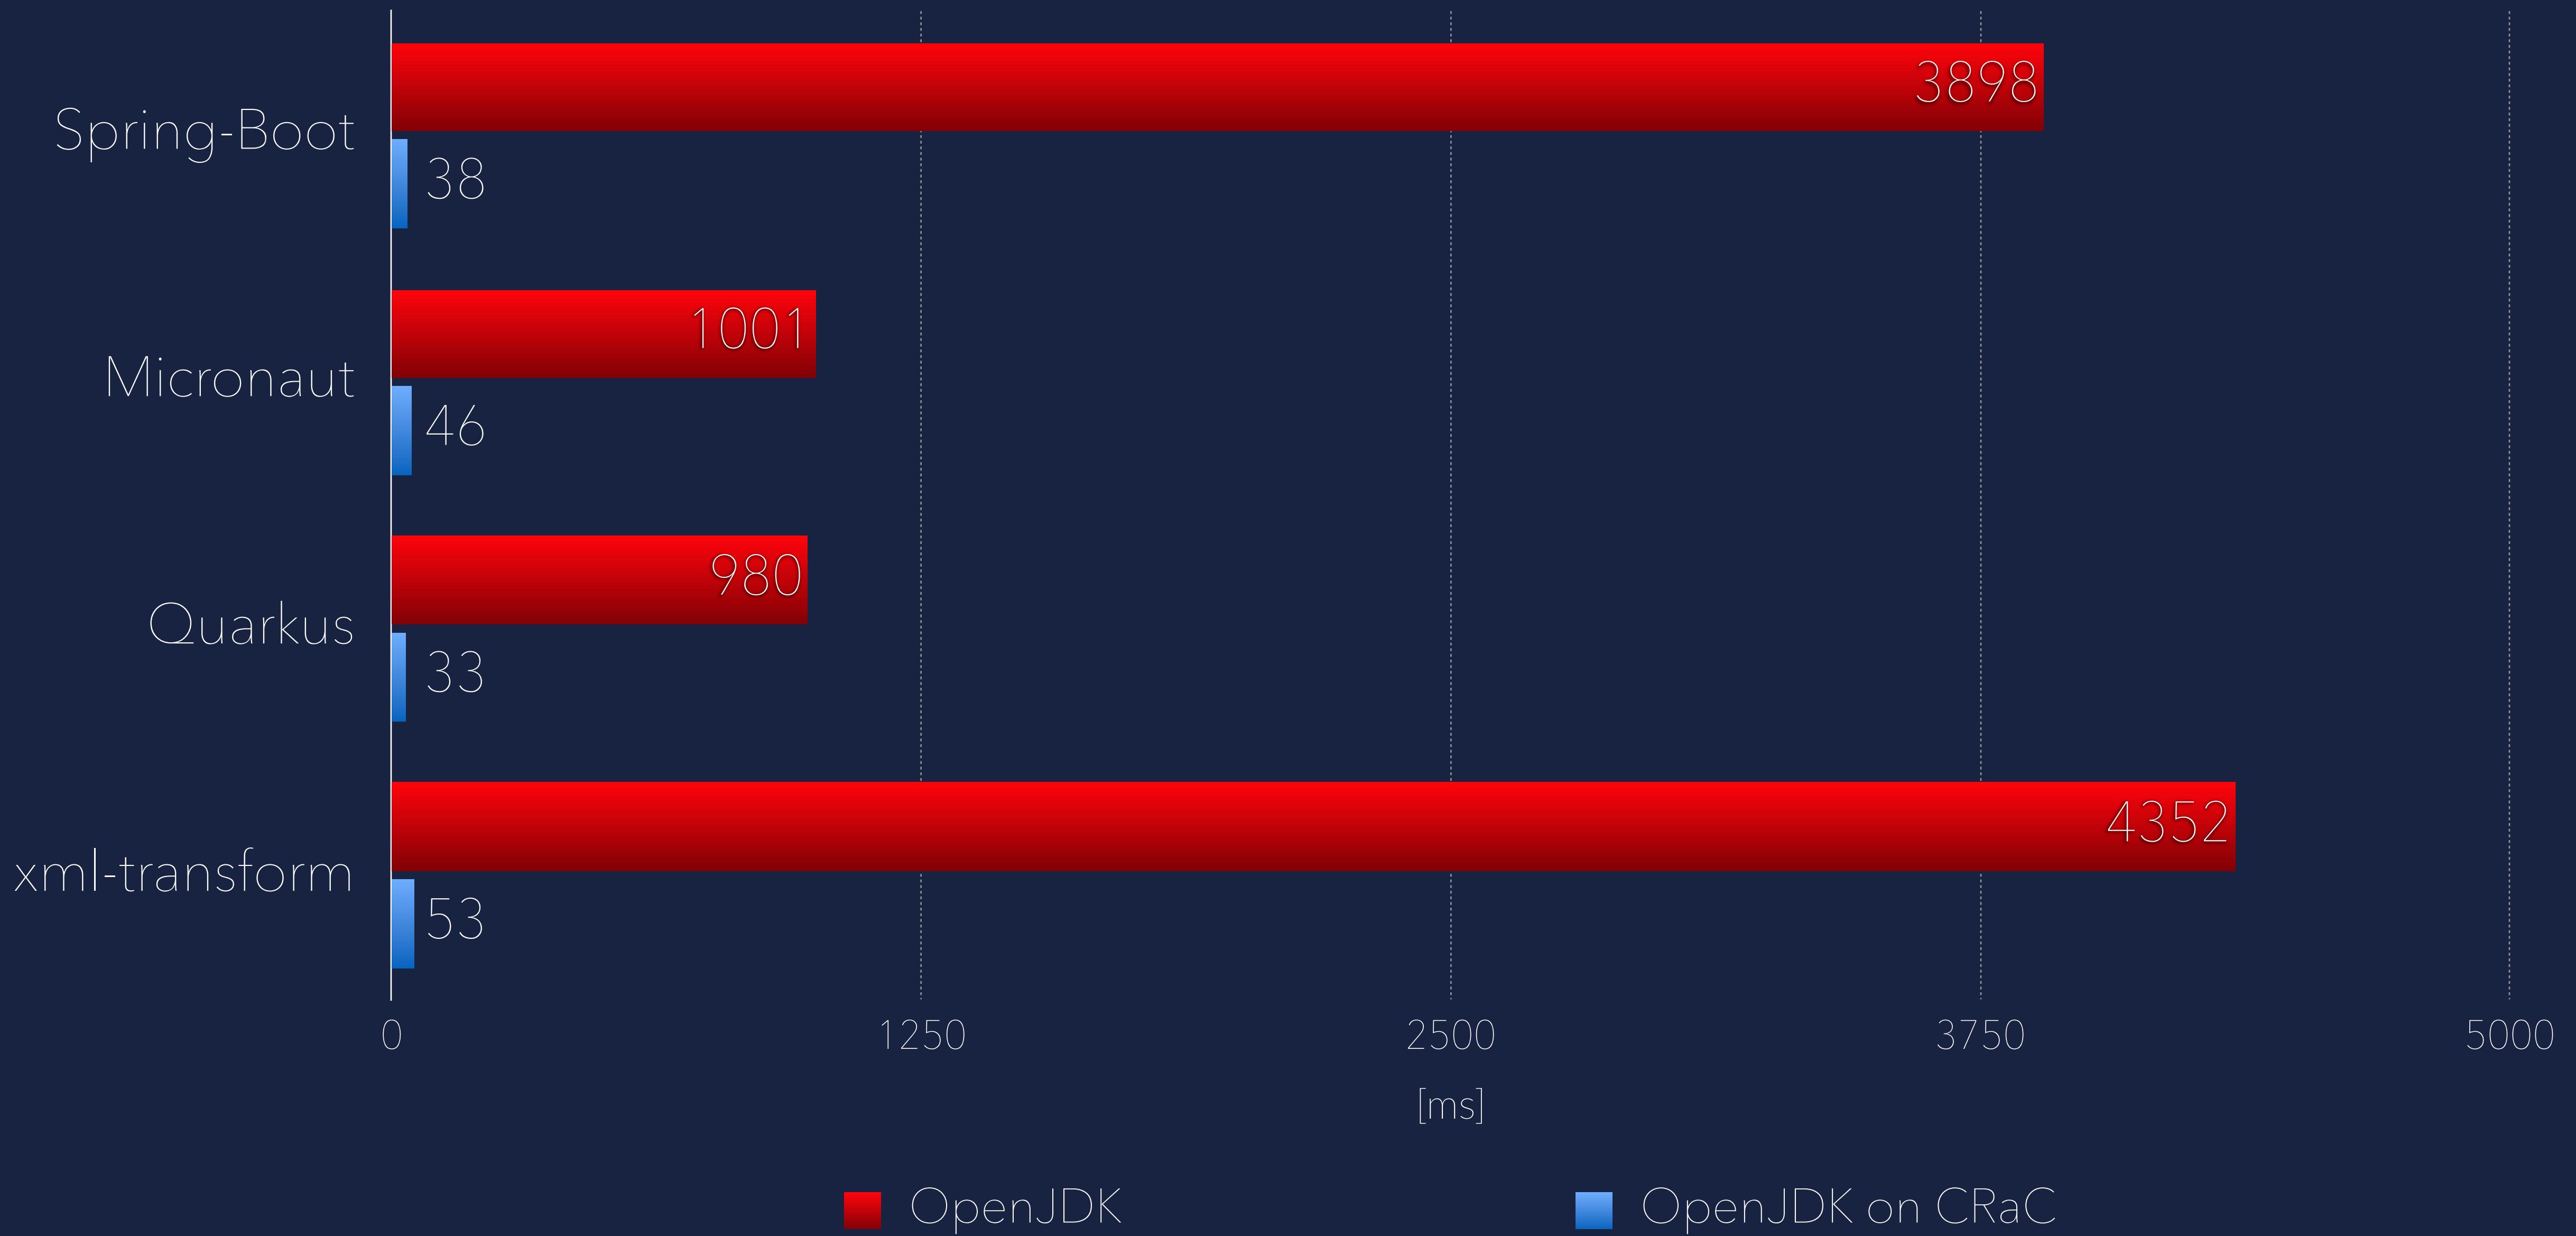
OK...BUT

HOW GOOD IS IT...?

Time to first operation



Time to first operation



SUMMARY

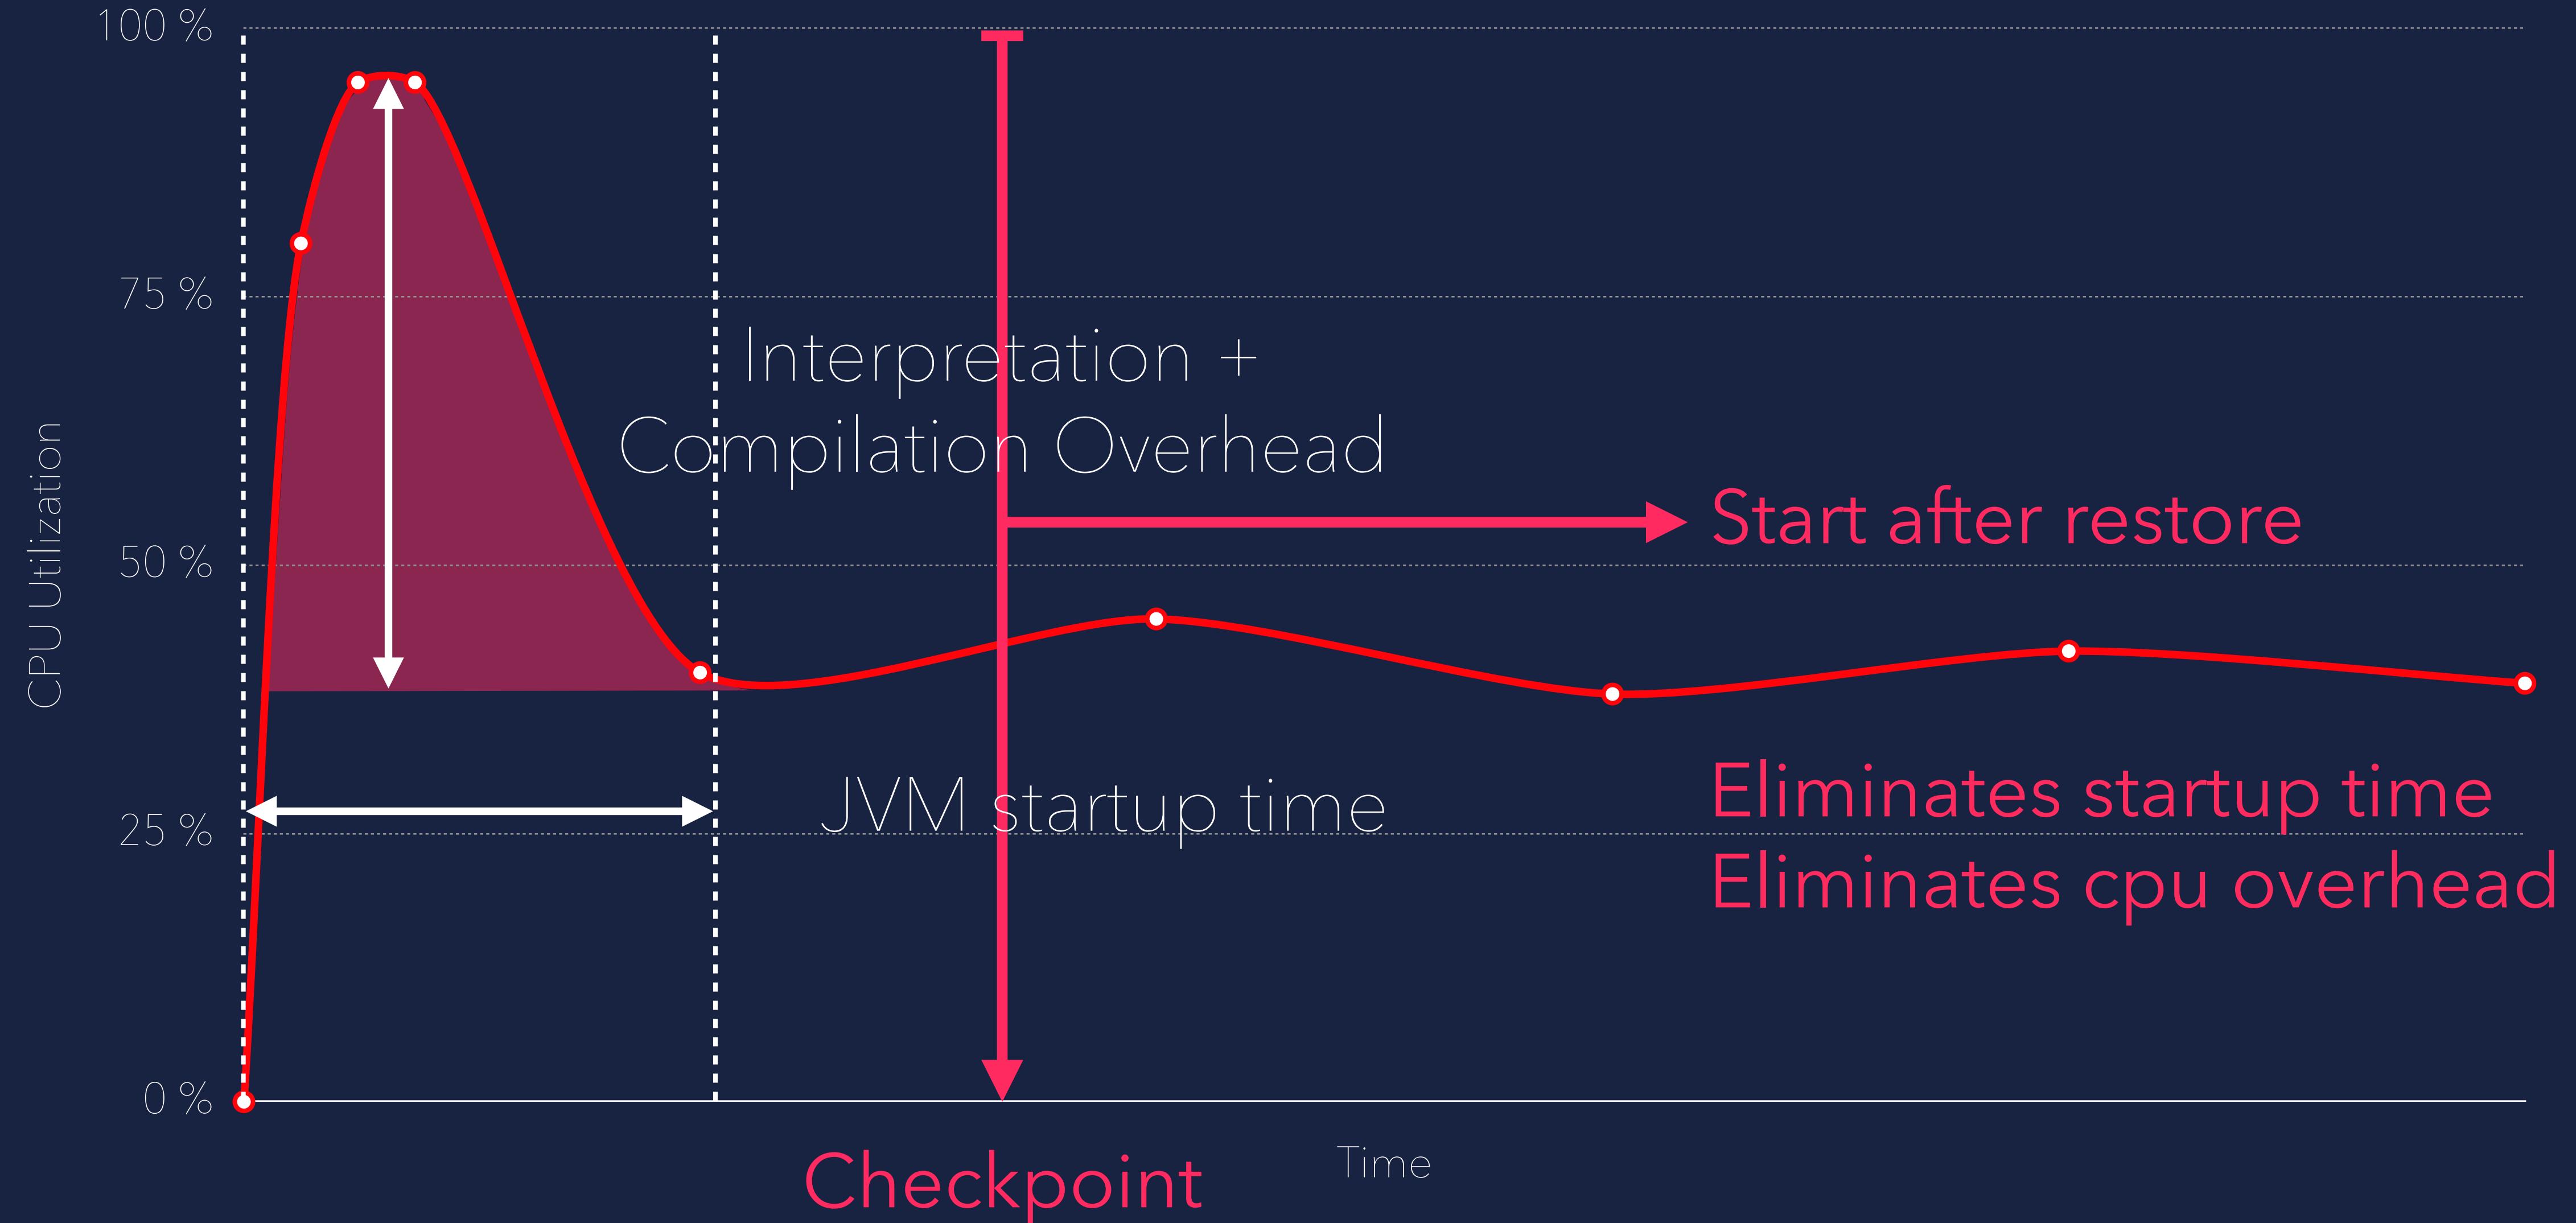


SUMMARY

ooo

- CRaC is a way to pause and restore a JVM based application
- It doesn't require a closed world as with a native image
- Benefit is potentially extremely fast time to full performance level
- Eliminates the need for hotspot identification, method compiles, recompiles and deoptimisations
- Improved throughput from start
- CRaC is an OpenJDK project
- CRaC can save infrastructure cost

INFRASTRUCTURE COST



github.com/CRaC



wiki.openjdk.org/display/cracker



DEMO

THANK



YOU